Phase mixing against a passive-scalar cascade: the Hermite spectrum of kinetic reduced magnetohydrodynamic turbulence

A. Kanekar

Draft, November 13, 2026

Abstract

In kinetic reduced magnetohydrodynamics (KRMHD) the compressive fluctuations are a kinetic field carried by the Alfvénic turbulence. Phase mixing moves their free energy to fine structure in parallel velocity and the turbulence moves it to small perpendicular scales, and the Hermite spectrum $W(m)$ records the competition between the two. We show that the compressive sector is exactly a continuum of passive scalars, one for each parallel phase velocity, each carried by the $\boldsymbol{E} \times \boldsymbol{B}$ flow and streaming along the perturbed magnetic field at its own speed. The nonlinearity therefore moves no free energy between phase velocities and, while the field lines stay frozen into the flow, none between field-aligned wavenumbers. We propose that its only effect on $W(m)$ is to remove free energy from every Hermite moment at the same rate, the rate at which the turbulence carries a passive scalar to the dissipation scale, while phase mixing proceeds at the field-aligned wavenumber at which the free energy was injected. This gives $W(m) \propto m^{-1/2}$ at small m and a decay exponential in $m^{1/2}$ beyond a break at $m_* \sim \beta_i/\chi^2$, where χ is the nonlinearity parameter. In simulations of forced KRMHD turbulence at $\beta_i = 1$, the model, with its two inputs measured independently, reproduces $W(m)$ and the Hermite flux along the field lines, and there is no return flux from high to low m; a backward flux appears only when the parallel wavenumber is measured along z rather than along the field. Lowering β_i moves the break into the resolved range, as predicted. With the Hermite resolution raised until the break is resolved, m_* grows about in proportion to β_i for $0.05 \leq \beta_i \leq 0.3$. Against χ, at $\beta_i = 0.05$, it grows faster than χ^{-2}, about as χ^{-4} below $\chi \approx 0.5$, because the perpendicular sink per outer eddy time weakens as the turbulence becomes weak. The model is too simple in one respect: the perpendicular sink is not the same at every m but grows with the time since the free energy was injected, in the same way at every β_i, so at $\beta_i = 0.3$ and 0.1 the spectrum lies above the model. The time in which the sink removes the free energy has not converged with the perpendicular resolution: it grows by about the same amount with each doubling up to 256^3 grid points, so whether the break becomes independent of the dissipation scale, as the model assumes, remains open.

1 Introduction

Kinetic reduced magnetohydrodynamics (KRMHD), the limit $k_\perp \rho_i \to 0$ of gyrokinetics derived by Schekochihin et al. (2009), hereafter S09,[1] describes anisotropic fluctuations at scales much larger than the ion gyroradius ρ_i in a strongly magnetised, weakly collisional plasma. Its Alfvénic fluctuations obey reduced magnetohydrodynamics (RMHD). Its compressive fluctuations, the slow waves and the entropy mode, obey a kinetic equation for the ion distribution in the parallel velocity

[1] Appendix E compares our notation with that of S09 and of Chandran et al. (2026).

$v_\parallel$; they are carried by the Alfvénic fluctuations and do not act back on them. Without collisions they are damped by phase mixing (Barnes, 1966; S09, §6.2), which moves their free energy to ever finer structure in $v_\parallel$, that is, to large Hermite order m in a Hermite expansion in $v_\parallel$ (Zocco and Schekochihin, 2011; Kanekar et al., 2015; Schekochihin et al., 2016). In turbulence the same free energy is also carried to small perpendicular scales. Which of the two routes it takes decides whether the compressive cascade heats ions through Landau damping or reaches the ion gyroscale (S09; Schekochihin et al., 2016; Meyrand et al., 2019; Kawazura et al., 2019, 2020), and the Hermite spectrum $W(m)$ is the most direct record of that competition.

Two pictures of $W(m)$ in turbulence have been proposed. In the first, nonlinear advection couples fluctuations with parallel wavenumbers of opposite sign and so returns free energy from high to low m, as a stochastic plasma echo (Kanekar, 2015; Schekochihin et al., 2016; Parker et al., 2016; Adkins and Schekochihin, 2018). Meyrand et al. (2019) found in KRMHD turbulence at $\beta_i = 1$ a spectrum somewhat steeper than m^{-1}, against $m^{-1/2}$ for linear phase mixing, and a net flux to high m much smaller than the injection rate away from the forcing scales, and attributed both to such echoes. In gyrokinetic simulations Kawazura et al. (2019, 2020) found $W(m) \propto m^{-1/2}$ at high β_i and close to m^{-1} at $\beta_i = 0.1$, and read the steeper spectrum as phase mixing suppressed by echoes. In the second picture (S09, §6.3), the compressive fluctuations have no parallel cascade: the field-aligned wavenumber at which they were injected stays fixed while the Alfvén waves carry them to smaller perpendicular scales, so at small enough scales the perpendicular cascade outruns the damping. The two pictures differ on whether the compressive fluctuations acquire structure along the moving field lines. Meyrand et al. (2019) argued that they do, because resistive reconnection relabels the field lines within about one eddy-turnover time and the compressive fluctuations then inherit the parallel cascade of the Alfvénic ones.

We make the second picture quantitative and test it. Section 3 shows that the compressive sector of KRMHD is exactly a continuum of passive scalars, or slices, labelled by the parallel phase velocity u, each carried by the $\boldsymbol{E} \times \boldsymbol{B}$ flow and streaming along the perturbed field at its own speed. The nonlinearity therefore moves no free energy between phase velocities, and free energy moves between Hermite moments only as the correlations between slices change (section 4). When the induction equation is ideal, the field lines are frozen into the flow that carries all the slices, and every volume-integrated single-point statistic, including $W(m)$, evolves as in linear phase mixing along the initial field lines (section 5). In turbulence the slices are also cascaded to small perpendicular scales and dissipated there. Section 6 proposes that this is the only effect of the nonlinearity on $W(m)$: every Hermite moment loses free energy at the same rate $1/T$, the rate at which the turbulence carries a passive scalar to the dissipation scale, while phase mixing proceeds at the field-aligned wavenumber κ of the injected free energy. The steady state is

$$W(m) = \frac{\varepsilon}{s\kappa\sqrt{2m}} \exp\Big(-\frac{\sqrt{2m}}{s\kappa T}\Big), \qquad m_* = \frac{(s\kappa T)^2}{2} \sim \frac{\beta_i}{\chi^2}, \tag{1}$$

for injection rate ε, $s = v_{\rm th}/v_A = \beta_i^{1/2}$ and nonlinearity parameter χ, summed over the distribution of κ. Two times set the spectrum: free energy injected at κ phase mixes up to moment m in the time $\sqrt{2m}/(s\kappa)$, which is shorter at higher β_i, and the perpendicular cascade removes it in the time T, which is set by the Alfvénic turbulence. The exponential in (1) is the fraction of the free energy that survives the sink while it phase mixes up to m, and m_* is the moment it reaches in the time T. Every quantity in (1) can be measured independently of $W(m)$, so the model has no free parameter. It predicts that the perpendicular dissipation of each Hermite moment is proportional to $W(m)$, that $W(m)$ and the Hermite flux computed along the field follow the model with no deficit at large m, that at fixed χ the break m_* grows with β_i, in proportion to it if T does not depend on β_i, and

that diagnostics built on the Eulerian k_z report a backward flux that the field-aligned flux does not have.

Section 7 tests these predictions in forced KRMHD turbulence with the GANDALF code (Kanekar, 2025). At $\beta_i = 1$ and $\chi \approx 1$ the perpendicular dissipation per Hermite moment is proportional to $W(m)$ to within 20%, the model reproduces $W(m)$ to within 8% for $m \leq 30$ and the field-aligned flux to within 7% for $m \leq 24$, and the flux split by the sign of k_z has a backward part of about 12% of the forward part. At $\beta_i = 1$ the break m_* lies beyond the $M = 32$ moments we resolve. At $\beta_i = 0.3$ and 0.1, with the same Alfvénic turbulence, it moves into the resolved range as predicted, and $W(m)$ steepens accordingly. There the measured spectrum lies above the model at every m, whereas a return flux would put it below the model at large m. The excess comes from the assumption of a single time T. The sink switches on over about five outer eddy-turnover times, as the cascade carries fresh free energy to the dissipation scale, so it removes less of the free energy at small m, which phase mixing brought there more recently. At lower β_i phase mixing is slower and the free energy at a given m is older, so the sink differs between the runs at a given m but is the same function of the time since injection in all of them. Doubling the number of Hermite moments leaves $W(m)$ unchanged where both resolutions resolve it. Raising the perpendicular resolution does not: up to 256^3 grid points T grows by about the same amount with each doubling, and the field-aligned wavenumber of g grows slowly through the longer inertial range, so the model's two approximations, a sink independent of the dissipation and a frozen wavenumber, hold only roughly at the resolutions we reach. With the T measured at each resolution the model still describes $W(m)$, and there is still no return flux. With the parallel flow forced in place of the density, as by Meyrand et al. (2019), the spectrum steepens at small m, by the amount the model gives with the injected distribution of κ. With the Hermite resolution raised until the break is resolved, m_* grows as $\beta_i^{1.09\pm0.04}$ for $0.05 \leq \beta_i \leq 0.3$ at $\chi \approx 1$. At $\beta_i = 0.05$ it grows somewhat faster than χ^{-2} down to $\chi \approx 0.56$ and about as χ^{-4} between $\chi = 0.56$ and 0.40. In units of the outer eddy time the sink weakens at the smallest χ, so the estimate $m_* \sim \beta_i/\chi^2$, which takes T proportional to the eddy time, holds only near $\chi \approx 1$.

The ingredients of (1) are not new, and we credit them where they are used. The rate comparison between Landau damping at a frozen parallel wavenumber and the perpendicular cascade is S09's (§6.3, eqs 217–220). The functional form is that of Zocco and Schekochihin (2011, eq. 112) for growing modes, with their growth rate replaced by the cascade rate, and Schekochihin et al. (2016, eqs 2.58–2.59) solved the same continuity equation in m with a collisional sink. The break scale is the threshold $k_\parallel v_{\rm th} \sim \sqrt{m}\, k_\perp u_\perp$ of Schekochihin et al. (2016, eq. 4.5), and Kawazura et al. (2019) stated $m \sim \beta_i$ for critically balanced turbulence. A Hermite spectrum decaying exponentially in $\sqrt{m}$, at a rate proportional to the ratio of the cascade rate to $k_\parallel v_{\rm th}$, was seen in simulations of a kinetic passive scalar advected by a two-dimensional flow (Kanekar, 2015, chapter 4). What is new here is the exact statement that in three-dimensional KRMHD with ideal induction the nonlinearity moves no free energy in m, the model with independently measured inputs that follows from it, and its test in turbulence with diagnostics computed along the field. The slice decomposition itself has precedents in the G-transform of Morrison and Pfirsch (1992) and Morrison (2000), its application to drift kinetics by Heninger and Morrison (2018), and the per-velocity invariants of Plunk and Parker (2014) and Plunk (2015).

Section 2 sets out the compressive sector of KRMHD, sections 3 to 5 give the exact results, section 6 the model and its predictions, section 7 the simulations, and section 8 discusses what the results do and do not establish.

2 The compressive sector of KRMHD

For a plasma with one ion species and isothermal electrons, KRMHD (Schekochihin et al., 2009) consists of reduced magnetohydrodynamics (RMHD) for the Alfvénic fluctuations and a kinetic equation for the compressive ones. We use the form of CMM (Chandran et al., 2026), whose section 3.2 and appendix B restate and extend S09; their equation numbers refer to arXiv:2607.27981v1. Appendix E compares our notation with that of CMM, S09 and the GANDALF code (Kanekar, 2025).

2.1 Alfvénic fields

We set the Alfvén speed $v_A = 1$ unless stated otherwise and take the mean field $\boldsymbol{B}_0$ along $\hat{\boldsymbol{z}}$. The stream function Φ and the flux function Ψ give the $\boldsymbol{E}\times\boldsymbol{B}$ velocity $\boldsymbol{u}_E = \hat{\boldsymbol{z}}\times\nabla_\perp\Phi$ and the perpendicular magnetic fluctuation $\boldsymbol{b}_\perp \equiv \delta\boldsymbol{B}_\perp/(4\pi n_i m_i)^{1/2} = \hat{\boldsymbol{z}}\times\nabla_\perp\Psi$, and to first order the unit vector along the field is $\hat{\boldsymbol{b}} = \hat{\boldsymbol{z}} + \boldsymbol{b}_\perp$; both $\boldsymbol{u}_E$ and $\hat{\boldsymbol{b}}$ are divergence-free. With the bracket $\{f,h\} = \partial_x f\,\partial_y h - \partial_y f\,\partial_x h = \hat{\boldsymbol{z}}\cdot(\nabla f\times\nabla h)$, advection by $\boldsymbol{u}_E$ is $\{\Phi,\cdot\}$ and the derivative along the field is $\nabla_\parallel \equiv \hat{\boldsymbol{b}}\cdot\nabla = \partial_z + \{\Psi,\cdot\}$. The RMHD equations are (Schekochihin et al., 2009, eqs 17–18)

$$\partial_t\Psi + \{\Phi,\Psi\} = \partial_z\Phi, \qquad \partial_t\nabla_\perp^2\Phi + \{\Phi,\nabla_\perp^2\Phi\} = \nabla_\parallel\nabla_\perp^2\Psi. \tag{2}$$

The first, the induction equation $\partial_t\Psi = \nabla_\parallel\Phi$, states that the electric field along the total magnetic field vanishes. Equations (2) are CMM's eqs (3.23)–(3.24) in their normalised potentials (3.2) and do not contain the compressive fluctuations. The Elsasser potentials $z^\pm = \Phi \pm \Psi$ (S09's $\zeta^\pm$) obey $\partial_t z^\pm = \pm\partial_z z^\pm$ linearly, so z^+ propagates toward $-z$ and z^- toward $+z$.[2]

2.2 Kinetic equation and the two branches

The ion thermal speed is $v_{\rm th} = (2T_i/m_i)^{1/2}$, so $s \equiv v_{\rm th}/v_A = \beta_i^{1/2}$; we write $\tau = T_i/T_e$ and $Z = q_i/e$. The compressive part of the ion distribution is the function g of S09 (CMM's g_{09}, their eq. 3.26), which obeys CMM's eq. (3.25) (S09 eq. 157 without collisions),

$$\frac{\mathrm{d}g}{\mathrm{d}t} + v_\parallel\nabla_\parallel\Big[g + \Big(\frac{Z}{\tau}\frac{\delta n}{n} + \frac{v_\perp^2}{v_{\rm th}^2}\frac{\delta B}{B}\Big)F_0\Big] = 0, \qquad \frac{\mathrm{d}}{\mathrm{d}t} = \partial_t + \{\Phi,\cdot\}, \tag{3}$$

where F_0 is the three-dimensional Maxwellian. Quasineutrality and perpendicular pressure balance give (CMM eqs 3.27–3.28; S09 eqs 158–159),

$$\overline{g} = \frac{\delta n}{n} - \frac{\delta B}{B}, \qquad \overline{(v_\perp^2/v_{\rm th}^2)\,g} = -\frac{Z}{\tau}\frac{\delta n}{n} - 2\Big(1+\frac{1}{\beta_i}\Big)\frac{\delta B}{B}, \qquad \overline{(\cdots)} \equiv \frac{1}{n_i}\int \mathrm{d}^3v\,(\cdots). \tag{4}$$

These relations are algebraic: $\delta n/n$ and $\delta B/B$ at a point are linear functionals of g at that point, with coefficients that do not depend on the wavenumber. At finite $k_\perp\rho_i$ they contain gyroaveraging operators.

We expand g in Laguerre polynomials of $v_\perp^2/v_{\rm th}^2$ with coefficients $G_l(v_\parallel)$ (S09 eq. 205; CMM eq. B8; appendix A). The source term in (3) is linear in $v_\perp^2$, so the coefficients with $l \geq 2$ obey $\mathrm{d}G_l/\mathrm{d}t + v_\parallel\nabla_\parallel G_l = 0$ (CMM eq. B10; S09 eqs 210 and 212). The coefficients with $l = 0, 1$ are coupled through $\delta n/n$ and $\delta B/B$; the combinations $G^\pm$ of CMM's eqs (B5)–(B6) (S09 eqs 183–184) obey CMM's eq. (B4) (S09 eq. 181),

$$\frac{\mathrm{d}G^\pm}{\mathrm{d}t} + v_\parallel\nabla_\parallel G^\pm = \frac{v_\parallel F_M(v_\parallel)}{\Lambda^\pm}\,\nabla_\parallel\int G^\pm(v_\parallel')\,\mathrm{d}v_\parallel', \tag{5}$$

[2] S09's §§2.3 and 2.5 write the solutions with $\zeta^\mp = 0$ as $\zeta^\pm = f^\pm(x,y,z\mp v_A t)$, the opposite direction. We checked ours against S09's eq. (21), CMM's equations and GANDALF's right-hand side.

where $F_M(v_\parallel) = \pi^{-1/2} v_{\rm th}^{-1} e^{-v_\parallel^2/v_{\rm th}^2}$ and

$$\Lambda^\pm = -\frac{\tau}{Z} + \frac{1}{\beta_i} \pm \Delta, \qquad \Delta = \Big[\Big(1+\frac{\tau}{Z}\Big)^2 + \frac{1}{\beta_i^2}\Big]^{1/2} \tag{6}$$

(CMM eq. B7; S09 eqs 182 and 204).[3] We call the right-hand side of (5), of strength $1/\Lambda^\pm$, the Landau coupling; in linear theory it produces the collisionless damping of the compressive fluctuations (S09, §6.2.2). The coefficients with $l \geq 2$ obey (5) with $1/\Lambda = 0$, and we call them the sector $\Lambda = \infty$. The compressive sector therefore consists of independent kinetic equations for G^+, G^- and each G_l with $l \geq 2$, all advected by the same Alfvénic fields. Since $\Delta > 1 + \tau/Z$ and $\Delta > 1/\beta_i$, $\Lambda^+ - 1 > \beta_i^{-1}$ and $\Lambda^- < -\tau/Z$, so $\Lambda^+ > 1$ and $\Lambda^- < 0$ for all $\beta_i, \tau/Z > 0$ (S09, below eq. 200); at $\beta_i = \tau = Z = 1$, $\Lambda^\pm = \pm\sqrt{5}$.

2.3 One branch in the variable u

We now drop the branch label, set $u = v_\parallel/v_{\rm th}$, and write $F_M(u) = \pi^{-1/2} e^{-u^2}$ and $M_k[G] = \int u^k G \, {\rm d}u$. Since $v_\parallel = su$ when $v_A = 1$, equation (5) becomes

$$\partial_t G + \{\Phi, G\} + s\,\nabla_\parallel (LG) = 0, \tag{7}$$

with the branch operator

$$LG = uG - \frac{uF_M}{\Lambda} M_0[G], \tag{8}$$

which acts on the velocity variable only; for $\Lambda = \infty$ it is multiplication by u. Because the field relations (4) are algebraic, L contains no wavenumber.

For given Φ and Ψ, equation (7) is linear in G. Its nonlinear terms are the advection $\{\Phi, G\}$ and the part $s\{\Psi, LG\}$ of the streaming, which follows the perturbed field lines. The Alfvénic fields obey (2) whatever G does, so the compressive sector is passive. At finite $k_\perp \rho_i$ this fails; the induction equation, for example, contains $\nabla_\parallel(\delta n/n)$ (CMM, §3.1).

The density and field-strength fluctuations are linear combinations of $M_0[G^+]$ and $M_0[G^-]$ (A3), the parallel ion flow involves $M_1[G^\pm]$, and the pressure perturbations involve $M_0[G^\pm]$ and $M_2[G^\pm]$ (appendix A); the sector $\Lambda = \infty$ contributes to none of these moments.

2.4 Hermite representation

Following CMM (eqs B27 and B30), we expand $G = \sum_{m\geq 0} g_m H_m(u) F_M(u)/\sqrt{2^m m!}$, so that $g_m = \int H_m \, G \, {\rm d}u/\sqrt{2^m m!}$ and $G/F_M^{1/2} = \sum_m g_m \psi_m(u)$ with the orthonormal Hermite functions ψ_m. The g_m are CMM's $\tilde{G}_m$, and $M_0[G] = g_0$, $M_1[G] = g_1/\sqrt{2}$. In this basis (7) is the hierarchy

$$\partial_t g + \{\Phi, g\} + s\,\nabla_\parallel (Ag) = 0, \qquad g = (g_0, g_1, g_2, \dots)^{\mathsf T}, \tag{9}$$

where the ladder matrix A, the matrix of L, is tridiagonal with zero diagonal,

$$A_{m,m+1} = A_{m+1,m} = \sqrt{\frac{m+1}{2}} \quad \text{except} \quad A_{10} = \frac{1 - 1/\Lambda}{\sqrt{2}}. \tag{10}$$

The weight matrix $\mathsf{P} = {\rm diag}(p_m)$, with $p_0 = 1 - 1/\Lambda$ and $p_m = 1$ for $m \geq 1$, symmetrises it, $\mathsf{P}A = (\mathsf{P}A)^{\mathsf T}$, and $\sum_m p_m g_m^2$ is the free-energy product $[G, G]$ of section 3. A truncation at $m = M$

[3] CMM's eq. (B16) prints the square root in (6) as $[(1+\tau/Z) + \beta_i^{-2}]^{1/2}$, without the square on $1+\tau/Z$; with the square, as in S09's eq. (204), their decompositions (B15) and (B17) hold, and without it they do not (appendix A).

needs a closure; by default we use the zero closure $g_{M+1} = 0$, which keeps the first $M+1$ rows and columns of A. GANDALF (Kanekar, 2025) integrates (2) in Elsasser form together with (9) for one branch, with optional dissipation, hypercollisions and forcing, truncated with the zero closure or with the copy closure $g_{M+1} = g_{M-1}$ (section 4.4). We checked to round-off that its right-hand side without dissipation and hypercollisions is (9)–(10), for both signs of Λ and both closures.

The rest of the paper concerns one equation (7) with $\Lambda \notin [0, 1]$, which we call a branch whether it describes $G^{\pm}$ or a Laguerre coefficient with $l \geq 2$. Sections 3 to 5 treat a collisionless and unforced branch except where stated; sections 6 and 7 add forcing, dissipation and hypercollisions. The Hermite spectrum of the free energy is $W(m) = p_m \langle g_m^2 \rangle$, where $\langle \cdot \rangle$ is the average over the periodic box, so that $\sum_m W(m) = \langle [G, G] \rangle$.

3 Phase-velocity slices

This section shows that each branch (7) is a continuum of passive scalars, or slices, labelled by the parallel phase velocity. The Alfvénic fields Φ and Ψ are given functions of $\boldsymbol{r}$ and t; (2) is not used.

3.1 Without the Landau coupling

For $\Lambda = \infty$ the operator L is multiplication by u, and (7) holds separately at every u:

$$\partial_t G + \{\Phi, G\} + su\, \nabla_\parallel G = 0.$$

Since $\nabla_\parallel = \partial_z + \{\Psi, \cdot\}$ and $\{f, G\} = (\hat{\boldsymbol{z}} \times \nabla_\perp f) \cdot \nabla G$, this is the advection equation $\partial_t G + \boldsymbol{U}_u \cdot \nabla G = 0$ with the velocity

$$\boldsymbol{U}_u = \hat{\boldsymbol{z}} \times \nabla_\perp (\Phi + su\Psi) + su\, \hat{\boldsymbol{z}} = \boldsymbol{u}_E + su\, \hat{\boldsymbol{b}}, \qquad \nabla \cdot \boldsymbol{U}_u = 0. \tag{11}$$

At each u, G is a passive scalar, carried by the $\boldsymbol{E} \times \boldsymbol{B}$ flow and streaming along the perturbed field at the speed su, which is the motion of the guiding centre of a free ion with $v_\parallel = u v_{\mathrm{th}}$. Values of G at different u do not interact. This is the exact equation for the Laguerre coefficients G_l with $l \geq 2$ (S09, eqs 210 and 212).

3.2 With the Landau coupling

For finite Λ, $LG = uG - (uF_M/\Lambda)\, M_0[G]$ ties every velocity to the density $M_0[G]$, through the parallel electric field and the mirror force, so G at fixed u is no longer a passive scalar. The slices survive in other variables. The reason is that L acts on velocity only and is the same at every point and at every wavenumber, because the field relations (4) are algebraic. A change of velocity variables that diagonalises L therefore does so everywhere at once, and it passes through the advection and the streaming unchanged.

The argument is clearest for the hierarchy (9) truncated at $m = M$ with the zero closure, in which L is the constant $(M+1) \times (M+1)$ matrix A. Let $\boldsymbol{w}_j$ be a left eigenvector of A, $\boldsymbol{w}_j^{\mathsf{T}} A = u_j \boldsymbol{w}_j^{\mathsf{T}}$. Since $\boldsymbol{w}_j$ does not depend on position or time, it passes through ∂_t, ∂_z and both brackets, and multiplying (9) by $\boldsymbol{w}_j^{\mathsf{T}}$ gives for $a_j = \boldsymbol{w}_j^{\mathsf{T}} g$

$$\partial_t a_j + \{\Phi, a_j\} + su_j\, \nabla_\parallel a_j = \partial_t a_j + \{\Phi + su_j \Psi,\, a_j\} + su_j\, \partial_z a_j = 0.$$

Each a_j is a passive scalar, advected with the stream function $\Phi + su_j\Psi$ and streaming along z at the speed su_j, as G at fixed u was for $\Lambda = \infty$, and nothing couples a_j to the other a_l. The

nonlinearity enters only through the bracket with $\Phi + s u_j \Psi$, and the Landau coupling does not appear, because it has been absorbed into $\boldsymbol{w}_j$. For the a_j to be a change of variables, the u_j must be real and the $\boldsymbol{w}_j$ linearly independent, so that the a_j determine g. Both hold because $\mathsf{P}A$ is symmetric (section 2.4): A is self-adjoint in the free-energy product $\sum_m p_m g_m h_m$, which is positive definite when $p_0 = 1 - 1/\Lambda > 0$, that is, for $\Lambda \notin [0, 1]$. Section 4.4 normalises the a_j and treats both closures.

At $M = \infty$ the same steps need the eigenfunctions of L itself. Self-adjointness carries over: L is self-adjoint in the free-energy product

$$[f, g] = \int \frac{f g}{F_M} \, \mathrm{d}u - \frac{M_0[f]\, M_0[g]}{\Lambda}, \tag{12}$$

which is $\sum_m p_m f_m g_m$ in the Hermite variables and is positive definite for $\Lambda \notin [0, 1]$, as for every compressive branch of KRMHD ($\Lambda^+ > 1$, $\Lambda^- < 0$; section 2). Unlike A, however, L has no eigenfunctions of finite free energy: its spectrum is the whole real line and is continuous. Its generalised eigenfunctions are the Case–Van Kampen modes F_u of the branch (van Kampen, 1955; Case, 1959), with $L F_u = u F_u$ for every real u. Each is a beam of particles at $v_\parallel = u v_{\mathrm{th}}$, a δ-function in velocity, dressed by the response of the Maxwellian, through the Landau coupling, to the density of the beam (appendix B). The projection $\mathcal{T}[G](u) = [F_u, G]$ on F_u takes the place of $\boldsymbol{w}_j^{\mathsf{T}} g$, and its unit-weight form is the slice amplitude

$$a_u = \sqrt{\varrho(u)}\, [F_u, G], \qquad \varrho(u) = \frac{F_M(u)}{|\mathcal{K}(u + \mathrm{i}0)|^2}, \tag{13}$$

where $\mathcal{K}$ is the dielectric function of the branch, which contains the plasma dispersion function, and ϱ is the spectral density of the branch, positive with $\int \varrho \, \mathrm{d}u = 1/p_0$. For $\Lambda = \infty$, F_u is the bare beam, $\mathcal{K} = 1$, $\varrho = F_M$ and $a_u = G/\sqrt{F_M}$. The map $G \mapsto a_u$ is a fixed linear change of variables in velocity. It is invertible, divides the free energy among phase velocities without cross terms, and turns L into multiplication by u:

$$[G, G] = \int a_u^2 \, \mathrm{d}u, \qquad G \mapsto a_u \ \text{ takes } \ LG \mapsto u\, a_u \tag{14}$$

(theorem 4, which also gives a_u in terms of G and its Hilbert transform in u). Its inverse is, up to normalisation, the G-transform that Morrison and Pfirsch (1992) and Morrison (2000) introduced for linearised Vlasov–Poisson; Heninger and Morrison (2018) applied it to electrostatic drift kinetics, where the transformed components remain coupled through the potential.

3.3 The slice theorem

Theorem 1. *Assume that (i) $\Lambda \notin [0, 1]$, the value $\Lambda = \infty$ included; (ii) $\Phi(\boldsymbol{r}, t)$ and $\Psi(\boldsymbol{r}, t)$ are real functions on $\mathbb{T}^3 \times [0, t_1]$, with $\mathbb{T}^3$ a periodic box, whose gradients are continuous and Lipschitz in $\boldsymbol{r}$, uniformly in t; they need not satisfy (2) and may obey resistive or forced RMHD, or be prescribed; (iii) G and uG are continuously differentiable functions of $(\boldsymbol{r}, t)$ with values in the space $\mathcal{W}$ of functions of u with finite $\int G^2 F_M^{-1} \, \mathrm{d}u$ (appendix B). Then G solves (7) if and only if, at every $(\boldsymbol{r}, t)$ and for almost every u,*

$$\partial_t a_u + \boldsymbol{U}_u \cdot \nabla a_u = \partial_t a_u + \{\Phi + s u \Psi,\, a_u\} + s u\, \partial_z a_u = 0, \tag{15}$$

with $\boldsymbol{U}_u$ of (11).

Proof. This is the argument for the truncated hierarchy in section 3.2, with the map $\mathcal{A} : G \mapsto a$ of (13) in place of $\boldsymbol{w}_j^{\mathsf{T}}$. The map $\mathcal{A}$ is bounded and injective (theorem 4) and acts on u only, through a kernel independent of $\boldsymbol{r}$ and t, so it commutes with ∂_t, with spatial derivatives and with multiplication by functions of $(\boldsymbol{r}, t)$, hence with $\{\Phi, \cdot\}$, $\{\Psi, \cdot\}$ and $\nabla_\parallel$. With $\mathcal{A}LG = u\,\mathcal{A}G$ this gives $\mathcal{A}[\partial_t G + \{\Phi, G\} + s\,\nabla_\parallel(LG)] = \partial_t a + \{\Phi, a\} + su\,\nabla_\parallel a$, and $u\,\nabla_\parallel a_u = u\,\partial_z a_u + \{u\Psi, a_u\}$ gives (15). Since $\mathcal{A}$ is linear and injective, G solves (7) if and only if a solves (15). □

With the Landau coupling, therefore, each slice is again a passive scalar carried by $\boldsymbol{U}_u$, in three dimensions, and slices with different u do not interact. The parallel electric field and the mirror force, which enter through the Landau coupling, only dress the slices: the distribution at fixed u is not a passive scalar, but its projection a_u on F_u is. Since $\boldsymbol{U}_u$ is divergence-free, each slice keeps its own free energy $\int a_u^2\,\mathrm{d}^3\boldsymbol{r}$, and by (14) the nonlinearity moves no free energy between phase velocities (section 4.1). In terms of the Elsasser potentials, the stream function of the slice at u is

$$\Phi + su\Psi = \tfrac{1}{2}\left[(1+su)\,z^+ + (1-su)\,z^-\right], \tag{16}$$

so the slices at $su = -1$ and $+1$ are advected by z^- alone and by z^+ alone, and the slice at $u = 0$, the distribution at $v_\parallel = 0$, by Φ alone. The sectors have different transforms, but $\boldsymbol{U}_u$ depends only on u, so at each u the slices of G^+, G^- and every G_l with $l \geq 2$ are carried by the same flow:

$$(\partial_t + \boldsymbol{U}_u \cdot \nabla)(a_u^+,\, a_u^-,\, a_u^{(2)},\, a_u^{(3)}, \dots) = 0. \tag{17}$$

The compressive sector of KRMHD is therefore a continuum of passive scalars labelled by phase velocity.

3.4 Dissipation, collisions and forcing

Let the right-hand side of (7) be $-\mathcal{B}G - \mathcal{C}G + F$, where the linear operators $\mathcal{B}$ and $\mathcal{C}$ act on $\boldsymbol{r}$ alone and on u alone, respectively, and F is a source. The proof of theorem 1 gives

$$\partial_t a_u + \boldsymbol{U}_u \cdot \nabla a_u = -\mathcal{B}\,a_u - \int \hat{\mathcal{C}}(u, u')\,a_{u'}\,\mathrm{d}u' + \sqrt{\varrho(u)}\;\mathcal{T}[F](u), \tag{18}$$

where $\hat{\mathcal{C}}$ is the kernel of $\mathcal{A}\,\mathcal{C}\,\mathcal{A}^{-1}$. Damping at a rate that depends only on $k_\perp$ or k_z, applied with the same coefficient to every Hermite moment, is of the form $\mathcal{B}$ and damps each slice separately. With $e_m = H_m(u)F_M(u)/\sqrt{2^m m!}$ and $r_m(u)$ the mth Hermite coefficient of F_u, a polynomial of degree m with $r_0 = 1$ and $r_1 = \sqrt{2}\,u$, one has $\mathcal{T}[e_m](u) = p_m\,r_m(u)$, and the functions $\phi_m = \sqrt{p_m \varrho}\;r_m$ form an orthonormal basis of $L^2(\mathbb{R}, \mathrm{d}u)$. For $\mathcal{C}e_m = \nu_m e_m$, as for hypercollisions $\nu_m = \nu(m/M)^n$, the kernel is $\hat{\mathcal{C}}(u, u') = \sum_m \nu_m\,\phi_m(u)\,\phi_m(u')$, which couples different slices unless all ν_m are equal; collisions are the only term in (18) that exchanges free energy between slices. Forcing of g_0 by a field $f(\boldsymbol{r}, t)$, $F = f\,e_0$, gives every slice the same spatial pattern, with amplitude $(1 - 1/\Lambda)\sqrt{\varrho(u)}\,f$.

3.5 Slices in a turbulent run

Figure 1 releases a density blob, $g(0) = f\,e_0$, into decaying Alfvénic turbulence with $\chi = 1.00$ and follows the 17 slices of the hierarchy truncated at $M = 16$. A density perturbation loads every slice with the same pattern f (section 3.4). Each slice is then carried by its own stream function $\Phi + su_j\Psi$, which at $su_j = -1.0218$, 0 and $+1.0218$ is mainly z^-, Φ alone and mainly z^+, and streams along z at su_j. The three slices end visibly different, and g_0 mixes them. The slices keep their values as they are carried: on the plane $z = 0$ the peaks of the slices at $su_j = \pm 1.0218$ are

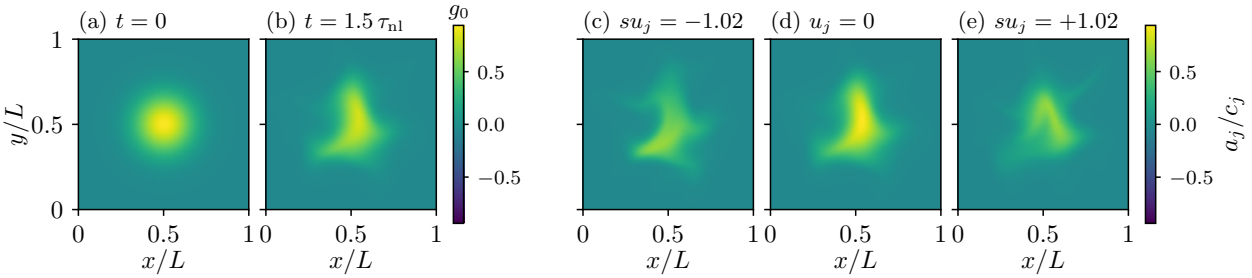


Figure 1: A density blob in decaying Alfvénic turbulence, on the plane $z = 0$ (coordinates in units of the box side $L = 1$). (a, b) The density moment g_0 at $t = 0$ and $t = 1.5\,\tau_{\rm nl}$. (c–e) Three of the 17 slices a_j of the run (section 3.2) at $t = 1.5\,\tau_{\rm nl}$, with $su_j = -1.0218$, $u_j = 0$ and $su_j = +1.0218$; their stream functions $\Phi + su_j\Psi$ are mainly z^-, Φ alone and mainly z^+. At $t = 0$ each slice is a constant multiple $c_j f$ of the blob f, and the panels show a_j/c_j, so all three started as panel (a). They differ at $t = 1.5\,\tau_{\rm nl}$ because each was carried by its own flow, and g_0 in (b) is a weighted sum of all 17 slices. Panels (c) and (e) are also fainter than (d), because those slices stream along z at su_j: the material on their plane $z = 0$ started on the planes $z = +0.21$ and $z = -0.21$, where f is about 0.7 of its value at $z = 0$. Each group of panels has one symmetric colour scale. GANDALF run on a 64^3 grid with $M = 16$ (zero closure), $\Lambda = \sqrt{5}$ and $\beta_i = 1$; balanced Gaussian random Elsasser fields with $1 \le k_\perp/k_1 \le 3$, $|k_z|/k_1 \le 2$ and $\chi = 1.00$; hyper-resistivity damps $z^\pm$ and every g_m; no forcing or hypercollisions; f is a Gaussian of standard deviation $1/8$ in x and y and $1/4$ in z; $\tau_{\rm nl} = 0.1357$ is the outer eddy time.

0.69 of the initial peak, close to the factor 0.71 by which f falls over their streaming distance of 0.21 along z, and the peak of the slice at $u_j = 0$ is 0.998 of it. GANDALF's $\partial_t g$ projected on the slices agrees with (26), computed with brackets that do not use GANDALF, to 5.2×10^{-16} of the largest slice rate, while advecting every slice by Φ alone gives relative errors of 0.30 to 1.46. Over the run each slice energy changes only by the hyper-resistive loss of that slice, to 1.2×10^{-6} of its initial energy, so no free energy passes between slices.

4 Free energy and the Hermite spectrum in slice variables

This section collects the consequences of theorem 1 for the free energy and its Hermite spectrum. Without collisions, dissipation or sources the energy of every slice is conserved, so the free energy is not moved between phase velocities, and free energy moves between Hermite moments only as the correlations between slices change.

4.1 Slice energies

Let $\mathcal{E}(u) = \langle a_u^2 \rangle$ be the slice energy, so that $\frac{1}{2} n_i T_i\, \mathcal{E}(u)\, \mathrm{d}u$ is the free energy per unit volume in the slices with phase velocities in $[u, u + \mathrm{d}u]$ and $\int \mathcal{E}\, \mathrm{d}u = \langle [G, G] \rangle = \sum_m W(m)$ by (14). Under the hypotheses of theorem 1, with no collisions, dissipation or sources acting on G, (11) and $\nabla \cdot \boldsymbol{U}_u = 0$ give $\partial_t a_u^2 + \nabla \cdot (\boldsymbol{U}_u\, a_u^2) = 0$, so $\mathcal{E}(u)$ is independent of time for every u. Every functional $\int \varphi(u)\, \mathcal{E}(u)\, \mathrm{d}u$ is then conserved. The free energy $W^\pm$ of each branch (S09) and the two further quadratic invariants found by Chandran et al. (2026), $\Gamma^\pm$ and the phase-space helicity $\mathcal{H}^\pm$ (their eqs B25 and B18), are the members $\varphi = 1$, u and $1/u$, up to normalisation; we do not use the last two here.

The nonlinearity of KRMHD therefore moves no free energy between phase velocities. Of the terms that act on G in a simulation, damping at a rate that depends only on $k_\perp$ or k_z, applied alike to every Hermite moment, removes free energy from each slice separately (18), and only velocity-dependent collisions exchange it between slices.

4.2 The two-slice law and the Hermite flux

Let $C(u,u') = \langle a_u\, a_{u'}\rangle$ be the two-slice correlation, whose diagonal is $\mathcal{E}(u)$.

Proposition 1 (Two-slice law)**.** *Under the hypotheses of theorem 1, with no collision operator, source or dissipation acting on G, and for slice amplitudes that are continuously differentiable in $\boldsymbol{r}$ and t with derivatives continuous in u,*

$$\frac{\mathrm{d}C(u,u')}{\mathrm{d}t} = -s\,(u-u')\,\langle a_{u'}\,\nabla_\parallel a_u\rangle. \tag{19}$$

Proof. By (15), $\partial_t a_u = -\{\Phi, a_u\} - su\,\nabla_\parallel a_u$, and the Leibniz rule gives $\partial_t(a_u a_{u'}) = -\{\Phi, a_u a_{u'}\} - su'\,\nabla_\parallel(a_u a_{u'}) - s\,(u-u')\,a_{u'}\nabla_\parallel a_u$. Since $\nabla\cdot\boldsymbol{u}_E = \nabla\cdot\hat{\boldsymbol{b}} = 0$, the first two terms are divergences, whose box averages vanish. □

The $\boldsymbol{E}\times\boldsymbol{B}$ flow carries the two slices of a pair together and cannot change their correlation, which changes only through their relative streaming along the perturbed field. Phase mixing and unmixing are therefore the decorrelation and recorrelation of slices along $\nabla_\parallel$, a definition that refers to no wavenumber.

Free energy moves between Hermite moments only through the streaming term. In the Hermite form (9) the bracket does not change $\langle g_m^2\rangle$, and since $\mathsf{P}A$ is symmetric and $\langle g_n\nabla_\parallel g_{n'}\rangle$ antisymmetric, only the coupling between m and $m+1$ survives in $E_{\le m} = \sum_{n\le m} W(n)$:

$$\frac{\mathrm{d}E_{\le m}}{\mathrm{d}t} = -\Pi_{m+1/2}, \qquad \Pi_{m+1/2} = s\,b_m\langle g_m\nabla_\parallel g_{m+1} - g_{m+1}\nabla_\parallel g_m\rangle, \qquad b_m = p_m A_{m,m+1}, \tag{20}$$

the free energy per unit time and volume that passes from moment m to $m+1$. In slice variables $\Pi_{m+1/2} = -\iint K_m(u,u')\,\partial_t C(u,u')\,\mathrm{d}u\,\mathrm{d}u'$, with the Christoffel–Darboux kernel $K_m(u,u') = \sum_{n\le m}\phi_n(u)\,\phi_n(u')$ of the functions ϕ_n of section 4.3, which by (19) is a weighted integral of the pair correlations $\langle a_{u'}\nabla_\parallel a_u\rangle$. The flux must be computed with the full $\nabla_\parallel$: the Hermite flux of Meyrand et al. (2019) uses it, whereas a flux computed with ∂_z alone, which we call the k_z-only flux, omits the part of the streaming that follows the tilt of the field lines and does not close the budget below.

With forcing of g_0 at the rate ε, dissipation that acts on every g_m in the same way and hypercollisions, the free energy in each moment obeys

$$\frac{\mathrm{d}W(m)}{\mathrm{d}t} = \Pi_{m-1/2} - \Pi_{m+1/2} + \varepsilon\,\delta_{m0} - D_\perp(m) - D_\nu(m), \qquad \Pi_{-1/2} = 0, \tag{21}$$

where $D_\perp(m)$ and $D_\nu(m)$ are the rates at which the spatial dissipation and the hypercollisions remove free energy from moment m. The brackets and the hyper-resistivity of GANDALF act on all moments alike, so (21) holds for its equations up to the error of the dealiased brackets, with the flux (20) evaluated with the same brackets; in the simulations of section 7 it closes to better than $5\times10^{-4}\,\varepsilon$.

4.3 The Hermite spectrum and the two-slice correlation

With the orthonormal basis $\phi_m = \sqrt{p_m \varrho}\, r_m$ of $L^2(\mathbb{R}, \mathrm{d}u)$ (section 3.4), $g_m = p_m^{-1/2} \int \phi_m\, a_u\, \mathrm{d}u$, and the Hermite spectrum is a quadratic form in the two-slice correlation,

$$W(m) = \iint C(u,u')\, \phi_m(u)\, \phi_m(u')\, \mathrm{d}u\, \mathrm{d}u'. \tag{22}$$

If the slices are completely decorrelated, $C(u,u') = e(u)\, \delta(u-u')$, then $W(m) = \int e\, \phi_m^2\, \mathrm{d}u$. At $\Lambda = \infty$, ϕ_m is the Hermite function ψ_m, and for $e = F_M$

$$\int F_M\, \psi_m^2\, \mathrm{d}u = \frac{\Gamma(m+\frac{1}{2})}{\sqrt{2}\,\pi\, m!} = \frac{1}{\pi\sqrt{2m}}\Big(1 - \frac{1}{8m} + O(m^{-2})\Big), \tag{23}$$

where $\Gamma(\cdot)$ is Euler's gamma function. This is the $m^{-1/2}$ spectrum of linear phase mixing. For one Fourier mode whose density moment is driven by white noise that injects free energy at the rate ϵ_k, the exact steady state of Kanekar et al. (2015, eq. 4.10) is the spectrum of decorrelated slices with $e = \pi(1-1/\Lambda)^2 \epsilon_k\, \varrho/(s|k_z|)$; at $\Lambda = \infty$

$$W(m) = \frac{\pi\epsilon_k}{s|k_z|}\, \frac{\Gamma(m+\frac{1}{2})}{\sqrt{2}\,\pi\, m!} \;\simeq\; \frac{\varepsilon}{s|k_z|\sqrt{2m}} \quad (m \gg 1), \tag{24}$$

and at every Λ the large-m form is $\varepsilon/(s|k_z|\sqrt{2m})$, with ε the free-energy injection rate, and the flux (20) is ε through every m. Density forcing loads every slice with the same spatial pattern (section 3.4), two slices loaded together dephase at the rate $s|k_z(u-u')|$, and the steady state superposes all past injections, so the slices are decorrelated in the mean. Free energy injected at $m = 0$ thus moves up the ladder at the speed $\mathrm{d}m/\mathrm{d}t = \Pi/W = s|k_z|\sqrt{2m}$ and reaches $m \simeq (sk_z t)^2/2$ after a time t.

Correlated slices give steeper spectra. If $C(u,u') = [e(u)e(u')]^{1/2}\, c(u-u')$ at $\Lambda = \infty$, with e slowly varying, the large-m form of ψ_m gives

$$W(m) \simeq \frac{\hat{c}(\sqrt{2m})}{\pi\sqrt{2m}} \int e\, \mathrm{d}u, \qquad \hat{c}(q) = \int c(\delta)\cos(q\delta)\, \mathrm{d}\delta, \tag{25}$$

so $\sqrt{2m}$ acts as a wavenumber in velocity space (Schekochihin et al., 2016; Adkins and Schekochihin, 2018). A power law $\hat{c} \propto q^{-p}$ gives $W \propto m^{-(p+1)/2}$, and a Lorentzian correlation $c(\delta) = \delta_c/[\pi(\delta^2+\delta_c^2)]$ of width δ_c gives $\hat{c} = e^{-q\delta_c}$ and $W(m) \propto m^{-1/2}\exp(-\sqrt{2m}\,\delta_c)$. A spectrum that falls exponentially in $\sqrt{m}$ is therefore one in which slices further apart than δ_c in phase velocity are decorrelated and nearer ones are not.

4.4 Truncated hierarchies

Codes such as GANDALF evolve the moments $g_0, \ldots, g_M$ of a branch and close the hierarchy at $m = M$. For the zero closure A is the leading $(M+1)\times(M+1)$ block of the ladder matrix (10); the copy closure $g_{M+1} = g_{M-1}$ changes, for $M \geq 2$, only $A_{M,M-1}$, to $\sqrt{M/2} + \sqrt{(M+1)/2}$. Section 3.2 gave the argument for part (ii) of the following theorem; the theorem adds the normalisation of the slices, the copy closure and the relation to the moments.

Theorem 2. *Let $\Lambda \notin [0,1]$, the value $\Lambda = \infty$ included, and use the zero closure or, for $M \geq 2$, the copy closure.*

(i) *A has $M+1$ real and simple eigenvalues $u_0 < \cdots < u_M$. The right eigenvector of u_j is $\boldsymbol{r}_j = (r_0(u_j), \ldots, r_M(u_j))^{\mathsf{T}}$, with r_m the polynomials of section 3.4; the left eigenvector is $\boldsymbol{w}_j = \mathsf{P}\boldsymbol{r}_j$, and $\boldsymbol{w}_j^{\mathsf{T}}\boldsymbol{r}_l = n_j\,\delta_{jl}$ with $n_j > 0$.*

(ii) *For arbitrary real Φ and Ψ, and brackets evaluated by any bilinear rule, possibly followed by linear operations, such as dealiasing, that act on every g_m alike, g solves (9) if and only if the slice amplitudes $a_j = \boldsymbol{w}_j^{\mathsf{T}} g/\sqrt{n_j} = \sum_m p_m R_{mj} g_m$, with $R_{mj} = r_m(u_j)/\sqrt{n_j}$, satisfy*

$$\partial_t a_j + \{\Phi + s u_j \Psi,\, a_j\} + s u_j\, \partial_z a_j = 0, \qquad j = 0, \ldots, M. \tag{26}$$

(iii) *$g_m = \sum_j R_{mj}\, a_j$, and $\sum_m p_m |g_m(\boldsymbol{k})|^2 = \sum_j |a_j(\boldsymbol{k})|^2$ for every Fourier mode.*

Proposition 2. *For the zero closure, the u_j are the zeros of r_{M+1}, and with $\lambda_j = 1/n_j$*

$$\sum_{j=0}^{M} \lambda_j\, f(u_j) = \int \varrho(u)\, f(u)\, \mathrm{d}u \qquad \text{for every polynomial } f \text{ of degree at most } 2M+1: \tag{27}$$

the u_j and λ_j are the nodes and weights of the Gauss rule of the spectral density $\varrho\,\mathrm{d}u$, and $a_j = \sqrt{\lambda_j}\,\mathcal{T}[G](u_j)$ for states in the span of the retained moments.

Appendix C gives the proofs. A truncated hierarchy is therefore exactly $M+1$ passive slices at the Gauss nodes of the Case–Van Kampen measure, and the results above hold for it with sums over j in place of integrals over u. At $\Lambda = \infty$ the nodes are the Gauss–Hermite nodes (Parker and Dellar, 2015; Tassi, 2022); in KRMHD the Landau coupling enters only A_{10} of a matrix that does not depend on the wavenumber, and the Alfvénic fields do not depend on g, so the decoupling is exact in three dimensions. The largest phase velocity $u_{\max} = u_M$ hardly depends on Λ; with $q_m = (2m+1)^{1/2}$,

$$\begin{aligned} u_{\max} &= q + 2^{-1/3}\alpha_1\, q^{-1/3} - \frac{2^{-2/3}\alpha_1^2}{10}\, q^{-5/3} + O(q^{-3}), && q = q_{M+1} \quad \text{(zero closure)},\\ u_{\max} &= q + 2^{-1/3}\alpha_1'\, q^{-1/3} - 2^{-5/3}\Big(\frac{4}{5\alpha_1'} + \frac{\alpha_1'^2}{5}\Big) q^{-5/3} + O(q^{-3}), && q = q_M \quad \text{(copy closure)}, \end{aligned} \tag{28}$$

where $\alpha_1 = -2.3381074$ and $\alpha_1' = -1.0187930$ are the first zeros of the Airy function and of its derivative (Szegő, 1975; Olver, 1959, theorem 6.32); $u_{\max} = 7.25$ at $M = 32$. Slice j is advected with the stream function $\Phi + s u_j \Psi$, so an explicit time step for the nonlinear term must resolve the speed $\max|\nabla_\perp \Phi| + s u_{\max} \max|\nabla_\perp \Psi|$, not $\max|\nabla_\perp \Phi| + s \max|\nabla_\perp \Psi|$; the simulations of section 7 set their time step in this way.

5 Frozen field lines

When the field lines are frozen into the $\boldsymbol{E} \times \boldsymbol{B}$ flow, which is common to all slices (11), each slice streams, in the frame of the flow, along field lines that do not move. This Lagrangian linearity is due to S09 (§6.3), who showed with the Cauchy solution for the perturbed field that in the frame of $\boldsymbol{u}_E$ the derivative $\hat{\boldsymbol{b}} \cdot \nabla$ is taken along the field lines at $t = 0$ (their eqs 214–216), and concluded that the compressive fluctuations do not cascade along the moving field lines; Kunz et al. (2015, §§2.6 and 4.1) restated it for pressure-anisotropic plasmas, and Lithwick and Goldreich (2001, §2.3.1) made the corresponding remark for a passive scalar in MHD turbulence. We give the solution operator and the hypotheses, and use volume preservation to obtain exact statements about volume-integrated statistics such as $W(m)$.

5.1 The frozen-line theorem

We treat one branch at a time, with state $g(\boldsymbol{r},t)$: the function $G(\boldsymbol{r},\cdot,t)$ of u, or $(g_0,\dots,g_M)$ in a truncated hierarchy, where A replaces L. With collisions $-\mathcal{C}g$ as in (18), (7) reads

$$\partial_t g + \{\Phi, g\} + s\,\nabla_\parallel(Lg) = -\mathcal{C}g. \tag{29}$$

Let $D_t = \partial_t + \{\Phi,\cdot\}$ and $\nabla_\parallel(t) = \partial_z + \{\Psi(\cdot,t),\cdot\}$, let $\boldsymbol{X}_t$ be the flow map of $\boldsymbol{u}_E$, with $\partial_t \boldsymbol{X}_t(\boldsymbol{\xi}) = \boldsymbol{u}_E(\boldsymbol{X}_t(\boldsymbol{\xi}),t)$ and $\boldsymbol{X}_0(\boldsymbol{\xi}) = \boldsymbol{\xi}$, and let $\mathcal{X}_t h = h \circ \boldsymbol{X}_t^{-1}$ turn a function of the labels $\boldsymbol{\xi}$ into a function of position. The first of equations (2) states that the residual $R = \partial_t\Psi + \{\Phi,\Psi\} - \partial_z\Phi$ vanishes. The hypotheses are

(A1) *Ideal induction.* On $[0,t_1]$ the residual R depends on z and t only. The vorticity equation may be forced, viscous or prescribed.

(A2) *Kinetic equation.* The branch has $\Lambda \notin [0,1]$, and g obeys (29), where $\mathcal{C}$ is a linear operator on the velocity variable, the same at every point and independent of the fields, and is bounded or dissipative. There is no spatial dissipation of g and no source.

(A3) *Smoothness.* Φ, Ψ and all their spatial derivatives are continuous on $\mathbb{T}^3\times[0,t_1]$, except that Φ may jump at finitely many times.

(30)

Under (A3), $\boldsymbol{X}_t$ is a volume-preserving diffeomorphism of $\mathbb{T}^3$. The frozen generator $L_0 = s\,\nabla_\parallel(0)\,L + \mathcal{C}$ describes the linear problem with the field lines fixed in their initial shape and no flow, and e^{-tL_0} is its solution operator.

Theorem 3 (Frozen-line theorem)**.** *Assume (30), and let g be a classical solution of (29) on $[0,t_1]$. Then*

$$g(t) = \mathcal{X}_t\, e^{-tL_0}\, g(0), \qquad 0 \le t \le t_1. \tag{31}$$

Proof. With $\hat{\boldsymbol{b}} = \hat{\boldsymbol{z}} + \hat{\boldsymbol{z}}\times\nabla_\perp\Psi$, $\partial_t\hat{\boldsymbol{b}} - \nabla\times(\boldsymbol{u}_E\times\hat{\boldsymbol{b}}) = \hat{\boldsymbol{z}}\times\nabla R$, which vanishes under (A1), so $\hat{\boldsymbol{b}}$ obeys the ideal induction equation and, since $\nabla\cdot\boldsymbol{u}_E = \nabla\cdot\hat{\boldsymbol{b}} = 0$, has Cauchy's solution (Webb et al., 2014, example 3) $\hat{\boldsymbol{b}}(\boldsymbol{X}_t(\boldsymbol{\xi}),t) = (\partial\boldsymbol{X}_t/\partial\boldsymbol{\xi})\,\hat{\boldsymbol{b}}(\boldsymbol{\xi},0)$. The chain rule then gives, for a differentiable function h of the labels,

$$\nabla_\parallel(t)\,\mathcal{X}_t h = \mathcal{X}_t\,\nabla_\parallel(0) h; \tag{32}$$

equivalently, by the Jacobi identity,

$$[D_t, \nabla_\parallel]\, f = \{R, f\}, \tag{33}$$

which vanishes under (A1). Write $g(t) = \mathcal{X}_t\hat{g}(t)$. By (29), $\partial_t\hat{g} = -[s\nabla_\parallel(t)(Lg) + \mathcal{C}g]\circ\boldsymbol{X}_t$. The operators L and $\mathcal{C}$ act at each point with kernels independent of $\boldsymbol{r}$ (for L because the field relations of KRMHD are algebraic), so they commute with composition by $\boldsymbol{X}_t$, and (32) gives $\partial_t\hat{g} = -L_0\hat{g}$ with $\hat{g}(0) = g(0)$. For $\Lambda \notin [0,1]$, $s\nabla_\parallel(0)L$ is antisymmetric in the free-energy norm, which gives uniqueness. □

Without collisions, (31) decouples into slices,

$$a_u(t) = \mathcal{X}_t\,\mathsf{S}_{sut}\,a_u(0), \qquad \mathsf{S}_\ell = e^{-\ell\nabla_\parallel(0)}, \tag{34}$$

where S_ℓ moves a function a distance ℓ along the field lines of $\hat{\boldsymbol{b}}(0)$: each slice is carried rigidly along the initial field lines at the speed su and then relabelled by the flow. This makes exact, for the slices, the remark of S09 (§6.2.4) that the ballistic response is also the nonlinear solution in Lagrangian variables. Whether solutions of the ideal three-dimensional RMHD equations stay smooth is not known in general, so (A3) is an assumption on $[0,t_1]$.

5.2 Single-point statistics and forcing

Let O be a function of the local state, all the velocity-space information at one point, and of finitely many of its derivatives along the field. We call

$$\mathcal{O}_t[f] = \int O(f, \nabla_\parallel(t) f, \nabla_\parallel(t)^2 f, \dots)\, \mathrm{d}^3 r \tag{35}$$

a single-point statistic. Volume preservation and $\nabla_\parallel(t)^k \mathcal{X}_t = \mathcal{X}_t \nabla_\parallel(0)^k$ give $\mathcal{O}_t[\mathcal{X}_t y] = \mathcal{O}_0[y]$.

Corollary 1 (Single-point statistics)**.** *Assume (30), and let $h(t) = e^{-tL_0} g(0)$. Then $\mathcal{O}_t[g(t)] = \mathcal{O}_0[h(t)]$ for every single-point statistic (35) and $0 \le t \le t_1$: every volume-integrated single-point statistic of g takes the value it has in the linear problem along the field lines frozen at $t = 0$, with no flow.*

The corollary covers $W(m,t)$, the slice correlations $C(u,u';t)$, the collisional dissipation rate and the Hermite flux (20) with the full $\nabla_\parallel$. It does not cover spectra in $k_\perp$ or k_z, fluxes resolved in $k_\perp$ shells or computed with ∂_z, which involve Fourier filters that do not commute with $\mathcal{X}_t$. With a source F, Duhamel's formula in the labels shows that the source injected at t' phase-mixes along the field lines frozen at t' and is then relabelled by the flow. For forcing that is white in time this gives the following.

Corollary 2 (White forcing)**.** *Let the forcing be $F\,\mathrm{d}t = \mathrm{d}Y$, with Y a Wiener process with spatially smooth realisations and covariance operator Σ whose increments are independent of the past of g and of the Alfvénic fields, let (A1) and (A3) hold for every realisation, and let g otherwise satisfy (A2). For every single-point statistic whose integrand is a quadratic form,*

$$\mathrm{E}\,\mathcal{O}_t[g(t)] = \mathrm{E}\,\mathcal{O}_0\big[e^{-tL_0} g(0)\big] + \int_0^t \mathrm{E}\,\mathcal{O}_{t'}\big[e^{-(t-t')L_{t'}} \xi\big]\, \mathrm{d}t', \tag{36}$$

where $L_{t'} = s\,\nabla_\parallel(t')\,L + \mathcal{C}$ and ξ is a random field with zero mean and covariance Σ, independent of all other variables. The mean of $W(m,t)$ is therefore the sum of the linear responses to the initial condition and to the individual injections, each phase-mixing along the field lines frozen at its injection time.

Sketch of proof. In the labels, $\mathrm{d}\hat{g} = -L_0 \hat{g}\,\mathrm{d}t + \mathcal{X}_t^{-1}\mathrm{d}Y$, whose solution is an Itô integral with an integrand fixed by the past. The cross term between the initial condition and the stochastic integral has zero mean, the Itô isometry turns the rest into $\int_0^t \mathrm{E}\,\mathcal{O}_0[e^{-(t-t')L_0}\mathcal{X}_{t'}^{-1}\chi]\,\mathrm{d}t'$, and relabelling with (32) gives (36). □

The corollary concerns the mean: a single realisation contains cross terms between injections at different times, and for forcing correlated in time the mean cross terms between injections closer than the correlation time survive.

5.3 Field-aligned wavenumbers

Since $\hat{\boldsymbol{b}}(t)$ is smooth and divergence-free, $K_\parallel(t) = -\mathrm{i}\,\nabla_\parallel(t)$ is self-adjoint on $L^2(\mathbb{T}^3)$ (e.g. Kato, 1966). Its spectral variable is the field-aligned wavenumber, and in a uniform field its spectral projections are the Fourier filters in k_z.

Corollary 3 (Field-aligned wavenumbers)**.** *Assume (30), and let $E_t(I)$ be the spectral projection of $K_\parallel(t)$ on a Borel set I. Then*

$$E_t(I)\, g(t) = \mathcal{X}_t\, e^{-tL_0}\, E_0(I)\, g(0). \tag{37}$$

For real g the parts with positive and negative field-aligned wavenumber evolve separately, so no part of g changes the sign of its field-aligned wavenumber, and without collisions the field-aligned spectrum of every slice is constant in time.

Proof. By (32) and the unitarity of $\mathcal{X}_t$, $E_t(I) = \mathcal{X}_t E_0(I) \mathcal{X}_t^{-1}$, and $E_0(I)$ commutes with e^{-tL_0}. Without collisions e^{-tL_0} multiplies the component of the slice at u at field-aligned wavenumber $k_\parallel$ by $e^{-\mathrm{i}suk_\parallel t}$. □

The Eulerian k_z of the same part of g can change sign. In the frozen-in shear $\Phi = -Szy$, $\Psi = -Sty$, an exact but not periodic solution of (2) with $\nabla_\parallel(t) = \partial_z + St\,\partial_x$, a Fourier mode of the labels becomes a Fourier mode of position with

$$k_z(t) = k_z - Stk_x, \qquad k_\parallel(t) = k_z(t) + Stk_x = k_z. \tag{38}$$

Its k_z changes sign at $t = k_z/(Sk_x)$ while its field-aligned wavenumber stays fixed, so a diagnostic that sorts the free energy by the sign of k_z counts the mode as reversed, although nothing has changed along the field.

5.4 What breaks the theorem

Resistivity or hyper-resistivity on Ψ, parallel dissipation of $z^\pm$, and forcing of Ψ in general make R depend on x and y, so that (33) does not vanish and the field lines are not carried by the flow. Dissipation of g that depends on $k_\perp$ or k_z does not commute with the relabelling and violates (A2). Fourier truncation does not commute with it either; its error vanishes as the resolution increases at fixed t, but not uniformly in t. In the equations that GANDALF integrates, the hyper-resistivity acts alike on $z^\pm$ and every g_m and violates (A1) and (A2), the hypercollisions satisfy (A2), and forcing that adds the same increment to z^+ and z^- keeps (A1).

5.5 A numerical check

The runs of figure 2 integrate one ideal initial state at $N = 16$ to 64 for about three eddy times and compare each run with the frozen-line prediction $e^{-tL_0} g(0)$, computed with GANDALF's right-hand side for g with $\Phi = 0$ and $\Psi = \Psi(0)$ held fixed, and with the k_z-only prediction ($\Phi = \Psi = 0$). For the statistics covered by corollary 1 the relative difference falls by two to three orders of magnitude from $N = 16$ to 64, to between 1.5×10^{-5} for $W(m)$ and 1.4×10^{-4} for the field-aligned flux, while for $W(k_\perp)$, $W(|k_z|)$ and the flux with ∂_z it stays between 0.13 and 0.38 at every N. The k_z-only prediction misses $W(m)$ by 15% and the field-aligned flux by 37%. In ideal turbulence forced through Φ, the response of $W(m)$ to each white injection into g converges with N to phase mixing along the field lines frozen at its injection time, as corollary 2 states. With one term added, the departure from the frozen-line prediction still falls with N for viscosity or forcing of Φ, but not for resistivity or forcing of Ψ or for damping of g in $k_\perp$ or k_z.

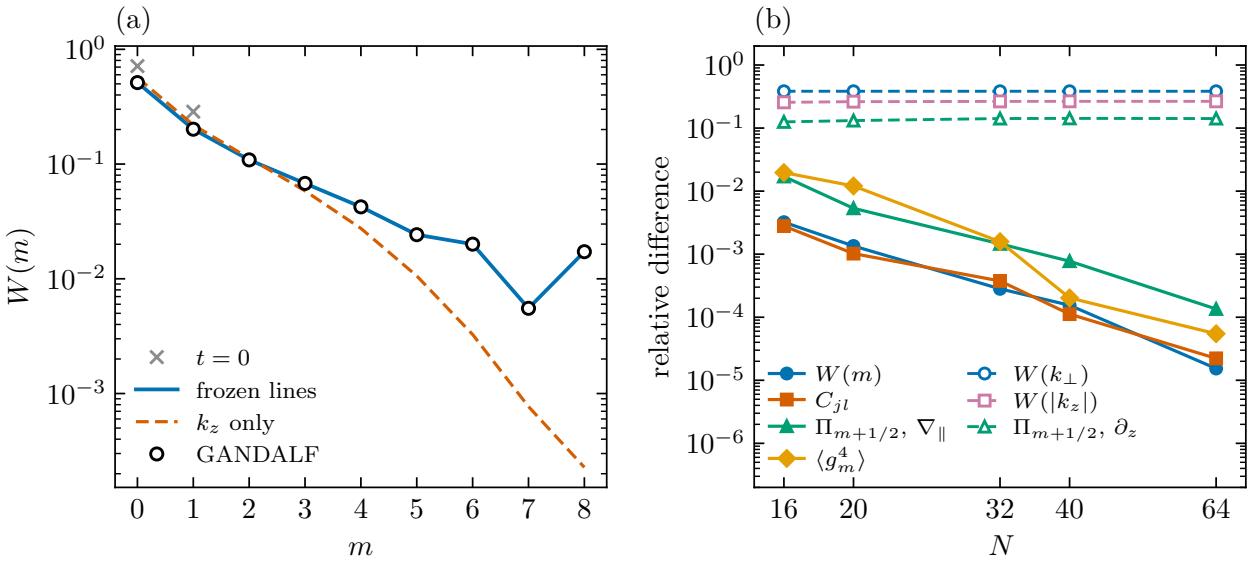


Figure 2: Test of the frozen-line theorem. Ideal GANDALF runs (no dissipation, forcing or collisions) from one initial state at N^3 grid points, with $M = 8$ (zero closure), $\Lambda = \sqrt{5}$, $\beta_i = 1$, random g_0 and g_1, $\chi = 2.21$, the term $\{\Phi, g_0\}$ 2.91 times the streaming term in rms, and $\Delta t = 0.0005$. (a) The Hermite spectrum $W(m)$ at $t = 0.125$ for $N = 64$: initial state (crosses), GANDALF (circles), frozen-line prediction (solid line) and k_z-only prediction (dashed line). (b) The relative difference between GANDALF and the frozen-line prediction at $t = 0.125$ against N, for the single-point statistics $W(m)$, the slice correlations $C_{jl} = \langle a_j a_l \rangle$ at the nodes of section 4.4, $\Pi_{m+1/2}$ with $\nabla_\parallel$ and $\langle g_m^4 \rangle$ (filled markers), and for $W(k_\perp)$, $W(|k_z|)$ and $\Pi_{m+1/2}$ with ∂_z (open markers).

5.6 Echoes

Under (30), $W(m, t)$ and the field-aligned Hermite flux evolve as in the linear problem $\partial_t h = -L_0 h$, so a return of free energy to low m must already be present there, as a linear echo of initial unmixing correlations, the recurrence of a truncated hierarchy, or slow phase mixing at small field-aligned wavenumbers; with white forcing the same holds for $\mathrm{E}\, W(m, t)$. In this sense ideal, smooth dynamics without spatial dissipation of g produces no stochastic echoes in $W(m)$. In drift kinetics with a straight guide field (Schekochihin et al., 2016; Parker et al., 2016) the particles stream along a fixed direction, $[D_t, \partial_z] f = -\{\partial_z \Phi, f\} \neq 0$, and the z-dependent flow couples positive and negative $k_\parallel$, which is the mechanism of the stochastic echo; a flow independent of z gives none (Schekochihin et al., 2016, appendix B; Kanekar, 2015, §5.2). In KRMHD the shear $\partial_z \boldsymbol{u}_E$ that would change $k_\parallel$ also bends the field lines together with g, so under (A1) the commutator (33) vanishes and no part of g moves between positive and negative field-aligned wavenumbers (corollary 3), although its Eulerian k_z can change sign (38).

The theorem does not contradict the simulations in which echoes were reported (Meyrand et al., 2019), which, like ours, dissipate both the magnetic field and the compressive distribution in space and so violate (A1) and (A2). An echo in $W(m)$ requires such terms, or a failure of deterministic flux freezing for rough fields as the resistivity goes to zero, which (A3) excludes; Eyink et al. (2013) reported such a failure at high conductivity, and Eyink (2018, §VIII.C) argued that in a turbulent inertial range the ideal Ohm's law does not imply flux freezing. The question for turbulence is therefore what the terms that break (30) do to $W(m)$, which the model of section 6 addresses.

6 A two-channel model of the Hermite spectrum

In a turbulent steady state the compressive distribution is forced and dissipated, and the hypotheses (30) fail in two ways: the induction equation is not ideal, and g is dissipated in space. The exact results still constrain what the nonlinearity can do. It moves no free energy between slices (section 4.1), and as long as the field lines stay frozen into the flow it moves none between field-aligned wavenumbers (corollary 3). This section builds a model on the assumption that the two failures matter only through the removal of free energy at small perpendicular scales, and lists the predictions by which it can be tested.

6.1 The two channels

Phase mixing at a frozen field-aligned wavenumber. With white forcing, the mean of $W(m)$ is a sum of linear responses to the individual injections, each phase mixing along the field lines frozen at its injection time (corollary 2). Each injection has, along those lines, a distribution of field-aligned wavenumbers κ, which it keeps (corollary 3). By (24), free energy injected at κ moves up the Hermite ladder at the speed $\mathrm{d}m/\mathrm{d}t = s\kappa\sqrt{2m}$ and carries the flux

$$\Pi_{m+1/2} \simeq s\kappa\sqrt{2m}\, W(m) \qquad (m \gg 1), \tag{39}$$

with no part returning to lower m.

A uniform sink. Each slice is a passive scalar, which the turbulence carries to the dissipation scale, where the spatial dissipation removes it, slice by slice (section 4.1). For passive scalars in turbulence the rate at which variance is dissipated stays finite as the diffusivity goes to zero (Shraiman and Siggia, 2000; Falkovich et al., 2001), and the time a scalar takes to reach the dissipation scale is set by the flow, not by the diffusivity. All slices share the $\boldsymbol{E}\times\boldsymbol{B}$ flow, and in its frame they differ only by their streaming along the field lines (theorem 3), which moves them along the lines but does not by itself refine them across the lines. We therefore assume that the turbulence cascades every slice in the same way. If the perpendicular spectrum of g_m then has the same shape for every m, the dissipation, which acts alike on every moment, removes free energy from each moment in proportion to its free energy,

$$D_\perp(m) = \frac{W(m)}{T}, \tag{40}$$

with one time T for all m. Equivalently, for single-point statistics such as $W(m)$, the nonlinearity of KRMHD acts as a damping at a rate $1/T$ that does not depend on m. Section 6.4 replaces it by a rate that depends on the age of the free energy, which the simulations require (section 7.5).

In the steady state of (21), treated as continuous in m and without collisions, the flux (39) changes with m only through the sink (40):

$$\frac{\mathrm{d}\Pi}{\mathrm{d}m} = -\frac{W(m)}{T} = -\frac{\Pi(m)}{s\kappa T\sqrt{2m}}. \tag{41}$$

Since $\mathrm{d}\sqrt{2m}/\mathrm{d}m = 1/\sqrt{2m}$, this integrates to $\ln[\Pi(m)/\Pi(0)] = -\sqrt{2m}/(s\kappa T)$. All the free energy enters the ladder through g_0, so $\Pi(0) = \varepsilon$, the rate of injection into g_0, and $W = \Pi/(s\kappa\sqrt{2m})$ gives

$$W(m) = \frac{\varepsilon}{s\kappa\sqrt{2m}} \exp\Big(-\frac{\sqrt{2m}}{s\kappa T}\Big). \tag{42}$$

The prefactor is the $m^{-1/2}$ spectrum of linear phase mixing, whose flux ε is the same at every m. The exponential can be read off by following the free energy up the ladder. By (39) it moves at $\mathrm{d}m/\mathrm{d}t = s\kappa\sqrt{2m}$, so $\mathrm{d}\sqrt{2m}/\mathrm{d}t = s\kappa$, and it reaches m at the age

$$\vartheta_m = \frac{\sqrt{2m}}{s\kappa}. \tag{43}$$

The sink removes it at the rate $1/T$ on the way, so the fraction that reaches m is $\exp(-\vartheta_m/T)$, which is the exponential in (42). Because the injected pattern is fixed in the Eulerian frame and the field lines are tilted, the injected free energy has a distribution $P(\kappa)$ of field-aligned wavenumbers, and

$$W(m) = \varepsilon \int P(\kappa)\, \frac{1}{s\kappa\sqrt{2m}} \exp\Big(- \frac{\sqrt{2m}}{s\kappa T}\Big)\, \mathrm{d}\kappa, \tag{44}$$

with each κ breaking at its own m, so that a broad $P(\kappa)$ turns the stretched exponential into a more gradual decline. For a single κ the local slope is $\mathrm{d}\ln W/\mathrm{d}\ln m = -\frac{1}{2}(1 + \sqrt{m/m_*})$, which passes through -1 at $m = m_*$ and $-3/2$ at $m = 4m_*$, so the spectrum is close to m^{-1} around the break without being a power law. A spread of κ makes the decline more gradual, because the free energy at large κ phase mixes fastest and survives to the largest m; for a distribution with a Gaussian tail of width σ, (44) falls at very large m as $\exp[-c\,(m/s^2T^2\sigma^2)^{1/3}]$ with c of order one, more slowly than the single-κ exponential. With the distribution of the injected pattern measured at $\beta_i = 1$ and $M = 64$ (section 7; mean $1.12\,k_1$ and rms $1.34\,k_1$, with k_1 the wavenumber of the box) and m_* evaluated at the mean, the local slope of (44) stays within 0.15 of -1 over a factor of five in m around m_* and reaches only -2 at $16\,m_*$ and -2.4 at $64\,m_*$, against -2.5 and -4.5 for a single κ. A spectrum measured over a limited range near m_* can therefore look like a power law somewhat steeper than m^{-1}. Hypercollisions add the sink $D_\nu(m)$ to (41). For the comparison with simulations we solve, for each κ, the forced linear hierarchy (9) with $\nabla_\parallel$ replaced by $\mathrm{i}\kappa$, with the damping rate $1/(2T)$ added to every amplitude and with the hypercollisions of the run, and sum over the measured $P(\kappa)$ (appendix D). This reproduces (44) at large m and is exact for the linear problem at small m. Once T and $P(\kappa)$ are measured, the model has no free parameter.

The form (42) is that of Zocco and Schekochihin (2011, eq. 112) for the Hermite spectrum of growing modes, with their 2γ replaced by $1/T$, and Schekochihin et al. (2016, eqs 2.58–2.59) solved the same continuity equation in m with the collisional sink $2\nu m$ in place of $1/T$. What the model adds is the identification of the sink with the passive-scalar cascade, which the exact results of sections 4 and 5 make the only effect of the nonlinearity on $W(m)$ once the field lines are frozen.

6.2 Scales

The exponential in (42) matters once the age (43) of the free energy reaches the cascade time T. The spectrum therefore follows $m^{-1/2}$ up to the break

$$m_* = \frac{(s\kappa T)^2}{2}, \tag{45}$$

and beyond it decays exponentially in $\sqrt{m}$, not as a power law.

The dependence of m_* on the plasma parameters follows from estimating the two times at the outer scale. The field-aligned wavenumber κ is S09's $k_{\parallel 0}$ (§6.3), the wavenumber along the perturbed field, which without a parallel cascade stays that of the large-scale stirring. It is therefore of the order of the parallel wavenumber $k_\parallel$ of the outer-scale turbulence. Phase mixing then carries free energy to moment m in the time

$$\vartheta_m \sim \frac{\sqrt{m}}{k_\parallel v_{\mathrm{th}}}, \tag{46}$$

and the perpendicular cascade removes it in a time of the order of the outer eddy-turnover time,

$$T \sim \frac{1}{k_\perp u_\perp}, \tag{47}$$

where $u_\perp$ is the $\boldsymbol{E} \times \boldsymbol{B}$ velocity at the outer scale. The free energy reaches m before the cascade removes it, $\vartheta_m \lesssim T$, as long as

$$\sqrt{m} \lesssim \frac{k_\parallel v_{\rm th}}{k_\perp u_\perp} = \frac{v_{\rm th}}{v_A}\,\frac{k_\parallel v_A}{k_\perp u_\perp} = \frac{\beta_i^{1/2}}{\chi}, \tag{48}$$

where $v_{\rm th}/v_A = \beta_i^{1/2}$ and $\chi = k_\perp u_\perp/(k_\parallel v_A)$ is the nonlinearity parameter. Hence

$$m_* \sim \frac{\beta_i}{\chi^2}. \tag{49}$$

This is the phase-mixing threshold $k_\parallel v_{\rm th} \sim \sqrt{m}\, k_\perp u_\perp$ of Schekochihin et al. (2016, eq. 4.5), written in KRMHD variables, with its coefficient fixed by (45) and the cascade time; for $\chi \sim 1$ it is the crossover $m \sim \beta_i$ stated by Kawazura et al. (2019). The estimate takes T to be a fixed multiple of the outer eddy time $1/(k_\perp u_\perp)$; section 7.6 finds that m_* grows somewhat faster than χ^{-2} down to $\chi \approx 0.56$ and about as χ^{-4} below it, because the sink per eddy time is not fixed. If the collisional cutoff m_c lies below m_*, most of the injected free energy reaches it and is dissipated by collisions; if it lies above, most is dissipated at small perpendicular scales, and the collisional share is about $\exp(-\sqrt{2m_c}/(s\kappa T))$.

By (25), the exponential factor in (42) is the cosine transform at $q = \sqrt{2m}$ of a Lorentzian correlation between slices with width

$$\delta u_c = \frac{1}{s\kappa T} = \frac{1}{\sqrt{2m_*}}. \tag{50}$$

Two slices loaded with the same pattern dephase at the rate $s\kappa|u - u'|$ and are removed at the rate $1/T$, so in the steady state slices closer than δu_c in phase velocity remain correlated.

6.3 Predictions

The model predicts the following, each measurable independently of the others.

(i) *Uniform sink.* $D_\perp(m)/W(m)$ is independent of m, and the budget (21) closes with the flux (20) computed with the full $\nabla_\parallel$, but not with the k_z-only flux.

(ii) *Shape.* With T from (i) and $P(\kappa)$ measured along field lines for the injected pattern, $W(m)$ and the field-aligned flux follow the model. A return flux would make both fall below it at large m.

(iii) *Dependence on β_i.* At fixed χ the Alfvénic turbulence does not depend on β_i, so $P(\kappa)$ is unchanged and, if T is too, $m_* \propto \beta_i$: the break moves into a fixed range of m as β_i decreases. With a sink that depends on the age of the free energy (section 6.4), the effective T falls somewhat as β_i decreases, because the same moments then hold older free energy.

(iv) *Anomaly.* T stays finite and does not depend on the dissipation coefficients as the resolution increases. Section 7.7 tests this at four resolutions.

(v) *Eulerian diagnostics.* A split of the flux by the sign of k_z assigns to the backward direction the forward flux of modes whose k_z and field-aligned wavenumber have opposite signs (38), and so reports a backward flux where the field-aligned flux has none. If the free energy at moment m phase mixes at field-aligned wavenumber κ while the split sorts it by $k_z = \kappa - \delta$, with $\delta = \boldsymbol{b}_\perp \cdot \boldsymbol{k}_\perp$ Gaussian with variance $\sigma^2 = \langle k_z^2 \rangle - \kappa^2$, the expected ratio of the unmixing to the phase-mixing part is $[\sigma\mathcal{N}(\zeta) - \kappa(1 - \mathcal{C}(\zeta))]/[\kappa\mathcal{C}(\zeta) + \sigma\mathcal{N}(\zeta)]$ with $\zeta = \kappa/\sigma$ and $\mathcal{N}$, $\mathcal{C}$ the standard normal density and distribution function; it grows with the tilt and so with the range of resolved $k_\perp$.

(vi) *Slice correlations.* The correlation between slices falls off over δu_c (50).

The echo picture (Meyrand et al., 2019) predicts instead a return flux along the field, which reduces the net field-aligned flux and $W(m)$ below the model at large m, and a field-aligned wavenumber of g that grows with $k_\perp$ through the inertial range as g inherits the parallel cascade of the Alfvénic fields.

6.4 Limits of the model

The model rests on two assumptions that the exact results do not guarantee. First, it assumes that the field lines stay frozen long enough for the injected κ to be kept, although the hyper-resistivity of any simulation, and resistivity in a plasma, reconnect them. Meyrand et al. (2019) argued that this relabelling gives g the parallel structure of the Alfvénic fields within about one eddy-turnover time. Second, it assumes that the perpendicular cascade has a dissipative anomaly. Drivas and Eyink (2017) showed that anomalous dissipation of a passive scalar requires Lagrangian spontaneous stochasticity, the non-uniqueness of trajectories as the diffusivity goes to zero, and Eyink (2018, §VIII.C) argued that the same stochasticity breaks flux freezing in a turbulent inertial range. If the two go together, the frozen κ of the first assumption must fail at the scales where the anomaly of the second operates. The model can survive this only if the free energy that phase mixes sits at scales where the field lines are still effectively frozen and the relabelling acts on free energy that the perpendicular sink is already removing. Section 7 measures the field-aligned wavenumber of g as a function of $k_\perp$ to see where this holds. The model also takes the cascade rate to be the same for every slice; slices with large $|u|$ stream fast along wandering field lines, which may refine them across the lines faster than slow slices, and prediction (i) measures how large this effect is. At 128^3 and 256^3 the field-aligned wavenumber of g_0 grows through the inertial range, roughly as $k_\perp^{1/3}$ to $k_\perp^{0.4}$, more slowly than that of the Alfvénic fields, and the sink rate falls by about 30% with each doubling of the resolution up to 256^3 (section 7.7). At the resolutions reached here the first assumption therefore holds only approximately and the second does not hold.

If the sink depends on the time ϑ since injection rather than being uniform, as section 7.5 finds, free energy injected at κ reaches m at the age $\vartheta = \sqrt{2m}/(s\kappa)$, and (42) becomes

$$W(m) = \frac{\varepsilon}{s\kappa\sqrt{2m}} \exp\Big(- \int_0^{\sqrt{2m}/(s\kappa)} \gamma(\vartheta)\, \mathrm{d}\vartheta \Big), \tag{51}$$

with $\gamma(\vartheta)$ the rate at which the perpendicular cascade removes free energy of age ϑ. For a passive scalar injected at the outer scale, γ is small at first, while the cascade carries the fresh pattern towards the dissipation scale, and approaches a constant once the scalar's spectrum has filled in. A uniform $1/T$ then removes too much free energy at small m and too little at large m. The function $\gamma(\vartheta)$ belongs to the Alfvénic turbulence, not to β_i, so at fixed χ one $\gamma(\vartheta)$ should describe every β_i.

7 Simulations

7.1 Set-up

We integrate (2) in Elsasser form with one compressive branch (9), the branch Λ^+ of (6) with $\tau = Z = 1$, truncated at $M = 32$ with the zero closure, using GANDALF (Kanekar, 2025) in double precision: version 0.7.0 on a CPU for the runs at $M = 32$, and version 0.8.1, whose solver is the same, on a GPU for the runs at $M = 64$ (section 7.5). The box is periodic with side $L = 1$ and 64^3 grid points, so $k_1 = 2\pi$; $v_A = 1$ and $\tau_A = L/v_A = 1$. This section reports the run at $\beta_i = 1$, where $\Lambda^+ = \sqrt{5}$; section 7.5 describes the runs at lower β_i.

The Alfvénic fluctuations are forced through Φ alone, so that z^+ and z^- receive the same increment, which keeps the forcing out of the induction equation (section 5). The forcing is white in time and acts on the Fourier modes with $1 \le k_\perp/k_1 \le 2$ and $|k_z| = k_1$, with injection rate $\varepsilon_A = 0.3$. Excluding $k_z = 0$ from the band avoids the two-dimensional condensate that a band including it drove at this resolution. The density moment g_0 is forced on the same band, white in time, with its amplitude rescaled at every step so that the expected injection rate is $\varepsilon = 0.3$; the band excludes $k_\perp = 0$, where injected free energy would phase mix with no perpendicular cascade. Every injection therefore has the single Eulerian parallel wavenumber $k_z = \pm k_1$, and its field-aligned wavenumbers differ from k_1 only through the tilt of the field lines.

GANDALF applies the hyper-resistivity $\eta\,(k_\perp^2/k_{\perp,\max}^2)^2$, with $\eta = 300$, alike to $z^\pm$ and to every g_m; there is no parallel dissipation. Hypercollisions damp g_m at the rate $\nu\,(m/M)^6$ with $\nu = 30$ for $m \ge 2$, and so remove its free energy at the rate $2\nu\,(m/M)^6$. Free energy moving up the ladder at $\mathrm{d}m/\mathrm{d}t = s\kappa\sqrt{2m}$ passes the moments $m \le M/2$ quickly enough that they remove only 3% of the injected free energy there; above $m = M/2$ they remove most of what arrives. The time step is GANDALF's Lawson–Runge–Kutta scheme (Lawson-RK4; Lawson, 1967), which integrates the streaming exactly through an integrating factor and treats the nonlinear terms with explicit RK4. Following section 4.4, $\Delta t = 0.8 \times 2\sqrt{2}/\rho$, where ρ is the spectral radius of the explicit slice-advection operator at the extreme nodes $u = \pm u_{\max}$, computed by a Lanczos iteration every 25 steps, or the Alfvénic limit $0.8\,\Delta x/\max|\nabla_\perp z^\pm|$ if that is smaller. The Alfvénic turbulence is first spun up with the same forcing and dissipation. The run then starts from $g = 0$ and lasts $53\,\tau_A$, and all averages are over the final $50\,\tau_A$, $3 \le t \le 53$; W reaches its mean level within about $3\,\tau_A$.

We characterise the outer scale by the part of Φ in the shells $k_\perp/k_1 \le 2.5$. Its energy-weighted mean $k_\perp$ and $|k_z|$ are $k_{\perp,\mathrm{out}} = 9.7 = 1.55\,k_1$ and $k_{\parallel,\mathrm{out}} = 6.97 = 1.11\,k_1$, and its rms $\boldsymbol{E} \times \boldsymbol{B}$ velocity is $u_{\mathrm{out}} = 0.71$. The nonlinearity parameter is $\chi = k_{\perp,\mathrm{out}} u_{\mathrm{out}}/k_{\parallel,\mathrm{out}} = 0.99$ (0.85 with the rms field-aligned wavenumber of the same field in place of $|k_z|$), and the outer eddy time is $\tau_{\mathrm{nl}} = 1/(k_{\perp,\mathrm{out}} u_{\mathrm{out}}) = 0.144$. The measured free-energy injection rate over the window is $\varepsilon = 0.307$.

7.2 Diagnostics

The energy change of every step is split into the parts due to the ideal dynamics, the forcing, the hyper-resistivity and the hypercollisions. This gives $D_\perp(m)$ and $D_\nu(m)$ in (21) exactly for every moment, and the remainder is the time-stepping error of the ideal part. Every 20 steps we record $W(m)$, its perpendicular spectrum $W(m, k_\perp)$, the field-aligned flux (20) and the k_z-only flux, the split of the k_z-only flux into phase-mixing and unmixing parts by the sign of k_z (Schekochihin et al., 2016), the slice amplitudes a_j at the $M+1$ nodes of theorem 2 with their correlation matrix $C_{jl} = \langle a_j a_l \rangle$, and the Alfvénic energy $E_A = \frac{1}{2}\langle|\nabla_\perp \Phi|^2 + |\nabla_\perp \Psi|^2\rangle$ and its spectrum.

Field-aligned wavenumbers are measured along traced field lines, as is done for the Alfvénic fields (Cho and Vishniac, 2000; Meyrand et al., 2019). On snapshots spaced by $5\,\tau_A$ within the

Table 1: Checks of the $\beta_i = 1$ run over the window $3 \le t \le 53$.

Check	Alfvénic	Compressive
Time-stepping error / injected energy	5.0×10^{-4}	3.8×10^{-8}
|injection − dissipation| / injection	1.7%	0.06%
Standard deviation / mean of E_A and W	0.117	0.141
Spread of block means, four blocks of $12.5\,\tau_A$	2.4% (E_A)	8.8% ($W(m)$), 6.9% (Π), 15.5% ($D_\perp$), 8.4% (D_ν)
Energy fraction beyond $k_\perp/k_1 = 21$	4×10^{-5}	2×10^{-5}
Free energy at $m > 3M/4$	–	1.1%
Share of D_ν at $m \ge M/2$	–	95%

averaging window ($t = 10, 15, \ldots, 50$ in this run; every $10\,\tau_A$ at 96^3 and above) we trace 4096 lines of $\hat{\boldsymbol{b}} = \hat{\boldsymbol{z}} + \hat{\boldsymbol{z}} \times \nabla_\perp \Psi$ over four box lengths in z, sample g_m, Φ and Ψ along them, the last three also filtered into the bands $k_\perp/k_1 \in [0,2]$, $[2,4]$, $[4,8]$ and $[8,16]$, and compute the spectra along the lines. We report κ_1, the first moment of $|\kappa|$. The distribution $P(\kappa)$ of the injected free energy is measured in the same way, on fresh realisations of the density forcing sampled along the same lines. Appendix D gives the details and the computation of the model.

7.3 Validation

Table 1 lists the checks, and figure 3 shows the run. Energy is conserved by the time step and balanced between injection and dissipation. The Alfvénic energy fluctuates by 12% about its mean, which at this box size is intrinsic to the forced turbulence. The block means of the averages used below, each divided by the injection rate of its block, agree to within 9% except $D_\perp$ at $m = 0$ (15.5%), so the averages are converged to about that level. The perpendicular spectrum $E_A(k_\perp)$ follows $k_\perp^{-5/3}$ only over $2 \le k_\perp/k_1 \le 8$ and is steeper than $k_\perp^{-4}$ from $k_\perp/k_1 = 9$, so the inertial range is short. The Hermite spectrum is marginally resolved: 1.1% of the free energy lies at $m > 3M/4$. We therefore treat this run as a pilot, which tests predictions (i), (ii), (v) and (vi) but its break m_* lies beyond the highest moment it keeps, $m = M$ (section 7.4).

7.4 Results at $\beta_i = 1$

(i) A uniform sink. Figure 4(b) shows $D_\perp(m)/W(m)$. It lies between 0.37 and 0.47 for $0 \le m \le 16$; over $1 \le m \le 16$ it varies by a factor of 1.20 between its extremes and by 6% in rms, and rises to 0.54 at $m = 24$, where the hypercollisions dominate. We take $1/T = 0.437$, its median over $2 \le m \le 16$, so $T = 2.29 = 15.9\,\tau_{\rm nl}$. The budget (21) closes to $4 \times 10^{-4}\,\varepsilon$ in every moment with the field-aligned flux, and to 8% with the k_z-only flux. The perpendicular dissipation is therefore close to an m-independent sink, as (40) assumes, with a systematic rise of 20% over $1 \le m \le 16$. Two causes are possible: slices with large $|u|$, which dominate the high moments, may be cascaded faster (section 6.4), or the free energy at higher m, which was injected longer ago, has had more time to reach small perpendicular scales. The runs at lower β_i (section 7.5) decide between them: at a fixed time since injection the streaming speed of the slices that dominate moment m, $s\sqrt{2m}$, scales as s^2 and differs by a factor of ten between $\beta_i = 1$ and 0.1, yet the sinks of all runs fall on one curve against that time (figure 7b), which favours the second.

Field-aligned wavenumbers. The injected pattern, which has $k_z = k_1$, has $\kappa_1 = 1.14\,k_1$ along the field lines; the rms of the tilt $\boldsymbol{b}_\perp \cdot \boldsymbol{k}_\perp$ over the pattern is $0.89\,k_1$, and the rms field-aligned

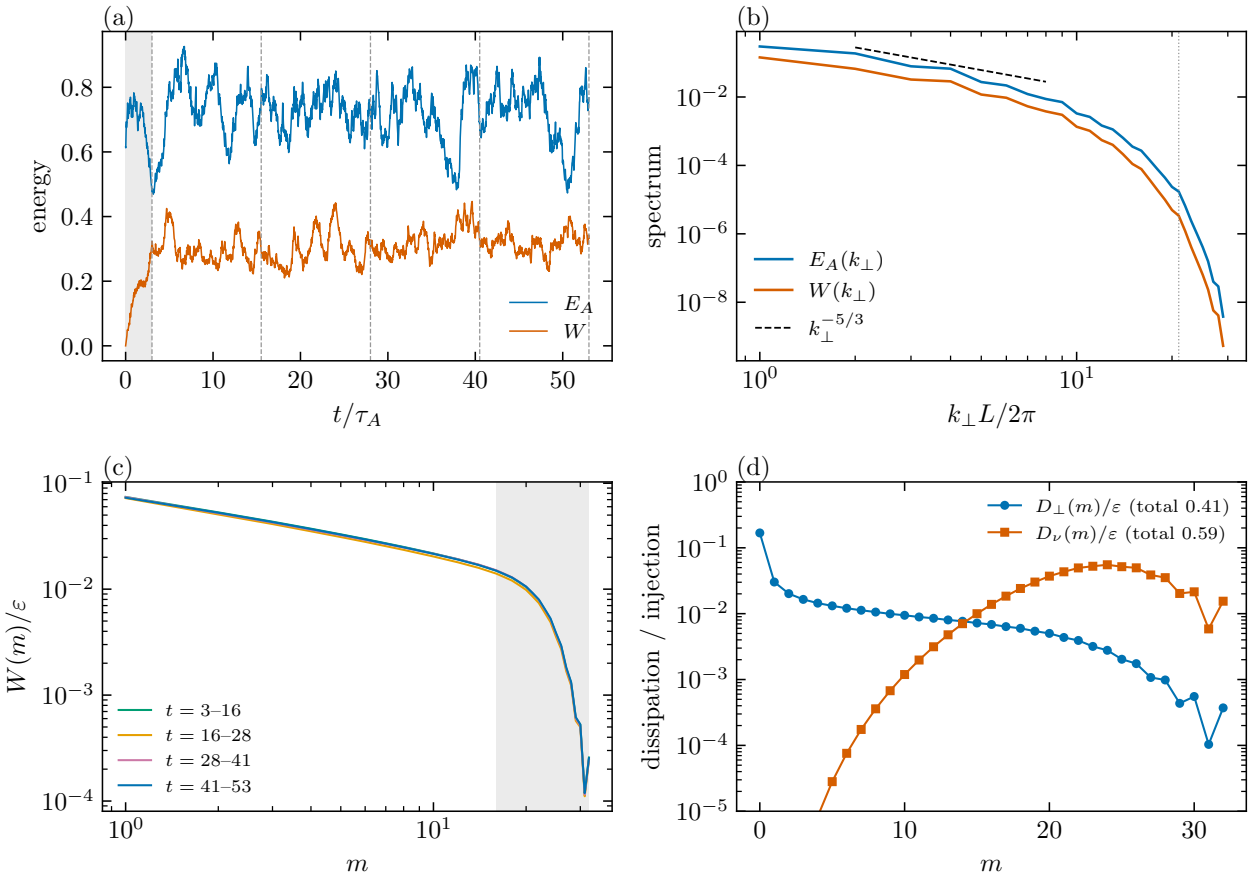


Figure 3: Overview of the run at $\beta_i = 1$, 64^3, $M = 32$. (a) E_A and W against t; the dashed lines bound the four averaging blocks, and the shading marks the excluded start. (b) $E_A(k_\perp)$ and $W(k_\perp)$ averaged over $3 \le t \le 53$, with $k_\perp^{-5/3}$ for reference; the dotted line is the edge of the dealiasing band. (c) $W(m)/\varepsilon$ in each block, divided by the injection rate of that block; the shading marks $m > M/2$, where the hypercollisions act. (d) The rates at which the hyper-resistivity, $D_\perp(m)$, and the hypercollisions, $D_\nu(m)$, remove free energy from moment m, as fractions of the injection rate; in total they remove 41% and 59% of it.

wavenumber is $1.34\,k_1$. The moments g_m have κ_1 between $1.07\,k_1$ and $1.19\,k_1$ for $1 \le m \le 21$, with no trend in m, and the wavenumber implied by the flux, $\Pi_{m+1/2}/[s\sqrt{2(m+1)}\,W(m)]$, lies between $1.02\,k_1$ and $1.10\,k_1$ for $3 \le m \le 23$. Figure 5 resolves κ_1 in $k_\perp$. In the bands centred on $k_\perp/k_1 = 1.5$, 3, 6 and 12, g_0 has $\kappa_1/k_1 = 0.91$, 0.97, 1.36 and 2.04, whereas Φ has 1.11, 1.55, 2.12 and 3.18 and Ψ has 1.13, 1.47, 1.94 and 2.60. The Alfvénic fields gain parallel structure from the outer scale on, though more slowly than $k_\perp^{2/3}$ at this resolution. The density moment keeps the injected field-aligned wavenumber up to $k_\perp = 3k_1$, exceeds it by 20% in the band $4 \le k_\perp/k_1 \le 8$ at the end of the inertial range, and reaches $2.04\,k_1$ only in the band $8 \le k_\perp/k_1 \le 16$, where E_A is steeper than $k_\perp^{-4}$. Measured with k_z, g_0 would appear to have $\kappa_1/k_1 = 0.78$, 0.93, 1.83 and 3.12, close to the Alfvénic values at small scales. In this run, then, g acquires substantial structure along the field only at the end of the inertial range and in the dissipation range, where the hyper-resistivity reconnects the field lines and the perpendicular sink is already removing it. At 128^3, with a longer inertial range, g_0 gains field-aligned structure more gradually and throughout it (section 7.7).

(ii) The shape of $W(m)$ and the flux. With T from (i) and $P(\kappa)$ of the injected pattern, the

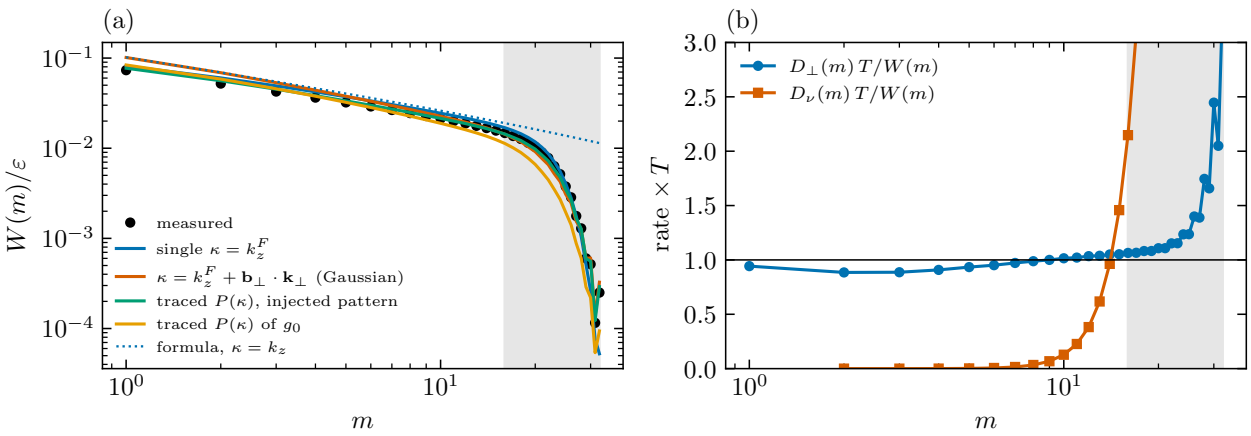


Figure 4: (a) The Hermite spectrum $W(m)/\varepsilon$ at $\beta_i = 1$ (black dots) and the two-channel model with the measured T and four choices of κ: a single $\kappa = k_z^F = k_1$, the wavenumber of the forcing (blue); $\kappa = k_z^F + \boldsymbol{b}_\perp \cdot \boldsymbol{k}_\perp$ with a Gaussian distribution of the tilt (red); the distribution $P(\kappa)$ of the injected pattern traced along field lines (green); and the distribution of g_0 itself traced along field lines (yellow). The dotted line is the formula (42) with $\kappa = k_z$ and no collisions. (b) The rates at which the perpendicular dissipation (circles) and the hypercollisions (squares) remove free energy from moment m, divided by $W(m)$ and multiplied by T. The shading marks $m > M/2$, where the hypercollisions act.

model reproduces $W(m)$ to within 8% for $1 \le m \le 30$, with an rms deviation of 3.3% in $\ln W$ over $2 \le m \le 16$ (figure 4a, green), and the field-aligned flux to within 3–7% for $m \le 24$ (figure 6a). The measured flux lies 4–7% above the model for $1 \le m \le 16$. Only in the last moments, $m \ge 27$, where the hypercollisions dominate, does it fall below the model, by up to 12%, and W by 14% at $m = 32$; a sink that grows with the age of the free energy (section 7.5) acts in this direction. The flux that $W(m)$ would carry if all of it phase mixed at its measured $\kappa_1(m)$ is close to the measured flux, so there is no room for a return flux of more than a few per cent. With a single $\kappa = k_1$ the model gives W 14–16% too high, because it omits the tilt. With the field-aligned distribution of g_0 itself in place of that of the injected pattern, it overdamps, giving W too low by a factor of up to 1.9 at $m = 24$: the free energy phase mixes at the wavenumber at which it was injected, not at the larger one that g_0 shows at small perpendicular scales.

The break. With these values the break lies far beyond $m = M = 32$: $m_* = (s\kappa T)^2/2$ is 103 for $\kappa = k_1$ and 185 for the rms field-aligned wavenumber, and $c_T^2\beta_i/(2\chi^2) = 127$ with $c_T = T/\tau_{\rm nl}$. Free energy injected at $m = 0$ reaches $m = 16$ after about $0.8\,\tau_A$ and $m = 32$ after $1.1\,\tau_A$, against $T = 2.3\,\tau_A$, and the hypercollisions remove 57–61% of the injected free energy (the range over the four blocks). At $\beta_i = 1$ the exponential factor in (42) therefore reduces W by only 29% at $m = 16$ and 39% at $m = 32$, and the test of the shape is correspondingly weak: it confirms that there is no deficit, not that the sink produces a break.

(v) Eulerian diagnostics. The k_z-only flux split by the sign of k_z has an unmixing part of -0.11 to -0.15 times its phase-mixing part for $1 \le m \le 16$, and -0.52 at $m = 0$ (figure 6a), whereas the field-aligned flux is reproduced by a model that has no return flux. A diagnostic of this kind would report a backward flux of about 12% of the forward one where, by the field-aligned measures, there is none. With the rms field-aligned and Eulerian wavenumbers of g_m measured in the run, about $1.2\,k_1$ and 1.8–$1.9\,k_1$, the estimate of prediction (v) gives a ratio of 0.09–0.11 for $m \ge 1$, against the measured 0.11–0.15: the backward flux of the split is accounted for by the tilt of the field lines,

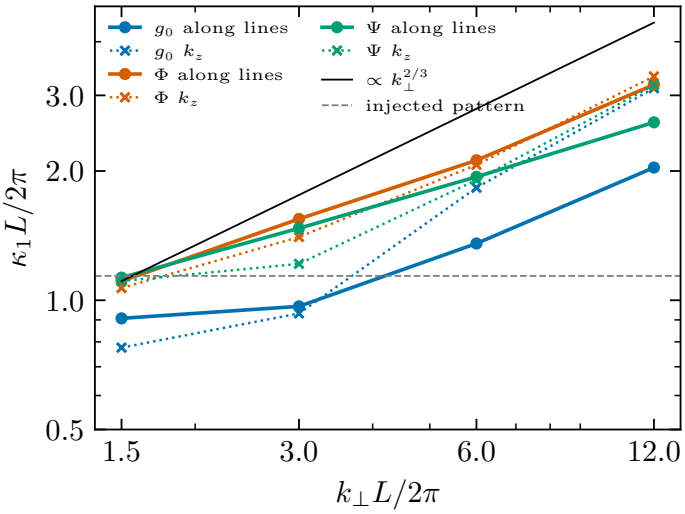


Figure 5: The mean field-aligned wavenumber κ_1 of g_0 (blue), Φ (red) and Ψ (green) in four bands of $k_\perp$, measured along traced field lines (solid) and from the Eulerian k_z (dotted), at $\beta_i = 1$. The dashed line is the injected pattern traced along the same lines; the thin line is proportional to $k_\perp^{2/3}$.

with no field-aligned return flux (figure 9).

(vi) Slice correlations. The width (50) is $\delta u_c = 0.05$–0.07, below the spacing of the nodes, about 0.4 near $u = 0$ at $M = 32$, so it cannot be resolved here. The model, evaluated at the nodes, reproduces the measured normalised correlations $C_{jl}/(C_{jj}C_{ll})^{1/2}$ to an rms difference of 0.027 (figure 6b).

7.5 Dependence on β_i at fixed χ

To bring the break into the resolved range of m, we ran the same set-up at $\beta_i = 0.3$ and 0.1, each at $M = 32$ and $M = 64$, and at $\beta_i = 1$ also at $M = 64$. The Alfvénic sector does not depend on β_i, so the same Alfvénic forcing gives the same χ (the outer-scale value is 0.93–0.99 in every run); the branch is $\Lambda^+ = 6.2206$ and 19.198. The attenuation of free energy that streams up the ladder by hypercollisions $\nu(m/M)^6$ depends on $\nu\sqrt{M}/s$, so we set $\nu = 30\,\beta_i^{1/2}$ at $M = 32$ and $\nu = 30\,\beta_i^{1/2}(32/M)^{1/2}$ at $M = 64$, which keeps the total attenuation fixed while the collisional range moves from $m > 16$ to $m > 32$. With T and $P(\kappa)$ from $\beta_i = 1$, (45) gives $m_* \approx 31$–56 at $\beta_i = 0.3$ and $m_* \approx 10$–19 at $\beta_i = 0.1$. The runs at $\beta_i = 0.1$ and at $M = 64$ use a safety factor of 0.4 on the Alfvénic limit of the time step, with both limits updated every step, and last $66\,\tau_A$, with averages over $16 \le t \le 66$.

Table 2 lists the runs. All conserve energy to within 2.1×10^{-3} of the injection and balance injection and dissipation to within 2%. The free energy fluctuates more slowly as β_i decreases, so its averages are less certain at low β_i; we quote the spread of the four block means as their uncertainty, about $\pm 12\%$ on $W(m)$ at $\beta_i = 0.1$. In the run at $\beta_i = 0.3$ and $M = 32$ the window starts at $t = 15.5$, after the free energy had settled. The spectra show no sign of numerical instability at small scales in any run.

The break enters the resolved range. At $\beta_i = 0.1$ the predicted break, $m_* \approx 6$–10 with the measured T, lies well inside the resolved range, and $W(m)$ falls below $m^{-1/2}$ from small m (figure 7a). The continuum formula (42) with $\kappa = k_1$ and the measured T follows the measured spectrum to about 10% up to $m \approx 16$, where the hypercollisions take over at $M = 32$; at $M = 64$ the spectrum

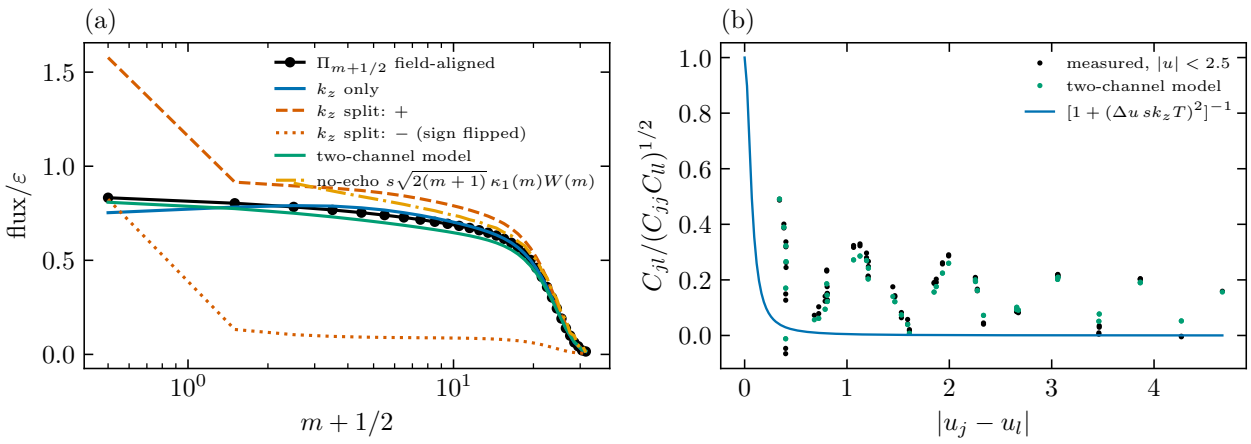


Figure 6: (a) Hermite fluxes divided by ε at $\beta_i = 1$: the field-aligned flux (20) (black), the k_z-only flux (blue), its phase-mixing part (dashed) and its unmixing part with the sign reversed (dotted), the two-channel model (green), and the flux $s\sqrt{2(m+1)}\,\kappa_1(m)\,W(m)$ that $W(m)$ would carry if all of it phase mixed at the measured field-aligned wavenumber (dash-dotted). (b) Correlations $C_{jl}/(C_{jj}C_{ll})^{1/2}$ between slices with $|u| < 2.5$ against $|u_j - u_l|$: measured (black) and two-channel model (green); the line is the Lorentzian $[1 + (\Delta u\, sk_zT)^2]^{-1}$ of width (50).

continues to fall steeply between $m \approx 15$ and 40. At $\beta_i = 0.3$ the spectrum bends at correspondingly larger m. To measure the break we fit $W(m)\sqrt{2m} \propto \exp(-\sqrt{2m}/\ell_*)$ over $2 \le m \le M/2$, below the collisional range. The fitted ℓ_* is 16.8, 8.3 and 3.8 at $\beta_i = 1$, 0.3 and 0.1 ($M = 32$), and 15.7, 7.2 and 3.6 at $M = 64$, against $s\kappa_1 T = 16.4$, 8.0, 3.8, 16.8, 7.6 and 3.7 from the measured T and the injected κ_1: the measured decay rate in $\sqrt{2m}$ is the predicted one to within 7% in every run. At $\beta_i = 1$ the fit range spans less than half a decay length ($\sqrt{2m} \le \sqrt{M} = 5.7$ and 8 against $\ell_* \approx 16$), so the break $m_* = \ell_*^2/2 \approx 140$ there extrapolates the fit well beyond $m = M$; at $\beta_i = 0.1$ and $M = 64$ the range spans 1.7 decay lengths and $m_* \approx 7$. Section 7.6 measures the break with M raised until it is resolved.

Comparison with the model at lower β_i. At $\beta_i = 1$ the measured $W(m)$ is within 7% of the model with the traced $P(\kappa)$ for $1 \le m \le M/2$, at both M. At lower β_i it lies above the model at every m, by up to 21% and 39% at $\beta_i = 0.3$ and 0.1 at $M = 32$ and by up to 20% and 42% at $M = 64$, and the field-aligned flux lies above the model as well. Neither falls increasingly below the model at large m in any run, so there is no sign of a return flux. The excess comes from the assumption of a uniform sink: $D_\perp(m)/W(m)$ rises with m, from 0.33–0.43 at $m \le 1$ in every run to 0.46, 0.53 and 0.62 at $m = 16$ for $\beta_i = 1$, 0.3 and 0.1, so a single rate $1/T$, the median over $2 \le m \le 16$, removes too much free energy at small m. The rise is steeper at lower β_i, where free energy reaches a given m later. Plotted against the time $\vartheta = \sqrt{2m}/(s\kappa_1)$ since injection, the sinks of all runs fall on one curve (figure 7b), with an rms scatter of 4.5% about a joint fit $D_\perp/W = 0.50 + 0.12 \ln \vartheta$, against 12.6% when plotted against m. The sink is about 0.4 for the youngest free energy and levels off at about 0.63 for ages beyond about $2.5\,\tau_A$. The perpendicular sink is therefore a property of the age of the free energy, the same at every β_i at fixed χ, as expected for a passive scalar whose variance reaches the dissipation scale only after the cascade has had time to act (section 6.4). The median rate $1/T$ rises from 0.44 to 0.59 between $\beta_i = 1$ and 0.1 for the same reason: at low β_i the moments $2 \le m \le 16$ hold older free energy. To test this directly we let each injection evolve by the linear hierarchy at its field-aligned wavenumber, with the run's hypercollisions, and remove it at a rate $\gamma(\vartheta)$ that depends only on its age ϑ, the time since injection, the same at every m; with

Table 2: The forced runs, on N^3 grid points with M Hermite moments. The window is the averaging interval in units of τ_A; std/mean is that of E_A and W over the window; the block spread is the largest relative spread of the four block means of $W(m)$, Π, $D_\perp(m)$ and $D_\nu(m)$; $W_>$ is the fraction of the free energy at $m > 3M/4$; $1/T$ is the median of $D_\perp(m)/W(m)$ over $2 \le m \le 16$, the same range in every run, below the collisional range $m > M/2$.

β_i	N	M	window	std/mean (E_A, W)	block spread	$W_>$	$1/T$	$T/\tau_{\rm nl}$
1	64	32	3–53	0.117, 0.141	0.155	1.1%	0.437	15.9
0.3	64	32	15.5–53	0.076, 0.147	0.174	1.0%	0.478	13.5
0.1	64	32	16–66	0.120, 0.180	0.229	0.65%	0.590	11.2
1	64	64	16–66	0.146, 0.143	0.312[a]	1.2%	0.421	15.4
0.3	64	64	16–66	0.158, 0.136	0.185	0.87%	0.508	13.2
0.1	64	64	16–66	0.114, 0.186	0.212	0.44%	0.604	11.0
0.3	96	64	16–66	0.133, 0.153	0.238[b]	1.02%	0.403	16.4
0.3	128	64	16–66	0.114, 0.156	0.236[a]	1.06%	0.352	18.4
0.3	256	64	10–60	0.112, 0.123	0.162	1.19%	0.257	24.6
1	128	64	16–66	0.112, 0.165	0.143	1.35%	0.284	22.6
1[c]	64	64	16–66	0.156, 0.080	0.321[a]	2.15%	0.427	15.3

[a] $D_\perp(m)$ at the forced moment: at $m = 0$ (31% and 24%; 20% at $m = 1$ for $\beta_i = 1$), and at $m = 1$ in the run with g_1 forced (32%, and 29% for $W(1)$); at most 19% for every other average and m. [b] $W(m)$ at $m = 47$ and $D_\nu(m)$ at $m = 51$, in the collisional range. [c] g_1 forced in place of g_0, on $|k_z| = k_1$ and $2k_1$ (section 7.8).

$S(\vartheta) = \exp(-\int_0^\vartheta \gamma)$, the steady state is $W(m) = \varepsilon \sum_\kappa P(\kappa) \int_0^\infty E_m(\vartheta;\kappa)\, S(\vartheta)\, \mathrm{d}\vartheta$, with $E_m(\vartheta;\kappa)$ the free energy at moment m of a unit injection at age ϑ, and $D_\perp(m)$ is the same integral with γS in place of S (for $\gamma = 1/T$ this is the uniform-sink model). Fitting $\gamma(\vartheta) = \gamma_0 + (\gamma_1 - \gamma_0)(1 - e^{-\vartheta/\vartheta_c})$ to $D_\perp(m)/W(m)$ at $\beta_i = 0.1$ ($M = 64$) gives $\gamma_0 = 0.03$, $\gamma_1 = 0.67$ and $\vartheta_c = 0.76\,\tau_A$, about five outer eddy times: fresh free energy is hardly removed, and the sink switches on as the cascade carries it to the dissipation scale. With this $\gamma(\vartheta)$ and no parameter fitted to them, the model gives $D_\perp(m)/W(m)$ at $\beta_i = 0.3$ and 1 to within 5% for $m \ge 2$ and $m \ge 4$, across a factor of ten in β_i (figure 8). It places $W(m)$ 5–18% above the measured values at small m, by about the same amount in all three runs, the fitted one included, so this offset does not come from the dependence on β_i but from the injected $P(\kappa)$ or the treatment of the lowest moments.

Diagnostics. The k_z-only flux split by the sign of k_z has an unmixing part of -0.10 to -0.15 times its phase-mixing part for $1 \le m \le M/2$ at every β_i, which the tilt of the field lines accounts for to within 0.03 for $m \ge 4$ (0.05 at $m = 1$) with no return flux (figure 9), while the per-m budget closes to better than $5 \times 10^{-4}\,\varepsilon$ with the field-aligned flux in every run, and to 2–8% with the k_z-only flux. The slice correlation width is not resolved: the node spacing near $u = 0$ is about 0.4 at $M = 32$ and 0.28 at $M = 64$, against the predicted widths 0.11–0.15 at $\beta_i = 0.3$ and 0.22–0.30 at $\beta_i = 0.1$, and the widths fitted at the nodes, 0.25 and 0.35 at $M = 64$, lie above them.

7.6 The break against β_i and χ

Of the runs in table 2, only those at $\beta_i = 0.1$ place the break well inside the collision-free range $m \le M/2$. To measure m_* against β_i and χ we made the runs of table 3, all on 64^3 points with the forcing of section 7.1, with M large enough in each that the break lies below $M/2$, and with $\nu = 30\,\beta_i^{1/2}(32/M)^{1/2}$ as in section 7.5. At $\chi \approx 0.95$ we ran $\beta_i = 0.05$ and 0.3, and we use the run at $\beta_i = 0.1$ and $M = 64$ of table 2. At $\beta_i = 0.05$ we lowered χ to 0.75, 0.56 and 0.40 by lowering

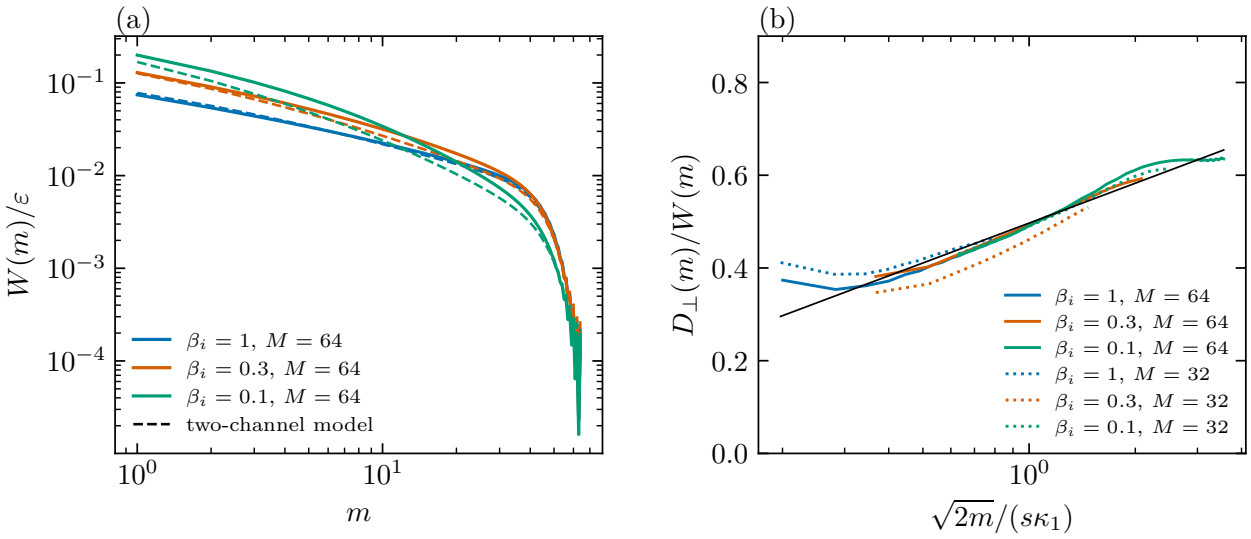


Figure 7: (a) $W(m)/\varepsilon$ at $\beta_i = 1$, 0.3 and 0.1 ($M = 64$) (solid) and the two-channel model with the traced $P(\kappa)$ of the injected pattern and the measured T of each run (dashed). (b) The perpendicular sink $D_\perp(m)/W(m)$ against the time $\sqrt{2m}/(s\kappa_1)$ since injection of the free energy at m, with κ_1 the mean field-aligned wavenumber of the injected pattern, for $1 \le m \le M/2$; the runs at $M = 64$ (solid) reach older free energy than those at $M = 32$ (dotted); the thin line is the joint fit $0.50 + 0.12 \ln[\sqrt{2m}/(s\kappa_1)]$.

the Alfvénic injection rate ε_A from 0.3 to 0.1, 0.031 and 0.0035. As χ decreases the hyper-resistive dissipation of E_A moves to larger scales unless η decreases with it, so we lowered η from 300 to 219 and 156, in proportion to χ, and to 33 at $\chi = 0.40$, where η and ε_A were set by short runs at $M = 2$. The median shell of the hyper-resistive dissipation of E_A is then $k_\perp/k_1 = 9$–11 in every run. At $\beta_i = 0.05$ the free energy fluctuates slowly, so these runs are averaged over 160–300 τ_A.

We measure the break from a fit of $W(m)\sqrt{2m} \propto \exp(-\sqrt{2m}/\ell_*)$, as in section 7.5, but over a range set by each run's own decay: from $m = 2$ to $\sqrt{2m} = 2 + 1.5\,\ell_*$, found by iterating from the whole range $2 \le m \le M/2$. Then $m_* = \ell_*^2/2$. The collision-free range holds 1.7 to 3.5 decay lengths in these runs (table 3), and at the top of the fit range the hypercollisional rate $2\nu_m$ is at most 13% of the perpendicular rate $1/T$. The model value is $(s\kappa_1 T)^2/2$, with T the median rate over $2 \le m \le 16$ of each run and κ_1 from lines traced on its snapshots.

Dependence on β_i. At $\chi \approx 0.95$ the break grows from $m_* = 3.7$ at $\beta_i = 0.05$ to 6.6 and 24.2 at $\beta_i = 0.1$ and 0.3 (figure 10a). A weighted fit gives $m_* \propto \beta_i^{1.09\pm0.04}$, close to the $m_* \propto \beta_i$ of (49). The model values follow $\beta_i^{1.31\pm0.04}$. They agree with the measured m_* to within 12% at $\beta_i \le 0.1$ and lie 30% above it at $\beta_i = 0.3$, where $T/\tau_{\rm nl}$ is larger, 13.6 against 10.9 and 11.0, because the moments $2 \le m \le 16$ that set T hold younger free energy at larger β_i (section 7.5).

Dependence on χ. At $\beta_i = 0.05$ the break grows from $m_* = 3.7$ at $\chi = 0.96$ to 7.0, 14.9 and 60.5 at $\chi = 0.75$, 0.56 and 0.40 (figure 10b). Between neighbouring runs the local exponent $-\mathrm{d}\ln m_*/\mathrm{d}\ln\chi$ is 2.5, 2.7 and 4.1. Down to $\chi \approx 0.56$ the break therefore grows somewhat faster than the χ^{-2} of (49), and below it the growth steepens towards χ^{-4}; no single power law describes the four runs. The model value follows the measured m_* to within 12% for $\chi \ge 0.56$, but at $\chi = 0.40$ it is 167, 2.8 times the measured value.

Both departures come from the perpendicular sink. With the age $\vartheta = \sqrt{2m}/(s\kappa_1)$ and the rate $D_\perp/W$ both measured in units of the outer eddy time $\tau_{\rm nl}$, the sinks of the runs at $\chi \ge 0.56$ agree

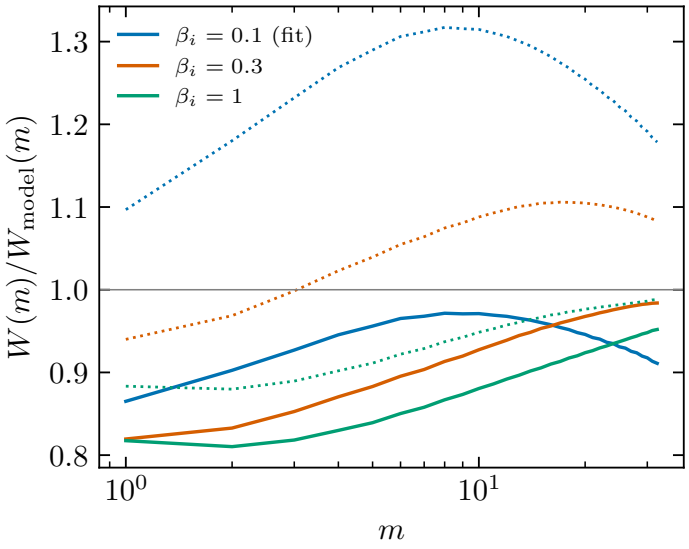


Figure 8: $W(m)$ divided by the two-channel model with the age-dependent sink $\gamma(\vartheta)$ fitted at $\beta_i = 0.1$ (solid) and with the uniform sink $1/T$ of each run (dotted), for the runs at $M = 64$.

to within about 10% for $\vartheta \geq 12\,\tau_{\mathrm{nl}}$: the rate switches on over the first 10–15 τ_{nl} and levels off at 0.087–0.094/τ_{nl}, slightly lower at smaller χ. The range $2 \leq m \leq 16$ that defines T covers the same ages in units of τ_A in every run at $\beta_i = 0.05$, but as χ decreases and τ_{nl} grows it holds free energy that is younger in units of τ_{nl}, from 8.5–24 τ_{nl} at $\chi = 0.96$ to 5–14 τ_{nl} at $\chi = 0.56$. The sink there is weaker, so T/τ_{nl} rises from 10.9 to 13.9 and m_* grows faster than χ^{-2}. At $\chi = 0.40$ the moments $2 \leq m \leq 16$ hold free energy of age 3–8 τ_{nl}, and the sink itself changes: it switches on as at larger χ, reaches 0.069/τ_{nl} between $\vartheta = 12$ and 16 τ_{nl} and falls to 0.055/τ_{nl} at 30 τ_{nl}, where the runs at larger χ are at 0.087–0.094/τ_{nl}. The cascade then removes the free energy more slowly per outer eddy time, as expected once the outer-scale turbulence is weak and its cascade time exceeds τ_{nl}. Over $1 \leq m \leq M/2$ the rate $D_\perp/W$ varies by a factor of 10 at $\chi = 0.40$, against 1.4, 1.7 and 2.5 at $\chi = 0.96$, 0.75 and 0.56, so the median over $2 \leq m \leq 16$ no longer represents the rate over the range $m \leq 171$ on which the break is measured, and $(s\kappa_1 T)^2/2$ overestimates m_*. With the rate averaged over each run's fit range instead, uniformly in $\sqrt{2m}$ as the fit weights it, $s\kappa_1 T$ agrees with the fitted ℓ_* to within 9% at $\chi = 0.40$ and to within 5% in every other run of table 3. The decay beyond the break is therefore still set by phase mixing at κ_1 against the perpendicular sink, as in (51). What departs from (49) at small χ is the assumption that the sink rate is a fixed multiple of $1/\tau_{\mathrm{nl}}$.

7.7 Dependence on resolution

Doubling M from 32 to 64 leaves $W(m)$ unchanged to within 5% for $1 \leq m \leq 10$ and changes $1/T$ by at most 6% at every β_i. What changes is the share of the free energy that reaches the hypercollisions. When the collisional range moves from $m > 16$ to $m > 32$ it falls from 59% to 50% at $\beta_i = 1$, from 43% to 27% at $\beta_i = 0.3$ and from 20% to 10% at $\beta_i = 0.1$, and the model with the traced $P(\kappa)$ gives the shares to within 0.06 at $M = 32$ and 0.02 at $M = 64$. The collisional share of the dissipation is therefore not a converged quantity at $M = 32$: free energy that the hypercollisions remove at a low cutoff is removed at small perpendicular scales when the cutoff is raised.

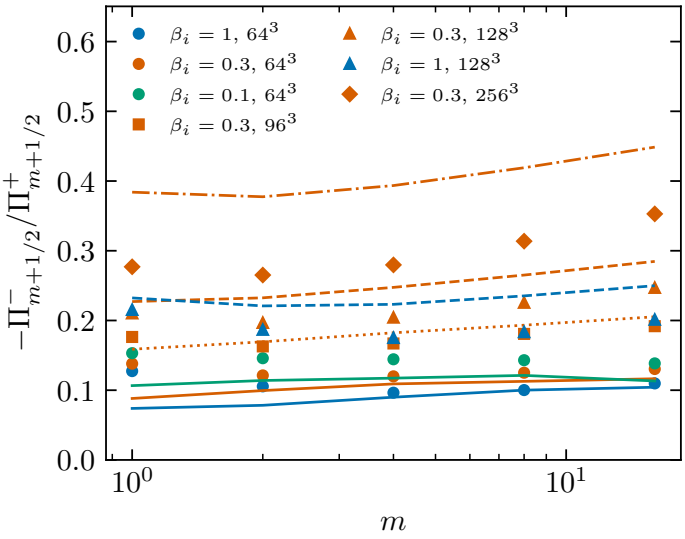


Figure 9: The unmixing part of the k_z-only Hermite flux split by the sign of k_z, relative to its phase-mixing part, measured (markers) and estimated from the tilt of the field lines with no return flux (lines; prediction (v)), for the runs at $M = 64$ with g_0 forced: $\beta_i = 1$ (blue), 0.3 (orange) and 0.1 (green) at 64^3 (circles, solid), 96^3 (squares, dotted), 128^3 (triangles, dashed) and 256^3 (diamonds, dash-dotted).

We repeated the runs at $M = 64$ on 96^3, 128^3 and 256^3 grid points at $\beta_i = 0.3$ and on 128^3 at $\beta_i = 1$, with the hyper-resistivity normalised to the grid, so that the inertial range lengthens with N: the slope of E_A over the fitted shells goes from -4.1 at 64^3 to -2.9 at 256^3. The outer scale is unchanged, with $\tau_{\rm nl} = 0.150$–0.158 and $\chi = 0.92$–0.97 in all six runs. The perpendicular sink is not (figure 11). $D_\perp(m)/W(m)$ falls at every m, by 18–35% with each doubling of N, and $T/\tau_{\rm nl}$ rises from 13.2 at 64^3 to 16.4, 18.4 and 24.6 at $\beta_i = 0.3$, and from 15.4 to 22.6 at $\beta_i = 1$. At $\beta_i = 0.3$, T grows by 5–6 outer eddy times with each doubling, with no sign of levelling off up to 256^3. The break moves out with it, from $m_* \approx 26$ at 64^3 to 86 at 256^3, the share of the free energy removed by hypercollisions rises from 27% to 46%, and $W(m)$ at $m \le 16$ rises by up to 43%.

The model follows these changes. At 128^3 and 256^3, with the T and $P(\kappa)$ measured there, it describes $W(m)$ to within -9% and $+13\%$ and the collisional shares to within 0.02, and the field-aligned flux lies 3–15% above it, with no deficit at large m. The measured decay rate in $\sqrt{2m}$ is within 5% of $1/(s\kappa_1 T)$, except at $\beta_i = 0.3$ at 128^3 and 256^3, where it is 14% and 17% faster.

The field-aligned structure of g changes more slowly. With the longer inertial range the field-aligned wavenumber of g_0 grows gradually across it, roughly as $k_\perp^{1/3}$ at 128^3 and $k_\perp^{0.4}$ at 256^3, against about $k_\perp^{1/2}$ for Φ, rather than only at the dissipation scale. For the moments as a whole the rms field-aligned wavenumber rises only from about $1.25\,k_1$ to $1.35\,k_1$, while the rms Eulerian wavenumber grows from about $1.8\,k_1$ at 64^3 to 3.6–$4.8\,k_1$ at 256^3. The unmixing part of the k_z-split flux grows accordingly, at $\beta_i = 0.3$ from about -0.13 times the phase-mixing part at 64^3 to -0.27 to -0.35 at 256^3. The tilt estimate of prediction (v) follows this growth but overshoots by up to 0.11 at 256^3, where the tilt is largest (figure 9). The field-aligned flux shows no deficit at any resolution.

At the resolutions we reach, then, neither the rate of the perpendicular sink nor the field-aligned wavenumber of g is independent of the dissipation scale, and prediction (iv) is not confirmed. Either the dissipation range of our runs reaches almost to the forcing scales and makes the sink artificially

Table 3: The runs of the m_* scan, on 64^3 points. η is the hyper-resistivity; the window, std/mean and block spread are as in table 2; $n_e = (\sqrt{M} - 2)/\ell_*$ is the number of decay lengths that fit in the collision-free range; $m_* = \ell_*^2/2$ from the fit over $2 \le \sqrt{2m} \le 2 + 1.5\,\ell_*$, with the standard error of the fits to the four block averages; T is the median of $D_\perp(m)/W(m)$ over $2 \le m \le 16$ and κ_1 the mean field-aligned wavenumber of the injected pattern.

β_i	χ	M	η	window	std/mean (E_A, W)	block spread	n_e	m_*	$(s\kappa_1 T)^2/2$	$T/\tau_{\rm nl}$
0.05	0.96	128	300	20–180	0.117, 0.172	0.47[a]	3.45	3.7 ± 0.3	3.3	10.9
0.1	0.96	64	300	16–66	0.114, 0.186	0.21	1.69	6.6 ± 0.4	6.9	11.0
0.3	0.94	192	300	28.5–66	0.124, 0.095	0.26[b]	1.73	24.2 ± 0.6	31.4	13.6
0.05	0.75	128	219	20–200	0.135, 0.139	0.18	2.65	7.0 ± 0.1	6.6	11.9
0.05	0.56	256	156	30–300	0.150, 0.143	0.12	2.81	14.9 ± 0.6	16.6	13.9
0.05	0.40	512	33	30–330	0.102, 0.107	0.16	1.87	60.5 ± 1.1	167	25.9

[a] $W(m)$ and $D_\nu(m)$ at $m = 65$, in the collisional range, from the slow fluctuation of W; on the fit range the four block fits of ℓ_* agree to within 15%. [b] $D_\perp$ at $m = 0$; every other average within 19%. The window starts after $28.5\,\tau_A$, because W was still settling before, as in the run at $\beta_i = 0.3$ and $M = 32$.

fast, in which case T levels off once the inertial range is long enough, or the perpendicular cascade of the slices has no dissipative anomaly, in which case T, and with it m_*, keeps growing as the dissipation weakens. Up to 256^3 we see no sign of levelling off, but our runs cannot exclude it at higher resolution.

7.8 Forcing the parallel flow

Meyrand et al. (2019) forced the compressive fluctuations through g_1, the parallel flow, on modes with $|k_z| \le 2k_1$. To see what this changes, we repeated the run at $\beta_i = 1$, 64^3 and $M = 64$ with g_1 forced in place of g_0, on the same band of $k_\perp$ but on $|k_z| = k_1$ and $2k_1$, at the same rate. The Alfvénic turbulence is the same, with $\chi = 0.95$ and $\tau_{\rm nl} = 0.153$, and so is the perpendicular sink: $1/T = 0.427$ against 0.421, with $D_\perp(m)/W(m)$ rising from 0.34 at $m \le 1$ to 0.45 at $m = 16$, as with g_0 forced. The injected pattern has $\kappa_1 = 1.63\,k_1$ along the field lines, against $1.12\,k_1$.

The spectrum is steeper at small m, with a local slope of -0.66 over $2 \le m \le 8$ against -0.54, and the same at larger m, -0.71 over $8 \le m \le 32$ against -0.72; it is nowhere close to m^{-1}. The model with the traced $P(\kappa)$, injected at $m = 1$, reproduces $W(m)$ to within 3.5% for $0 \le m \le 32$, and the field-aligned flux lies 2–4% above it for $m \ge 1$, with no deficit at large m, so forcing g_1 on two parallel wavenumbers adds no return flux. At $m = 16$ the field-aligned flux carries 75% of the injected free energy, against 66% with g_0 forced, and the hypercollisions remove 61% of it, against 50%, because the larger injected κ carries the free energy to the collisional range sooner. The model gives 60% and 48%.

The k_z split behaves differently. Its unmixing part is -0.39 times the phase-mixing part at the forced moment, -0.17 at $m = 2$ and -0.08 to -0.11 for $4 \le m \le 16$, while the tilt estimate of prediction (v) gives 0.03–0.04 for $m \ge 2$. The estimate assumes a single field-aligned wavenumber and a Gaussian tilt independent of it, and with this forcing it accounts for less than half of the measured ratio for $2 \le m \le 16$. Since the field-aligned flux follows the model, the excess of the split is not a return flux along the field.

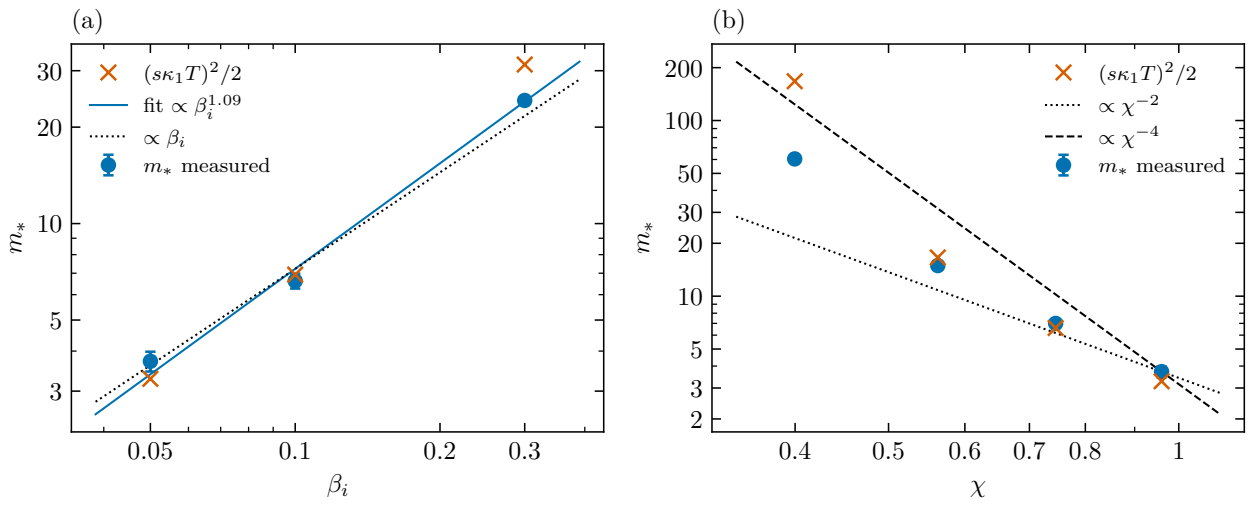


Figure 10: The break $m_* = \ell_*^2/2$ from the fit to $W(m)$ (circles, with the standard error of the four block fits) and the model value $(s\kappa_1 T)^2/2$ (crosses). (a) Against β_i at $\chi \approx 0.95$, with a weighted power-law fit (solid) and $m_* \propto \beta_i$ (dotted). (b) Against the outer-scale χ at $\beta_i = 0.05$, with χ^{-2} (dotted) and χ^{-4} (dashed) through the run at $\chi = 0.96$.

8 Discussion

Two results of this paper are exact. The compressive sector of KRMHD is a continuum of passive slices labelled by phase velocity, for each branch and in truncated hierarchies, for any smooth Alfvénic fields and without collisions (theorems 1 and 2), so the nonlinearity moves no free energy between phase velocities. With an ideal induction equation, smooth fields and no spatial dissipation of the compressive distribution, every volume-integrated single-point statistic, including $W(m)$ and the field-aligned Hermite flux, evolves as in linear phase mixing along the initial field lines (theorem 3), so such dynamics produces no stochastic echoes in $W(m)$.

The two-channel model extends these results to turbulence, where they do not apply directly, by treating spatial dissipation as a sink that removes free energy from every Hermite moment at the passive-scalar cascade rate. In the pilot simulation at $\beta_i = 1$ and $\chi \approx 1$ the model holds to the accuracy of the measurement. The perpendicular dissipation is close to an m-independent sink, and $W(m)$ and the field-aligned flux are those of phase mixing at the injected field-aligned wavenumber with that sink, with no deficit at large m. At 64^3 the density moment acquires substantial structure along the field only at the perpendicular scales where the hyper-resistivity acts. In this run the relabelling of field lines that Meyrand et al. (2019) invoke happens, but only to free energy that is being removed at the same scales, and the free energy that phase mixes keeps its injected wavenumber.

Our spectrum at $\beta_i = 1$ differs from that of Meyrand et al. (2019), who found $W(m)$ somewhat steeper than m^{-1} at the same β_i and a net flux to high m much smaller than the injection rate. Their set-up differs from ours in four ways that matter here. They forced the compressive fluctuations at $m = 1$, the parallel flow, on modes with $|k_x|, |k_y|, |k_z| \leq 2$, where we force g_0 at $|k_z| = k_1$; they damped g with $\mu k_\perp^8$ and νm^6, where we use a fourth power of $k_\perp$ normalised to the grid and $\nu(m/M)^6$; they ran at 256^3 to 512^3 with $M = 32$ to 128, against our 64^3 to 256^3 with $M = 32$ and 64; and their Hermite flux is resolved in $k_\perp$ shells, which only in its sum over shells is a pure flux in m. In our model each of the first three differences raises the break rather than lowering it: forcing up to $|k_z| = 2k_1$ raises the injected κ, and a longer inertial range with sharper dissipation lengthens

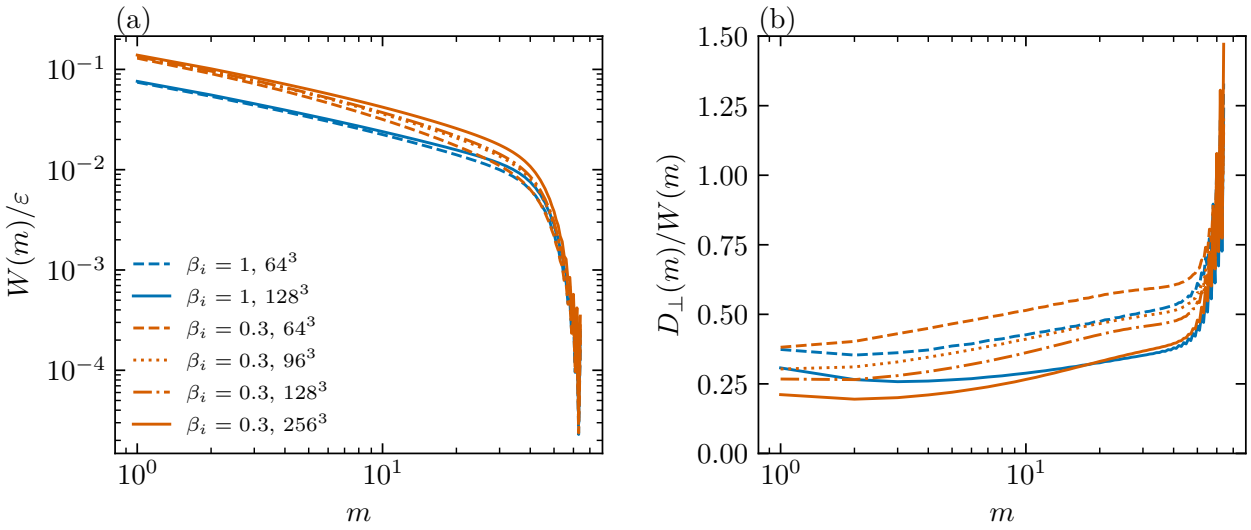


Figure 11: Resolution in $k_\perp$ at $M = 64$, with the hyper-resistivity normalised to the grid: $\beta_i = 1$ (blue) at 64^3 (dashed) and 128^3 (solid), and $\beta_i = 0.3$ (orange) at 64^3 (dashed), 96^3 (dotted), 128^3 (dash-dotted) and 256^3 (solid). (a) $W(m)/\varepsilon$. (b) $D_\perp(m)/W(m)$.

T, as our runs at 96^3 to 256^3 show, and $m_* = (s\kappa T)^2/2$ grows with both. The model therefore does not account for their steeper spectrum by these differences alone. We tested the forcing directly (section 7.8). With g_1 forced on $|k_z| = k_1$ and $2k_1$ at 64^3, $W(m)$ steepens at small m, from $m^{-0.54}$ to $m^{-0.66}$ over $2 \le m \le 8$, keeps a slope of about -0.7 beyond, and is reproduced by the model with the injected $P(\kappa)$ to within 4%, with no return flux; at $m = 16$ the net flux still carries 75% of the injection. The forcing at $m = 1$ therefore does not produce their spectrum at our resolution. That leaves the different dissipation of g and effects that appear only at their higher resolution, such as echoes that our runs do not resolve; in our runs a higher resolution makes the spectrum shallower, not steeper. Running our field-aligned diagnostics on their set-up would decide it. A diagnostic built on the Eulerian k_z reports a backward flux of about 12% of the forward flux in our runs at 64^3, about 20% at 128^3 and about 30% at 256^3, where the field-aligned measures show none, so splits and fluxes of the free energy should be computed along the perturbed field. And the field-aligned wavenumber of g measured with k_z grows with $k_\perp$ much as that of the Alfvénic fields does, whereas measured along the field lines it grows more slowly, at 256^3 from $1.0\,k_1$ to $1.6\,k_1$ between $k_\perp = 1.5\,k_1$ and $6\,k_1$, against $1.1\,k_1$ to $2.2\,k_1$ for Φ.

At fixed χ the model predicts that m_* grows with β_i, as $\beta_i T^2$. At $\beta_i = 1$ the break lies beyond the highest Hermite moment we keep, $m = M$; at $\beta_i = 0.1$ it lies well below it and $W(m)$ steepens from small m, as predicted, with the same Alfvénic turbulence in every run. The trend that Kawazura et al. (2019, 2020) found in gyrokinetic simulations, $m^{-1/2}$ at high β_i and close to m^{-1} at $\beta_i = 0.1$, is the trend the model gives without any echo, because m_* falls with β_i. Fitted to the measured spectra, the decay rate in $\sqrt{2m}$ is the predicted $1/(s\kappa_1 T)$ to within 7% at every β_i and both M. With M raised until the break is resolved, m_* grows as $\beta_i^{1.09\pm0.04}$ at $\chi \approx 0.95$ for $0.05 \le \beta_i \le 0.3$, close to the $m_* \propto \beta_i$ of (49). Extending this measurement to $\beta_i = 1$, where resolving the break needs M of several hundred, is ongoing work. The comparison also showed where the model is too simple: the perpendicular sink is not uniform in m but grows with the age of the free energy, and with that refinement (51) the sinks of all runs collapse onto one function of age. Against χ the estimate $m_* \sim \beta_i/\chi^2$ holds only near $\chi \approx 1$. At $\beta_i = 0.05$ the local exponent

of m_* in χ is -2.5 and -2.7 down to $\chi \approx 0.56$ and -4.1 between 0.56 and 0.40. In units of the outer eddy time the sink is close to the same function of age at $\chi \geq 0.56$ and weaker at $\chi = 0.40$, as expected when the outer-scale turbulence becomes weak. With the rate averaged over the range where the break is fitted, the decay length is still $s\kappa_1 T$ to within 9%, so what fails at small χ is the assumption $T \propto \tau_{\rm nl}$, not the competition between phase mixing and the sink. The model also predicts that the share of the free energy dissipated by collisions falls as the collisional cutoff m_c rises, as $\exp(-\sqrt{2m_c}/(s\kappa T))$; the runs at $M = 32$ and $M = 64$ show this fall, and with the traced $P(\kappa)$ the model gives the shares to within 0.06. If this holds as $m_c \to \infty$, the free energy that phase mixes is in the end removed at small perpendicular scales, and the fraction dissipated by collisions in velocity space at fixed β_i and χ is set by where the collisions act; in a real plasma the perpendicular cascade of the slices stops at the ion gyroscale, beyond the reach of KRMHD.

The main open question is the limit of high resolution. The model needs a perpendicular cascade whose rate does not depend on the dissipation coefficients and field lines that stay frozen for the free energy that phase mixes. Drivas and Eyink (2017) and Eyink (2018) give reasons to expect that the first requires the failure of the second at small enough scales (section 6.4). Our runs do not reach a cascade rate independent of the dissipation. Up to 256^3 grid points T grows by about the same amount with each doubling of the resolution, from 13 to 25 outer eddy times at $\beta_i = 0.3$, and the density moment gains field-aligned structure gradually through the longer inertial range, although the field-aligned wavenumber of the moments as a whole rises only from about $1.25\,k_1$ to $1.35\,k_1$ (section 7.7). With the T and κ measured at each resolution the model still describes $W(m)$ to within -9% and $+13\%$, and there is still no return flux. The model therefore describes the spectrum at each resolution with the T measured there, but in our runs T, and with it m_*, has no value independent of the dissipation. Whether T converges once the inertial range is long enough, or keeps growing so that the perpendicular cascade loses ground to phase mixing as the dissipation weakens, needs higher resolution than we reach.

The results hold for $k_\perp \rho_i \to 0$. At finite $k_\perp \rho_i$ the Landau coupling depends on the wavenumber, the advection depends on $v_\perp$ and the compressive fluctuations act back on the Alfvénic ones (Chandran et al., 2026), so the slices no longer decouple. Velocity-dependent collisions are the only terms in KRMHD that exchange free energy between slices (18).

A The two compressive branches

We derive the branches of section 2.2 following S09 (§6.2) and CMM (appendix B), with $\delta_n \equiv \delta n/n$, $\delta_B \equiv \delta B/B$ and $\langle\cdot\rangle$ the volume average. We checked every identity symbolically, and the moments and splits also by quadrature over $(v_\perp, v_\parallel)$ for random distribution functions.

A.1 Laguerre projection and the branches

With $x = v_\perp^2/v_{\rm th}^2$, CMM (eq. B8) and S09 (eq. 205) write

$$g = \frac{n_i}{\pi v_{\rm th}^2}\, e^{-x} \sum_{l\geq 0} L_l(x)\, G_l(v_\parallel), \qquad F_0 = \frac{n_i}{\pi v_{\rm th}^2}\, e^{-x} F_M(v_\parallel),$$

with L_l orthonormal with weight e^{-x} on $(0,\infty)$. Since $x = L_0 - L_1$ and $x^k \perp L_l$ for $l > k$, the source term of (3) projects only on L_0 and L_1, and (3) and (4) become

$$\frac{\mathrm{d}G_0}{\mathrm{d}t} + v_\parallel \nabla_\parallel \Big[G_0 + F_M\Big(\frac{Z}{\tau}\delta_n + \delta_B\Big)\Big] = 0, \qquad \frac{\mathrm{d}G_1}{\mathrm{d}t} + v_\parallel \nabla_\parallel (G_1 - F_M \delta_B) = 0,$$
$$\frac{\mathrm{d}G_l}{\mathrm{d}t} + v_\parallel \nabla_\parallel G_l = 0 \quad (l \ge 2),$$
$$\int G_0 \,\mathrm{d}v_\parallel = \delta_n - \delta_B, \qquad \int (G_0 - G_1)\,\mathrm{d}v_\parallel = -\frac{Z}{\tau}\delta_n - 2\Big(1 + \frac{1}{\beta_i}\Big)\delta_B.$$

With $D = Z/\tau + 2(1 + \beta_i^{-1})$, the distributions of CMM's eqs (B2)–(B3) (S09 eqs 176–177),

$$G_n = -\frac{1}{D}\Big[(G_0 - G_1) - 2\Big(1 + \frac{1}{\beta_i}\Big)G_0\Big], \qquad G_B = -\frac{1}{D}\Big[(G_0 - G_1) + \frac{Z}{\tau}G_0\Big],$$

have $\int G_n \,\mathrm{d}v_\parallel = \delta_n$ and $\int G_B \,\mathrm{d}v_\parallel = \delta_B$, invert to

$$G_0 = G_n - G_B, \qquad G_1 = \Big(1 + \frac{Z}{\tau}\Big)G_0 + D\,G_B, \tag{A1}$$

and obey S09's eqs (179)–(180),

$$\frac{\mathrm{d}}{\mathrm{d}t}\begin{pmatrix} G_n \\ G_B \end{pmatrix} + v_\parallel \nabla_\parallel \begin{pmatrix} G_n \\ G_B \end{pmatrix} = v_\parallel F_M \,\nabla_\parallel K \begin{pmatrix} \delta_n \\ \delta_B \end{pmatrix}, \qquad K = \frac{1}{D}\begin{pmatrix} -\frac{Z}{\tau}\Big(1 + \frac{2}{\beta_i}\Big) & -\frac{2}{\beta_i} \\ \frac{Z}{\tau}\Big(1 + \frac{Z}{\tau}\Big) & 2 + \frac{Z}{\tau} \end{pmatrix}.$$

If $\ell K = \Lambda^{-1}\ell$, then $G = \ell \cdot (G_n, G_B)$ has $\int G \,\mathrm{d}v_\parallel = \ell \cdot (\delta_n, \delta_B)$ and obeys (5). The two left eigenvectors give CMM's branches (B5)–(B6) (S09 eqs 183–184),

$$G^+ = G_B + \frac{1 + Z/\tau}{\sigma}\,G_n, \qquad G^- = G_n + \frac{2\tau}{\sigma Z \beta_i}\,G_B, \qquad \sigma = 1 + \frac{\tau}{Z} + \frac{1}{\beta_i} + \Delta,$$

and the eigenvalues are $1/\Lambda^\pm$ with $\Lambda^\pm$ and Δ of (6).

A.2 Moments

The map $(G_n, G_B) \mapsto (G^+, G^-)$ has determinant $-2\Delta/\sigma \neq 0$ and inverse

$$G_n = \frac{\sigma G^- - (2\tau/Z\beta_i)\,G^+}{2\Delta}, \qquad G_B = \frac{\sigma G^+ - (1 + Z/\tau)\,G^-}{2\Delta}, \tag{A2}$$

so that, with $M_0^\pm = \int G^\pm \,\mathrm{d}v_\parallel$,

$$\delta_n = \frac{\sigma M_0^- - (2\tau/Z\beta_i)\,M_0^+}{2\Delta}, \qquad \delta_B = \frac{\sigma M_0^+ - (1 + Z/\tau)\,M_0^-}{2\Delta}. \tag{A3}$$

These are one half of the right-hand sides of S09's eqs (202)–(203), as CMM found (footnote to their eq. B15). The parallel flow and pressure perturbations of the non-Alfvénic part $g_i = g + x\,\delta_B F_0$ of the ion distribution (CMM eq. 3.26) are

$$u_\parallel = \int v_\parallel G_0 \,\mathrm{d}v_\parallel, \qquad \frac{\delta p_\perp}{p} = -\frac{Z}{\tau}\delta_n - \frac{2}{\beta_i}\delta_B, \qquad \frac{\delta p_\parallel}{p} = \frac{2}{v_{\mathrm{th}}^2}\int v_\parallel^2 G_0 \,\mathrm{d}v_\parallel + \delta_B,$$

with $p = n_i T_i$. By (A1) and (A2) these are linear in $G^\pm$, and the G_l with $l \ge 2$ enter only moments $\int \mathrm{d}^3 v\, v_\parallel^j v_\perp^{2k} g_i$ with $k \ge 2$.

A.3 Free energy

With $x = L_0 - L_1$, the orthonormality of the L_l and the field relations, CMM's compressive free energy (their eq. 3.31) becomes

$$\begin{aligned} W_{\mathrm{compr}} &= \frac{n_i T_i}{2}\Big\langle \frac{Z}{\tau}\delta_n^2 + \frac{2}{\beta_i}\delta_B^2 + \frac{1}{n_i}\int \mathrm{d}^3 v\, \frac{g_i^2}{F_0}\Big\rangle \\ &= \frac{n_i T_i}{2}\Big\langle \int \frac{G_0^2+G_1^2}{F_M}\,\mathrm{d}v_\parallel + \frac{Z}{\tau}\delta_n^2 - \frac{2Z}{\tau}\delta_n\delta_B - 2\Big(1+\frac{1}{\beta_i}\Big)\delta_B^2 + \sum_{l\ge 2}\int \frac{G_l^2}{F_M}\,\mathrm{d}v_\parallel\Big\rangle. \end{aligned}$$

The part with $l = 0, 1$ splits between the branches by the identities

$$G_0^2+G_1^2 = c_+\,(G^+)^2 + c_-\,(G^-)^2, \qquad \frac{Z}{\tau}\delta_n^2 - \frac{2Z}{\tau}\delta_n\delta_B - 2\Big(1+\frac{1}{\beta_i}\Big)\delta_B^2 = -c_+\frac{(M_0^+)^2}{\Lambda^+} - c_-\frac{(M_0^-)^2}{\Lambda^-}, \tag{A4}$$

the first of which holds at each $v_\parallel$, with

$$c_+ = \Big[1 + \frac{1+\tau/Z}{\Delta}\Big](\Lambda^+)^2 > 0, \qquad c_- = \frac{Z^2}{2\tau^2}\Big[1 + \frac{1}{\Delta\beta_i}\Big](\Lambda^-)^2 > 0. \tag{A5}$$

This gives CMM's eq. (B15) with their eqs (B13)–(B14),

$$W_{\mathrm{compr}} = \tilde{W}_{\mathrm{compr}} + c_+ W^+_{\mathrm{compr}} + c_- W^-_{\mathrm{compr}}, \qquad W^\pm_{\mathrm{compr}} = \frac{n_i T_i}{2}\Big\langle \int \frac{(G^\pm)^2}{F_M}\,\mathrm{d}v_\parallel - \frac{(M_0^\pm)^2}{\Lambda^\pm}\Big\rangle, \tag{A6}$$

with $\tilde{W}_{\mathrm{compr}}$ the part with $l \ge 2$. Each part is conserved (S09, §§6.2.3 and 6.2.5; CMM eq. B12) and positive definite (section 3).

The coefficients (A5) use Δ of (6) (S09 eq. 204) and are one quarter of those in S09's eq. (213), as CMM state. CMM's eq. (B16) prints this square root as $[(1+\tau/Z) + \beta_i^{-2}]^{1/2}$, which gives $\sqrt{3}$ instead of $\sqrt{5}$ at $\beta_i = \tau = Z = 1$. Their eqs (B6)–(B7) are correct, so the misprint affects only $c_\pm$ in (B15) and (B17) and our (A6), which then fails. For random distribution functions the relative error in W_{compr} with the printed square root is up to 7.2×10^{-2} over the parameters we tried (2.5×10^{-2} at $\beta_i = \tau = Z = 1$), and below 5×10^{-14} with the square restored.

B The slice transform and the spectral theory of the branch operator

B.1 Definitions and statement

Let $\mathcal{W}$ be the space of real functions $G(u)$ with $\langle G, G\rangle < \infty$, where $\langle f, g\rangle = \int f g\, F_M^{-1}\,\mathrm{d}u$, so that $\mathsf{Q} = F_M\, M_0$ is the orthogonal projector onto F_M. The free-energy product (12) is $[f, g] = \langle f, Pg\rangle$ with $P = I - \mathsf{Q}/\Lambda$, and with $G = M_0[G]\, F_M + G_\perp$, $[G, G] = \langle G_\perp, G_\perp\rangle + p_0\, M_0[G]^2$, which is positive definite, with a norm equivalent to that of $\langle\cdot,\cdot\rangle$, if and only if $\Lambda \notin [0, 1]$. With U the multiplication by u,

$$L = UP, \tag{B1}$$

so $[f, Lg] = \langle Pf, UPg\rangle$ is symmetric in f and g. For $\mathrm{Im}\, z > 0$,

$$(z - L)^{-1}G = \frac{1}{z-u}\left[G - \frac{uF_M}{\Lambda\,\mathcal{K}(z)}\int \frac{G(v)}{z-v}\,\mathrm{d}v\right], \qquad \mathcal{K}(\zeta) = 1 - \frac{1+\zeta Z(\zeta)}{\Lambda}, \tag{B2}$$

where Z is the plasma dispersion function of Fried and Conte (1961), not the charge number; we write $\mathcal{K}$ also for its continuation into the lower half plane, whose zeros are the Landau roots. On the real axis

$$\mathcal{K}(u+\mathrm{i}0) = \mathcal{K}_R(u) + \mathrm{i}\,\mathcal{K}_I(u), \qquad \mathcal{K}_R = 1 - \frac{1 + u\,\mathrm{Re}\,Z(u)}{\Lambda}, \qquad \mathcal{K}_I = -\frac{\pi\,u\,F_M(u)}{\Lambda}, \tag{B3}$$

and $\mathsf{H}[f](u) = \pi^{-1}\mathrm{P}\!\int f(v)\,(v-u)^{-1}\,\mathrm{d}v$. By the argument of Penrose (1960), $\mathcal{K}$ has no zeros in the closed upper half plane for $\Lambda \notin [0,1]$ (section B.5). The spectral density

$$\varrho(u) = \frac{F_M(u)}{|\mathcal{K}(u+\mathrm{i}0)|^2} \tag{B4}$$

is then positive, smooth and even, with $\int \varrho\,\mathrm{d}u = 1/p_0$. For real u the equation $(L-u)F = 0$ has the Case–Van Kampen solutions

$$F_u(v) = \frac{1}{\Lambda}\,\mathrm{P}\frac{vF_M(v)}{v-u} + \mathcal{K}_R(u)\,\delta(v-u), \qquad M_0[F_u] = 1, \qquad M_1[F_u] = u, \tag{B5}$$

and the slice transform is the projection on them in the free-energy product,

$$\mathcal{T}[G](u) = [F_u, G] = \frac{\mathcal{K}_R\,G - \mathcal{K}_I\,\mathsf{H}[G]}{F_M}(u), \qquad \mathcal{T}[LG](u) = u\,\mathcal{T}[G](u), \tag{B6}$$

of which the slice amplitude (13) is the unit-weight form, $a_u = \sqrt{\varrho(u)}\,\mathcal{T}[G](u)$; $\mathcal{A}$ denotes the map $G \mapsto a$.

Theorem 4. *Let $\Lambda \notin [0,1]$, the value $\Lambda = \infty$ included. Then L, with the domain $D(U)$ of U, is self-adjoint in $(\mathcal{W}, [\cdot,\cdot])$; it has no eigenvalues, and its spectrum is the real line, purely absolutely continuous and simple; and $\mathcal{A}$ is a unitary map from $(\mathcal{W}, [\cdot,\cdot])$ onto $L^2(\mathbb{R}, \mathrm{d}u)$ with $\mathcal{A}LG = u\,\mathcal{A}G$ for $G \in D(U)$. In particular (14) holds, $[G, \varphi(L)G] = \int \varphi(u)\,a_u^2\,\mathrm{d}u$ for every bounded Borel function φ, and*

$$G = \int \varrho(u)\,\mathcal{T}[G](u)\,F_u\,\mathrm{d}u = \mathcal{K}_R\,c + \mathcal{K}_I\,\mathsf{H}[c], \qquad c = \sqrt{\varrho}\;a. \tag{B7}$$

The rest of this appendix proves the theorem and the statements about the eigenvalues of L and the zeros of $\mathcal{K}$. For $0 < \Lambda < 1$, $\mathcal{K}$ has one zero in the upper half plane, on the imaginary axis, L has the eigenvalues $\pm\mathrm{i}\gamma$, and the slice decomposition is incomplete.

B.2 Preliminaries

We work in the complexification of $\mathcal{W}$, with $\langle f, g\rangle = \int \bar{f}\,g\,F_M^{-1}\,\mathrm{d}u$. The set $\mathcal{D}_0$ of functions pF_M, with p a polynomial, is dense in $\mathcal{W}$ (it contains the orthonormal basis e_m), lies in $D(U)$, and is mapped into itself by U, P and L. Let $\mathcal{S}$ be the set of smooth f with $|f(v)| \le C(1+|v|)^k F_M(v)$ for some C and k; it contains $\mathcal{D}_0$ and the smooth functions with compact support. For $f \in \mathcal{S}$ the Cauchy integral

$$\chi_f(\zeta) = \frac{1}{\pi}\int \frac{f(v)}{v-\zeta}\,\mathrm{d}v \qquad (\mathrm{Im}\,\zeta > 0)$$

is analytic in the upper half plane, extends continuously to the real axis with $\chi_f(u+\mathrm{i}0) = \mathsf{H}[f](u) + \mathrm{i}f(u)$ (the Plemelj formula), and is $O(1/\zeta)$ as $|\zeta| \to \infty$ in the closed upper half plane. The following converse is used twice.

Lemma 1. *Let Y be analytic in the upper half plane, continuous up to the real axis, and $O(1/\zeta)$ as $|\zeta| \to \infty$ in the closed upper half plane. Then $Y = \chi_y$ with $y(u) = \operatorname{Im} Y(u + \mathrm{i}0)$.*

Proof. For $\operatorname{Im}\zeta > 0$, Cauchy's formula on a large half disc gives $Y(\zeta) = (2\pi\mathrm{i})^{-1} \int Y(v + \mathrm{i}0)\,(v - \zeta)^{-1}\,\mathrm{d}v$, and Cauchy's theorem for $\overline{Y(\bar\zeta)}$, analytic in the lower half plane with the same decay, gives $\int \overline{Y(v + \mathrm{i}0)}\,(v - \zeta)^{-1}\,\mathrm{d}v = 0$. Subtracting $(2\pi\mathrm{i})^{-1}$ times the second integral from the first gives $Y = \chi_y$. □

B.3 Self-adjointness

Let $\Lambda \notin [0, 1]$, so that $p_0 > 0$. Since $P = (I - \mathsf{Q}) + p_0\mathsf{Q}$, the operator $P^{1/2} = I + (\sqrt{p_0} - 1)\mathsf{Q}$ is bounded and symmetric, with bounded inverse, and $[f, g] = \langle P^{1/2}f, P^{1/2}g\rangle$, so $V : G \mapsto P^{1/2}G$ is unitary from $(\mathcal{W}, [\cdot,\cdot])$ onto $(\mathcal{W}, \langle\cdot,\cdot\rangle)$. Since $V - I$ has its range in the span of $F_M \in D(U)$, V maps $D(U)$ onto itself, and by (B1)

$$VLV^{-1} = P^{1/2}UP^{1/2} = U + \mathsf{K}, \qquad \mathsf{K} = (\sqrt{p_0} - 1)(\mathsf{Q}U + U\mathsf{Q}), \tag{B8}$$

where $(\sqrt{p_0} - 1)^2\mathsf{Q}U\mathsf{Q}$ is absent because $\mathsf{Q}U\mathsf{Q}G = M_1[F_M]\,M_0[G]\,F_M = 0$. The operator $\mathsf{K}G = (\sqrt{p_0}-1)\big[\langle uF_M, G\rangle\,F_M + \langle F_M, G\rangle\,uF_M\big]$ is symmetric, of rank two and bounded, since $\langle F_M, F_M\rangle = 1$ and $\langle uF_M, uF_M\rangle = 1/2$. By the Kato–Rellich theorem (Kato, 1966, chapter V), $U + \mathsf{K}$ is then self-adjoint on $D(U)$ in $(\mathcal{W}, \langle\cdot,\cdot\rangle)$, so $L = V^{-1}(U + \mathsf{K})V$ is self-adjoint on $D(U)$ in $(\mathcal{W}, [\cdot,\cdot])$, which is part (i) of theorem 4.

B.4 Eigenvalues

For every $\Lambda \neq 0$, a function $G \in D(U)$ is an eigenfunction of L with eigenvalue λ if and only if

$$(u - \lambda)\,G = \frac{uF_M}{\Lambda}\,M_0[G], \qquad G \neq 0.$$

If λ is not real, $G = M_0[G]\,uF_M/[\Lambda(u - \lambda)]$, and taking M_0 gives $M_0[G]\,\mathcal{K}(\lambda) = 0$ if $\operatorname{Im}\lambda > 0$ and $M_0[G]\,\overline{\mathcal{K}(\bar\lambda)} = 0$ if $\operatorname{Im}\lambda < 0$, because the integral in (B2) at a point of the lower half plane equals $\overline{\mathcal{K}(\bar\lambda)}$ (for the same reason the resolvent for $\operatorname{Im} z < 0$ is (B2) with $\overline{\mathcal{K}(\bar z)}$ in place of $\mathcal{K}(z)$). Since $M_0[G] = 0$ would give $G = 0$, a non-real λ is an eigenvalue if and only if λ or $\bar\lambda$ is a zero of $\mathcal{K}$ in the upper half plane. For real $\lambda \neq 0$, $uF_M/(u - \lambda)$ is not square integrable near $u = \lambda$, so $M_0[G] = 0$ and $G = 0$. For $\lambda = 0$, $G = M_0[G]\,F_M/\Lambda$ requires $\Lambda = 1$.

B.5 Zeros of the dielectric function

The function $\mathcal{K}$ is analytic in the upper half plane and continuous up to the real axis, where it takes the values (B3). Since $1 + \zeta Z(\zeta) = -1/(2\zeta^2) + O(\zeta^{-4})$ uniformly in $\operatorname{Im}\zeta \geq 0$, $\mathcal{K} \to 1$ as $|\zeta| \to \infty$ there, so if $\mathcal{K}$ has no real zeros, the number of its zeros in the upper half plane is the winding number about 0 of the closed curve $u \mapsto \mathcal{K}(u + \mathrm{i}0)$. A real zero requires $\mathcal{K}_I(u) = 0$, hence $u = 0$, and $\mathcal{K}_R(0) = p_0$ vanishes only at $\Lambda = 1$. For $\Lambda \neq 1$, $\mathcal{K}(u + \mathrm{i}0) = 1 - \mathcal{N}(u)/\Lambda$ with $\mathcal{N}(u) = 1 + uZ(u + \mathrm{i}0)$, and the winding number of $\mathcal{K}$ about 0 equals that of $\mathcal{N}$ about the point Λ. The curve $\mathcal{N}$ has $\operatorname{Im}\mathcal{N}(u) = \pi u F_M(u)$, $\mathcal{N}(0) = 1$ and $\mathcal{N}(u) \to 0$ as $|u| \to \infty$: it leaves 0 at $u = -\infty$, runs through the lower half plane to the point 1, reached at $u = 0$, and returns to 0 through the upper half plane, meeting the real axis only at 0 and 1. It therefore winds once around every point of $(0, 1)$ and around no real point outside $[0, 1]$, so

$$\#\{\zeta :\ \mathcal{K}(\zeta) = 0,\ \operatorname{Im}\zeta \geq 0\} = \begin{cases} 0, & \Lambda < 0 \text{ or } \Lambda > 1, \\ 1, & 0 < \Lambda < 1, \end{cases} \tag{B9}$$

which is the criterion of Penrose (1960) for this dielectric function. By section B.4, L has no eigenvalues for $\Lambda \notin [0,1]$.

The plasma dispersion function continues to the entire function $Z(\zeta) = \mathrm{i}\sqrt{\pi}\, e^{-\zeta^2}\operatorname{erfc}(-\mathrm{i}\zeta)$, and the zeros of the continued $\mathcal{K}$ in the lower half plane are the Landau roots. They are not eigenvalues of L, because in the lower half plane the resolvent contains $\overline{\mathcal{K}(\bar{\zeta})}$, which by (B9) has no zeros there for $\Lambda \notin [0,1]$.

B.6 Spectral measure

For $\operatorname{Im} z > 0$ and real $G \in \mathcal{W}$, let

$$C_G(z) = \int \frac{G(v)}{z - v}\,\mathrm{d}v, \qquad I_G(z) = \int \frac{G(v)^2}{(z-v)\,F_M(v)}\,\mathrm{d}v.$$

By (B2), $M_0[(z-L)^{-1}G] = C_G(z)/\mathcal{K}(z)$, and since $\int uG(u)\,(z-u)^{-1}\,\mathrm{d}u = zC_G(z) - M_0[G]$, the resolvent form is

$$R_G(z) \equiv [G, (z-L)^{-1}G] = I_G(z) - \frac{z}{\Lambda}\,\frac{C_G(z)^2}{\mathcal{K}(z)}.$$

It is the Stieltjes transform $\int (z-x)^{-1}\,\mathrm{d}\mu_G(x)$ of the spectral measure μ_G of L in the state G, which Stone's formula recovers:

$$\mu_G((u_1,u_2)) + \tfrac{1}{2}\mu_G(\{u_1\}) + \tfrac{1}{2}\mu_G(\{u_2\}) = -\frac{1}{\pi}\lim_{\eta\to 0+}\int_{u_1}^{u_2} \operatorname{Im} R_G(u+\mathrm{i}\eta)\,\mathrm{d}u. \tag{B10}$$

Let $G \in \mathcal{D}_0$ be real and $\Lambda \notin [0,1]$. Then $C_G = -\pi\chi_G$ and $I_G = -\pi\chi_{G^2/F_M}$ extend continuously to the real axis, with $C_G(u+\mathrm{i}0) = -\pi(\mathsf{H}[G] + \mathrm{i}G)(u)$ and $\operatorname{Im} I_G(u+\mathrm{i}0) = -\pi G(u)^2/F_M(u)$. By (B9), $1/\mathcal{K}$ is continuous on the closed upper half plane, so $R_G(u+\mathrm{i}\eta) \to R_G(u+\mathrm{i}0)$ as $\eta \to 0+$, uniformly for u in compact sets. With $\mathcal{K}_I = -\pi u F_M/\Lambda$, the imaginary part of the boundary value is a square:

$$-\frac{1}{\pi}\operatorname{Im} R_G(u+\mathrm{i}0) = \frac{G^2}{F_M} - \frac{\mathcal{K}_I}{F_M}\,\frac{2\mathcal{K}_R\,G\,\mathsf{H}[G] - \mathcal{K}_I(\mathsf{H}[G]^2 - G^2)}{|\mathcal{K}|^2} = \frac{\left(\mathcal{K}_R G - \mathcal{K}_I \mathsf{H}[G]\right)^2}{F_M\,|\mathcal{K}|^2} = \varrho\,\mathcal{T}[G]^2 = a_u^2.$$

By (B10) and the uniform convergence, μ_G has no atoms and $\mathrm{d}\mu_G = a_u^2\,\mathrm{d}u$.

Three consequences follow. First, the absolutely continuous subspace of L is closed and contains the dense set $\mathcal{D}_0$, so it is $\mathcal{W}$, and L has no eigenvalues and no singular continuous spectrum. Second, K in (B8) has rank two and is therefore of trace class, so by the Kato–Rosenblum theorem (Kato, 1966, chapter X) the absolutely continuous parts of $U + \mathsf{K} = VLV^{-1}$ and of U are unitarily equivalent; since both operators are purely absolutely continuous, L is unitarily equivalent to U, multiplication by u in $L^2(\mathbb{R}, \mathrm{d}u/F_M)$, whose spectrum is the real line with multiplicity one. This is part (ii) of theorem 4. Third, applying $\mathrm{d}\mu_G = a_u^2\,\mathrm{d}u$ to $G+G'$ and $G-G'$ gives, for real $G, G' \in \mathcal{D}_0$ and every bounded Borel function φ,

$$[G', \varphi(L)G] = \int \varphi(u)\,a_u[G']\,a_u[G]\,\mathrm{d}u, \qquad a[G] = \mathcal{A}G, \tag{B11}$$

and by linearity the same holds for complex $G, G' \in \mathcal{D}_0$ with $a_u[G']$ replaced by its complex conjugate.

B.7 The slice transform

The formula (13), with H the Hilbert transform on $L^2(\mathbb{R}, \mathrm{d}u)$, defines a bounded operator $\mathcal{A}$ from $\mathcal{W}$ into $L^2(\mathbb{R}, \mathrm{d}u)$: for $\Lambda \notin [0,1]$, $|\mathcal{K}(u+\mathrm{i}0)|$ is bounded below, since it has no real zeros and tends to 1 at infinity, and since $|\mathcal{K}_R| \le |\mathcal{K}|$, the first term of (13) has norm at most $\langle G, G\rangle^{1/2}$. In the second, $\mathcal{K}_I \mathsf{H}[G]/(\sqrt{F_M}\,|\mathcal{K}|) = -\pi u \sqrt{F_M}\,\mathsf{H}[G]/(\Lambda|\mathcal{K}|)$, the factor $u\sqrt{F_M}/|\mathcal{K}|$ is bounded, H is an isometry, and $\int |G|^2\,\mathrm{d}u \le \pi^{-1/2}\langle G, G\rangle$. By (B11) with $\varphi = 1$, $\mathcal{A}$ is isometric on the dense set $\mathcal{D}_0$, so it is an isometry of $(\mathcal{W}, [\cdot,\cdot])$ into $L^2(\mathbb{R}, \mathrm{d}u)$, and its range is closed.

The following argument is the counterpart for L of the invertibility proof of Morrison (2000). To show that the range of $\mathcal{A}$ is dense, let b be a smooth function with compact support, $c = \sqrt{\varrho}\,b$, and

$$G = \mathcal{K}_R c + \mathcal{K}_I \mathsf{H}[c] = \operatorname{Im}\bigl[\mathcal{K}(u+\mathrm{i}0)\,\chi_c(u+\mathrm{i}0)\bigr].$$

Since c has compact support, $\mathcal{K}_I = -\pi u F_M/\Lambda$ and $\mathsf{H}[c]$ is smooth and bounded, G belongs to $\mathcal{S}$ and to $\mathcal{W}$. The product $\mathcal{K}\chi_c$ is analytic in the upper half plane, continuous up to the real axis and $O(1/\zeta)$, so by lemma 1 it equals χ_G, and $\chi_c(u+\mathrm{i}0) = \chi_G(u+\mathrm{i}0)/\mathcal{K}(u+\mathrm{i}0)$. Taking imaginary parts, with the Plemelj formula for χ_G,

$$c = \frac{\operatorname{Im}\bigl[(\mathsf{H}[G] + \mathrm{i}G)\,\bar{\mathcal{K}}\bigr]}{|\mathcal{K}|^2} = \frac{\mathcal{K}_R G - \mathcal{K}_I \mathsf{H}[G]}{|\mathcal{K}|^2} = \varrho\,\mathcal{T}[G],$$

so that $\mathcal{A}G = b$. The smooth functions with compact support are dense in $L^2(\mathbb{R}, \mathrm{d}u)$ and the range of $\mathcal{A}$ is closed, so $\mathcal{A}$ is onto, with the inverse

$$\mathcal{A}^{-1} a = \mathcal{K}_R \sqrt{\varrho}\,a + \mathcal{K}_I\,\mathsf{H}\bigl[\sqrt{\varrho}\,a\bigr], \tag{B12}$$

which is (B7): since $|\mathcal{K}_R|\sqrt{\varrho/F_M} \le 1$ and $|\mathcal{K}_I|/\sqrt{F_M}$ and ϱ are bounded, the right-hand side is a bounded operator from $L^2(\mathbb{R}, \mathrm{d}u)$ into $\mathcal{W}$ that equals $\mathcal{A}^{-1}$ on a dense set. Conversely, for $G \in \mathcal{D}_0$, $Y = \chi_G/\mathcal{K}$ is analytic in the upper half plane, because $\mathcal{K}$ has no zeros there, and $\operatorname{Im} Y(u+\mathrm{i}0) = \varrho\mathcal{T}[G] = c \in \mathcal{S}$; by lemma 1, $Y = \chi_c$, and $G = \operatorname{Im}\chi_G(u+\mathrm{i}0) = \operatorname{Im}[\mathcal{K}\,Y](u+\mathrm{i}0) = \mathcal{K}_R c + \mathcal{K}_I \mathsf{H}[c]$. The first direction does not use the zeros of $\mathcal{K}$ in the upper half plane; the second uses their absence, and fails by the residue of Y at the zero $\mathrm{i}\gamma$ that $\mathcal{K}$ has for $0 < \Lambda < 1$, where $1 - \sqrt{\pi}\,\gamma\,e^{\gamma^2}\operatorname{erfc}(\gamma) = \Lambda$ and L has the eigenvalues $\pm\mathrm{i}\gamma$ (section B.4); the slices are then not complete. The KRMHD branches never fall in this range, because $\Lambda^+ > 1$ and $\Lambda^- < 0$.

Part (iii) of theorem 4 follows. By (B11) with $\varphi(u) = (z-u)^{-1}$, $\operatorname{Im} z \ne 0$, $[G', (z-L)^{-1}G] = \int \overline{a_u[G']}\,a_u[G]\,(z-u)^{-1}\,\mathrm{d}u$ for $G, G' \in \mathcal{D}_0$, and by continuity for all $G, G' \in \mathcal{W}$. Since $\mathcal{A}$ is unitary, the left-hand side is $\int \overline{a_u[G']}\,a_u[(z-L)^{-1}G]\,\mathrm{d}u$, and since $\mathcal{A}G'$ runs over $L^2(\mathbb{R}, \mathrm{d}u)$, $\mathcal{A}(z-L)^{-1}G = (z-u)^{-1}\mathcal{A}G$. For $G \in D(U)$, applying this to $(z-L)G$ shows that $u\,\mathcal{A}G \in L^2(\mathbb{R}, \mathrm{d}u)$ and $\mathcal{A}LG = u\,\mathcal{A}G$. The relations (14) are then the functional calculus of multiplication by u.

By (B5), $\int c(u)\,F_u(v)\,\mathrm{d}u = \mathcal{K}_R(v)\,c(v) + \mathcal{K}_I(v)\,\mathsf{H}[c](v)$, so (B12) is the completeness relation $G = \int \varrho\,\mathcal{T}[G]\,F_u\,\mathrm{d}u$ of the modes, and $\varrho\,\mathcal{T}[\mathcal{A}^{-1}a] = \sqrt{\varrho}\,a$ is the orthogonality relation $[F_u, F_{u'}] = \delta(u-u')/\varrho(u)$, in the sense of distributions.

C The truncated hierarchy

C.1 Weights, closures and eigenvectors

Let J be the ladder matrix of the zero closure at $\Lambda = \infty$, the symmetric Jacobi matrix of the Hermite polynomials h_m orthonormal with weight F_M, and J_c that of the copy closure, which for $M \ge 2$ differs from J only in $(J_c)_{M,M-1} = \sqrt{M/2} + \sqrt{(M+1)/2}$. A finite Λ changes only

$A_{10} = (1-1/\Lambda)\,J_{10}$, and column 0 of J and of J_c has the single non-zero entry $J_{10} = 1/\sqrt{2}$. Hence, with $P_\Lambda = \mathrm{diag}(1-1/\Lambda, 1, \ldots, 1)$,

$$A = J\,P_\Lambda \quad \text{(zero closure)}, \qquad A = J_c\,P_\Lambda \quad \text{(copy closure, } M \geq 2\text{)}, \tag{C1}$$

the Hermite form of $L = UP$ (B1). For the zero closure $\mathsf{P} = P_\Lambda$, and $\mathsf{P}A = P_\Lambda J P_\Lambda$ is symmetric. For the copy closure $D_c = \mathrm{diag}(1, \ldots, 1, p_M)$ satisfies $D_c J_c = J$, because $p_M (J_c)_{M,M-1} = \sqrt{M/2}$, so that with $\mathsf{P} = D_c P_\Lambda$ again $\mathsf{P}A = P_\Lambda J P_\Lambda$.

Since $LF_u = uF_u$ and the untruncated ladder matrix is the matrix of L in the Hermite basis, the Hermite coefficients $r_m(u)$ of F_u satisfy $\sum_n A_{mn} r_n = u\,r_m$ row by row, with $r_0 = M_0[F_u] = 1$:

$$r_1 = \sqrt{2}\,u, \qquad r_2 = u\,r_1 - \frac{1-1/\Lambda}{\sqrt{2}}\,r_0, \qquad \sqrt{\frac{m+1}{2}}\,r_{m+1} = u\,r_m - \sqrt{\frac{m}{2}}\,r_{m-1} \quad (m \geq 2). \tag{C2}$$

Thus r_m is a polynomial of degree m with $r_m(-u) = (-1)^m r_m(u)$, and $r_m = h_m$ at $\Lambda = \infty$. Rows 0 to $M-1$ of the truncated matrix coincide with those of the infinite matrix for both closures, and row M is the condition on the eigenvalues,

$$r_{M+1}(u_j) = 0 \quad \text{(zero closure)}, \qquad r_{M+1}(u_j) = r_{M-1}(u_j) \quad \text{(copy closure)}, \tag{C3}$$

the second from row M of the copy closure, $u\,r_M = [\sqrt{M/2} + \sqrt{(M+1)/2}\,]\,r_{M-1}$, and the recurrence at $m = M$.

Proof of theorem 2. (i) A is similar to the symmetric tridiagonal matrix $\mathsf{P}^{1/2} A\,\mathsf{P}^{-1/2}$, whose off-diagonal entries are non-zero because $A_{01}A_{10} = (1-1/\Lambda)/2 > 0$, so its eigenvalues are real and simple and the vectors $\mathsf{P}^{1/2}\boldsymbol{r}_j$ are orthogonal. By (C2) the vector $(r_0(u), \ldots, r_M(u))^\mathsf{T}$ satisfies rows 0 to $M-1$ of $A\boldsymbol{r} = u\boldsymbol{r}$ for every u, (C3) fixes the u_j, and $\mathsf{P}A = A^\mathsf{T}\mathsf{P}$ gives $\boldsymbol{w}_j^\mathsf{T} A = u_j \boldsymbol{w}_j^\mathsf{T}$. (ii) The combination $\boldsymbol{w}_j^\mathsf{T} g$ commutes with the brackets, with ∂_z and with the linear operations, and $\boldsymbol{w}_j^\mathsf{T} A = u_j \boldsymbol{w}_j^\mathsf{T}$ turns $s\nabla_\parallel(Ag)$ into $s u_j \nabla_\parallel(\boldsymbol{w}_j^\mathsf{T} g)$; the converse follows from (iii). (iii) By (i), $R^\mathsf{T}\mathsf{P}R = I$, so $g = Ra$ and $\sum_m p_m |g_m|^2 = a^\dagger R^\mathsf{T}\mathsf{P}R\,a = \sum_j |a_j|^2$. □

C.2 Gauss quadrature

We prove proposition 2. For the basis functions e_m, $\mathcal{T}[e_m](u) = [F_u, e_m] = p_m\,r_m(u)$. (i) By theorem 4 the slice transform is unitary, and the slice amplitude of e_m is $\sqrt{\varrho}\,p_m r_m$, so $p_m \delta_{mn} = [e_m, e_n] = \int \varrho\,p_m r_m\,p_n r_n\,\mathrm{d}u$. (ii) The polynomials $\sqrt{p_m}\,r_m$ are therefore orthonormal for $\varrho\,\mathrm{d}u$, with the Jacobi matrix $\mathsf{P}^{1/2} A\,\mathsf{P}^{-1/2}$ by (C2), and the zero closure keeps its first $M+1$ rows and columns. By the theorem of Golub and Welsch (1969), the eigenvalues of this truncated matrix are the zeros of r_{M+1} and the nodes of the Gauss rule, and the weights are $\mu_0 O_{0j}^2$, with O its orthonormal eigenvectors and $\mu_0 = \int \varrho\,\mathrm{d}u = 1/p_0$. Since $O_{mj} = \sqrt{p_m}\,R_{mj}$, the weights are $R_{0j}^2 = 1/n_j$, and exactness to degree $2M+1$ is Gauss's theorem. (iii) By definition $\boldsymbol{r}_j$ holds the Hermite coefficients of F_{u_j}, and $\mathcal{T}[G](u_j) = \sum_m g_m p_m r_m(u_j) = \boldsymbol{w}_j^\mathsf{T} g = \sqrt{n_j}\,a_j$.

C.3 The largest phase velocity

For the zero closure $A = J + \delta A$ exactly, with $\delta A = -(\Lambda\sqrt{2})^{-1}\boldsymbol{e}_1\boldsymbol{e}_0^\mathsf{T}$ and $\boldsymbol{e}_m$ the mth unit vector. The eigenvalues of J are the zeros x_j of h_{M+1}, and its normalised eigenvectors have components $(\lambda_j^\infty)^{1/2} h_m(x_j)$, with λ_j^∞ the Gauss–Hermite weights. Since $h_0 = 1$ and $h_1(x) = \sqrt{2}\,x$, to first order $u_j = x_j - \lambda_j^\infty x_j/\Lambda + O(\Lambda^{-2})$. The weight λ_M^∞ is of order $e^{-x_M^2}$ up to an algebraic factor, with $x_M^2 = 2M + 3 - 2^{2/3}|\alpha_1|(2M+3)^{1/3} + \ldots$, so the shift of $u_{\max}$ is about 10^{-49} at $M = 64$. Up to such

terms, $u_{\max}$ is the largest zero of ψ_{M+1} for the zero closure and the largest root of $\psi_{M+1} = \psi_{M-1}$ for the copy closure.

The function ψ_{M+1} solves $\psi'' + (q^2 - x^2)\psi = 0$ with $q = q_{M+1}$. With $x = q + (2q)^{-1/3}\xi$ and $\sigma = (2q)^{-4/3}$, the uniform Airy (Langer) approximation $\psi \propto (\chi')^{-1/2}\mathrm{Ai}(\chi)$, with $\chi = \xi + \sigma\xi^2/5 + O(\sigma^2)$, places the largest zero at $\chi = \alpha_1$, which gives the first line of (28), the classical expansion of Szegő (1975, theorem 6.32) and Olver (1959) (DLMF, 2026, eq. 12.11.9). For the copy closure, $x\psi_M = \sqrt{M/2}\,\psi_{M-1} + \sqrt{(M+1)/2}\,\psi_{M+1}$ and $\psi'_M = \sqrt{M/2}\,\psi_{M-1} - \sqrt{(M+1)/2}\,\psi_{M+1}$ turn row M, $\psi_{M+1} = \psi_{M-1}$, into the Robin condition $\psi'_M + x(\sqrt{M}+\sqrt{M+1})^{-2}\psi_M = 0$ on ψ_M, which solves the same equation with $q = q_M$ and to leading order reads $\psi_\xi + \sigma\psi = 0$. The same approximation turns this condition into

$$\mathrm{Ai}'(\chi) + \tfrac{4}{5}\,\sigma\,\mathrm{Ai}(\chi) = 0, \qquad \text{so that} \qquad \chi = \alpha'_1 - \frac{4\sigma}{5\alpha'_1} + O(\sigma^2),$$

and $\xi = \chi - \sigma\chi^2/5$ gives the second line of (28); $u_{\max}$ up to $M = 16384$ confirms both lines.

D Numerical methods

D.1 Energy accounting

Each step of the simulations of section 7 consists of GANDALF's Lawson-RK4 step with dissipation and hypercollisions, followed by the forcing increment. The hyper-resistivity and the hypercollisions are linear and diagonal in Fourier and Hermite space, so the free energy they remove from each moment during a step follows from the post-step fields and their damping factors, and the energy before dissipation is reconstructed exactly. The difference between that energy and the energy at the start of the step is the change due to the ideal part of the step, which the exact dynamics conserves; summed over the window it is the time-stepping error of table 1. The injected energy of each step is the change across the forcing increment. This gives $D_\perp(m)$, $D_\nu(m)$ and the injection rate in (21) per moment without finite differences in time, and the field-aligned flux is computed from (20) with GANDALF's dealiased bracket.

D.2 Field-line tracing

Field lines of $\hat{\boldsymbol{b}} = \hat{\boldsymbol{z}} + \hat{\boldsymbol{z}} \times \nabla_\perp \Psi$ are integrated in z with the classical fourth-order Runge–Kutta method and step $\Delta z = L/128$ ($L/256$ at 256^3), with $\nabla_\perp \Psi$ interpolated by cubic B-splines on a grid refined by a factor of two by zero padding. We start 4096 lines at random points of the plane $z = 0$ and follow each over four box lengths. Along each line we sample the field of interest and compute its spectrum in the arc length, which equals z to first order in $b_\perp$. The mean field-aligned wavenumber κ_1 is the first moment of $|\kappa|$ of the spectrum averaged over lines, and we average it over the snapshots of the run. For the injected pattern we draw four fresh realisations of the density forcing with the run's generator and trace them on the same lines; their field-aligned spectrum, averaged over lines, draws and snapshots, is $P(\kappa)$.

D.3 The two-channel model

For one field-aligned wavenumber κ the model is the forced linear hierarchy

$$\frac{\mathrm{d}g}{\mathrm{d}t} = -\mathsf{L}g + \xi\,\boldsymbol{e}_0, \qquad \mathsf{L} = \mathrm{i}s\kappa A + \mathrm{diag}\Big(\frac{1}{2T} + \nu\Big(\frac{m}{M}\Big)^n [m \geq 2]\Big), \tag{D1}$$

where A is the truncated ladder matrix of the run, $\boldsymbol{e}_0$ the unit vector of g_0 and ξ white noise of intensity q. The damping $1/(2T)$ of every amplitude removes free energy from every moment at the rate $1/T$. The stationary covariance $\mathsf{C} = \langle g g^\dagger \rangle$ solves the Lyapunov equation $\mathsf{L}\mathsf{C} + \mathsf{C}\mathsf{L}^\dagger = q\,\boldsymbol{e}_0\boldsymbol{e}_0^\mathsf{T}$. Then $W(m) = p_m \mathsf{C}_{mm}$, the flux is $\Pi_{m+1/2} = -2s\kappa\, p_m A_{m,m+1}\,\mathrm{Im}\,\mathsf{C}_{m+1,m}$, the injection rate is $p_0 q$, and the slice covariance is $\mathsf{R}^\mathsf{T}\mathsf{P}\,\mathrm{Re}\,\mathsf{C}\,\mathsf{P}\mathsf{R}$ with R the matrix of theorem 2. The model spectrum is the sum of these over the measured $P(\kappa)$, normalised so that the total injection rate equals the measured ε. Free energy injected at $\kappa = 0$ does not phase mix and is removed at $m = 0$ at the rate $1/T$. At large m and without hypercollisions, the sum reproduces (44).

D.4 Code and data

The configuration files, the run script and the analysis scripts are in the public repository `https://github.com/anjor/krmhd-research`, in the folders `configs/`, `scripts/` and `analysis/` of `paper/phase-velocity-slices/`; the analysis scripts read only saved data and regenerate every number and figure of section 7. The self-checking scripts behind sections 3 to 5 are in `derivations/` of the same folder. The time series, snapshots and analysis outputs of the runs are available from the author on request.

E Notation

Table 4 compares our notation with that of CMM (Chandran et al., 2026), S09 (Schekochihin et al., 2009) and the GANDALF code (Kanekar, 2025). Three symbols mean different things in the three papers. Our κ is the wavenumber of the compressive fluctuations along the perturbed field, which S09 (§6.3) call $k_{\parallel 0}$; CMM use κ for the square root that we call Δ. Our χ is the nonlinearity parameter, whereas in S09 and CMM χ is the gyrokinetic potential. Our $W(m)$ is the Hermite spectrum of one branch, whereas in S09 and CMM W is the free energy of the whole plasma, which S09 call the generalized energy. S09 also use ζ twice, for the Elsasser potentials $\zeta^{\pm}$ and for the phase velocity ζ_i of a linear mode.

Table 4: Notation of this paper, CMM (Chandran et al., 2026), S09 (Schekochihin et al., 2009) and GANDALF (Kanekar, 2025), with equation or section numbers of the respective papers; "same" means as in this paper.

This paper	CMM	S09	GANDALF
Parameters and velocity			
$v_{\rm th} = (2T_i/m_i)^{1/2}$, $s = v_{\rm th}/v_A = \beta_i^{1/2}$	same	same	$s = \beta_i^{1/2}$
$u = v_\parallel/v_{\rm th}$, also the phase velocity of a slice; $F_M(u) = \pi^{-1/2}e^{-u^2}$	$\hat{v}_\parallel$; F_M in $v_\parallel$ (below B4)	$\zeta_i = \omega/\lvert k_\parallel\rvert v_{\rm th}$ for a linear mode (§6.2.2); F_M as CMM	same
$\Lambda^\pm$ of (6); $\Lambda = \infty$ for $l \geq 2$	same (B7)	same (182)	input
Δ in (6)	κ (B16, printed without the square)	written out (182)	—
$\chi = k_\perp u_\perp/(k_\parallel v_A)$, the nonlinearity parameter	$\chi = \phi - \boldsymbol{v}\cdot\boldsymbol{A}/c$, the gyrokinetic potential (2.7)	χ as CMM; critical balance is our $\chi \sim 1$	—
Alfvénic fields			
Φ, with $\boldsymbol{u}_E = \hat{\boldsymbol{z}}\times\nabla_\perp\Phi$	$\varphi = 2\Phi/(\rho_i v_{\rm th})$ (3.2); $\boldsymbol{u}_\perp = \boldsymbol{u}_E$	$\Phi = c\phi/B_0$ (135); $\boldsymbol{u}_\perp = \boldsymbol{u}_E$	$\Phi = (z^+ + z^-)/2$
$\Psi = -A_\parallel/(4\pi n_i m_i)^{1/2}$	$A = -\Psi/(\rho_i v_A)$ (3.2)	same (135)	$\Psi = (z^+ - z^-)/2$
$\nabla_\parallel = \hat{\boldsymbol{b}}\cdot\nabla = \partial_z + \{\Psi,\cdot\}$, with $\boldsymbol{b}_\perp = \hat{\boldsymbol{z}}\times\nabla_\perp\Psi$	$\partial_z + \boldsymbol{b}_\perp\cdot\nabla_\perp/v_A = \partial_z - \rho_i\{A,\cdot\}$ (3.4)	$\hat{\boldsymbol{b}}\cdot\nabla$	same
$z^\pm = \Phi \pm \Psi$; z^+ propagates toward $-z$	same by (3.23)–(3.24)	$\zeta^\pm$ (21)	same
κ (sections 6 and 7) and $k_\parallel$ (section 5): wavenumber along the perturbed field; k_z: along $\hat{\boldsymbol{z}}$	—	$k_{\parallel 0}$ for the compressive fluctuations, $k_{\parallel A}$ for the Alfvén waves (§6.3)	k_z only
Compressive fields			
g	g_{09} (3.26), not their g_i	same (157)	—
$G^\pm$; G_l for $l \geq 2$	same (B5, B8)	same (183–184, 205)	one branch, set by Λ
$g_m = (2^m m!)^{-1/2}\int H_m(u)\,G\,{\rm d}u$	$\tilde{G}_m^\pm$ (B27, B30)	—	same
Free energy and invariants of one branch			
$W(m) = p_m\langle g_m^2\rangle$, the Hermite spectrum; $\sum_m W(m) = \int\mathcal{E}(u)\,{\rm d}u$	$\sum_m W(m) = 2W^\pm_{\rm compr}/(n_i T_i)$ (B13, B31); W (2.13) is the gyrokinetic free energy	W is the generalized energy (§3.4)	—
$\int u\,\mathcal{E}(u)\,{\rm d}u$	$2\Gamma^\pm/v_{\rm th}$ (B25, B32)	—	—
$\mathcal{H}^\pm \propto \int\mathcal{E}(u)\,{\rm d}u/u$	$H^\pm_{\text{ph-sp}}$ (B18); $\mathcal{H}$ (2.15) is the gyrokinetic helicity	—	—

References

T. Adkins and A. A. Schekochihin. A solvable model of Vlasov-kinetic plasma turbulence in Fourier–Hermite phase space. *J. Plasma Phys.*, 84(1):905840107, 2018. doi: 10.1017/S0022377818000089.

A. Barnes. Collisionless damping of hydromagnetic waves. *Phys. Fluids*, 9(8):1483–1495, 1966. doi: 10.1063/1.1761882.

K. M. Case. Plasma oscillations. *Ann. Phys.*, 7(3):349–364, 1959. doi: 10.1016/0003-4916(59)90029-6.

B. D. G. Chandran, A. Mallet, and R. Meyrand. The second quadratic invariant of astrophysical gyrokinetics. arXiv:2607.27981, 2026.

J. Cho and E. T. Vishniac. The anisotropy of magnetohydrodynamic Alfvénic turbulence. *Astrophys. J.*, 539:273–282, 2000.

DLMF. NIST digital library of mathematical functions. https://dlmf.nist.gov/, Release 1.2.8 of 2026-09-15, 2026. F. W. J. Olver, A. B. Olde Daalhuis, D. W. Lozier, B. I. Schneider, R. F. Boisvert, C. W. Clark, B. R. Miller, B. V. Saunders, H. S. Cohl and M. A. McClain, eds.

T. D. Drivas and G. L. Eyink. A Lagrangian fluctuation–dissipation relation for scalar turbulence. Part I. Flows with no bounding walls. *J. Fluid Mech.*, 829:153–189, 2017.

G. Eyink, E. Vishniac, C. Lalescu, H. Aluie, K. Kanov, K. Bürger, R. Burns, C. Meneveau, and A. Szalay. Flux-freezing breakdown in high-conductivity magnetohydrodynamic turbulence. *Nature*, 497(7450):466–469, 2013. doi: 10.1038/nature12128.

G. L. Eyink. Cascades and dissipative anomalies in nearly collisionless plasma turbulence. *Phys. Rev. X*, 8 (4):041020, 2018. doi: 10.1103/PhysRevX.8.041020.

G. Falkovich, K. Gawędzki, and M. Vergassola. Particles and fields in fluid turbulence. *Rev. Mod. Phys.*, 73 (4):913–975, 2001. doi: 10.1103/RevModPhys.73.913.

B. D. Fried and S. D. Conte. *The Plasma Dispersion Function: The Hilbert Transform of the Gaussian.* Academic Press, New York, 1961. doi: 10.1016/C2013-0-12176-9.

G. H. Golub and J. H. Welsch. Calculation of Gauss quadrature rules. *Math. Comp.*, 23(106):221–230, 1969. doi: 10.1090/S0025-5718-69-99647-1.

J. M. Heninger and P. J. Morrison. An integral transform technique for kinetic systems with collisions. *Phys. Plasmas*, 25(8):082118, 2018. doi: 10.1063/1.5046194.

A. Kanekar. GANDALF: a hardware-agnostic spectral solver for kinetic reduced MHD turbulence. arXiv:2511.21891, 2025.

A. Kanekar, A. A. Schekochihin, W. Dorland, and N. F. Loureiro. Fluctuation-dissipation relations for a plasma-kinetic Langevin equation. *J. Plasma Phys.*, 81(1):305810104, 2015. doi: 10.1017/S0022377814000622.

A. V. Kanekar. *Phase mixing in turbulent magnetized plasmas.* PhD thesis, University of Maryland, College Park, 2015. URL https://doi.org/10.13016/M20W5H.

T. Kato. *Perturbation Theory for Linear Operators.* Springer, Berlin, Heidelberg, 1966. doi: 10.1007/978-3-642-53393-8.

Y. Kawazura, M. Barnes, and A. A. Schekochihin. Thermal disequilibration of ions and electrons by collisionless plasma turbulence. *Proc. Natl Acad. Sci. USA*, 116:771–776, 2019.

Y. Kawazura, A. A. Schekochihin, M. Barnes, J. M. TenBarge, Y. Tong, K. G. Klein, and W. Dorland. Ion versus electron heating in compressively driven astrophysical gyrokinetic turbulence. *Phys. Rev. X*, 10(4):041050, 2020. doi: 10.1103/PhysRevX.10.041050.

M. W. Kunz, A. A. Schekochihin, C. H. K. Chen, I. G. Abel, and S. C. Cowley. Inertial-range kinetic turbulence in pressure-anisotropic astrophysical plasmas. *J. Plasma Phys.*, 81(5):325810501, 2015. doi: 10.1017/S0022377815000811.

J. D. Lawson. Generalized Runge–Kutta processes for stable systems with large Lipschitz constants. *SIAM J. Numer. Anal.*, 4(3):372–380, 1967. doi: 10.1137/0704033.

Y. Lithwick and P. Goldreich. Compressible magnetohydrodynamic turbulence in interstellar plasmas. *Astrophys. J.*, 562(1):279–296, 2001. doi: 10.1086/323470.

R. Meyrand, A. Kanekar, W. Dorland, and A. A. Schekochihin. Fluidization of collisionless plasma turbulence. *Proc. Natl Acad. Sci. USA*, 116(4):1185–1194, 2019. doi: 10.1073/pnas.1813913116.

P. J. Morrison. Hamiltonian description of Vlasov dynamics: action-angle variables for the continuous spectrum. *Transp. Theory Stat. Phys.*, 29(3–5):397–414, 2000. doi: 10.1080/00411450008205881.

P. J. Morrison and D. Pfirsch. Dielectric energy versus plasma energy, and Hamiltonian action-angle variables for the Vlasov equation. *Phys. Fluids B*, 4(10):3038–3057, 1992. doi: 10.1063/1.860415.

F. W. J. Olver. Uniform asymptotic expansions for Weber parabolic cylinder functions of large orders. *J. Res. Natl Bur. Stand. B*, 63B(2):131–169, 1959. doi: 10.6028/jres.063B.014.

J. T. Parker and P. J. Dellar. Fourier–Hermite spectral representation for the Vlasov–Poisson system in the weakly collisional limit. *J. Plasma Phys.*, 81(2):305810203, 2015. doi: 10.1017/S0022377814001287.

J. T. Parker, E. G. Highcock, A. A. Schekochihin, and P. J. Dellar. Suppression of phase mixing in drift-kinetic plasma turbulence. *Phys. Plasmas*, 23(7):070703, 2016. doi: 10.1063/1.4958954.

O. Penrose. Electrostatic instabilities of a uniform non-Maxwellian plasma. *Phys. Fluids*, 3(2):258–265, 1960. doi: 10.1063/1.1706024.

G. G. Plunk. On the nonlinear stability of a quasi-two-dimensional drift kinetic model for ion temperature gradient turbulence. *Phys. Plasmas*, 22(4):042305, 2015. doi: 10.1063/1.4917233.

G. G. Plunk and J. T. Parker. Irreversible energy flow in forced Vlasov dynamics. *Eur. Phys. J. D*, 68(10): 296, 2014. doi: 10.1140/epjd/e2014-50157-8.

A. A. Schekochihin, S. C. Cowley, W. Dorland, G. W. Hammett, G. G. Howes, E. Quataert, and T. Tatsuno. Astrophysical gyrokinetics: kinetic and fluid turbulent cascades in magnetized weakly collisional plasmas. *Astrophys. J. Suppl.*, 182(1):310–377, 2009. doi: 10.1088/0067-0049/182/1/310.

A. A. Schekochihin, J. T. Parker, E. G. Highcock, P. J. Dellar, W. Dorland, and G. W. Hammett. Phase mixing versus nonlinear advection in drift-kinetic plasma turbulence. *J. Plasma Phys.*, 82(2):905820212, 2016. doi: 10.1017/S0022377816000374.

B. I. Shraiman and E. D. Siggia. Scalar turbulence. *Nature*, 405(6787):639–646, 2000. doi: 10.1038/35015000.

G. Szegő. *Orthogonal Polynomials*, volume 23 of *American Mathematical Society Colloquium Publications*. American Mathematical Society, Providence, RI, 4th edition, 1975.

E. Tassi. Poisson brackets and truncations in nonlinear reduced fluid models for plasmas. *Physica D*, 437: 133338, 2022. doi: 10.1016/j.physd.2022.133338.

N. G. van Kampen. On the theory of stationary waves in plasmas. *Physica*, 21(6–10):949–963, 1955. doi: 10.1016/S0031-8914(55)93068-8.

G. M. Webb, B. Dasgupta, J. F. McKenzie, Q. Hu, and G. P. Zank. Local and nonlocal advected invariants and helicities in magnetohydrodynamics and gas dynamics I: Lie dragging approach. *J. Phys. A: Math. Theor.*, 47(9):095501, 2014. doi: 10.1088/1751-8113/47/9/095501.

A. Zocco and A. A. Schekochihin. Reduced fluid-kinetic equations for low-frequency dynamics, magnetic reconnection, and electron heating in low-beta plasmas. *Phys. Plasmas*, 18(10):102309, 2011. doi: 10.1063/1.3628639.